

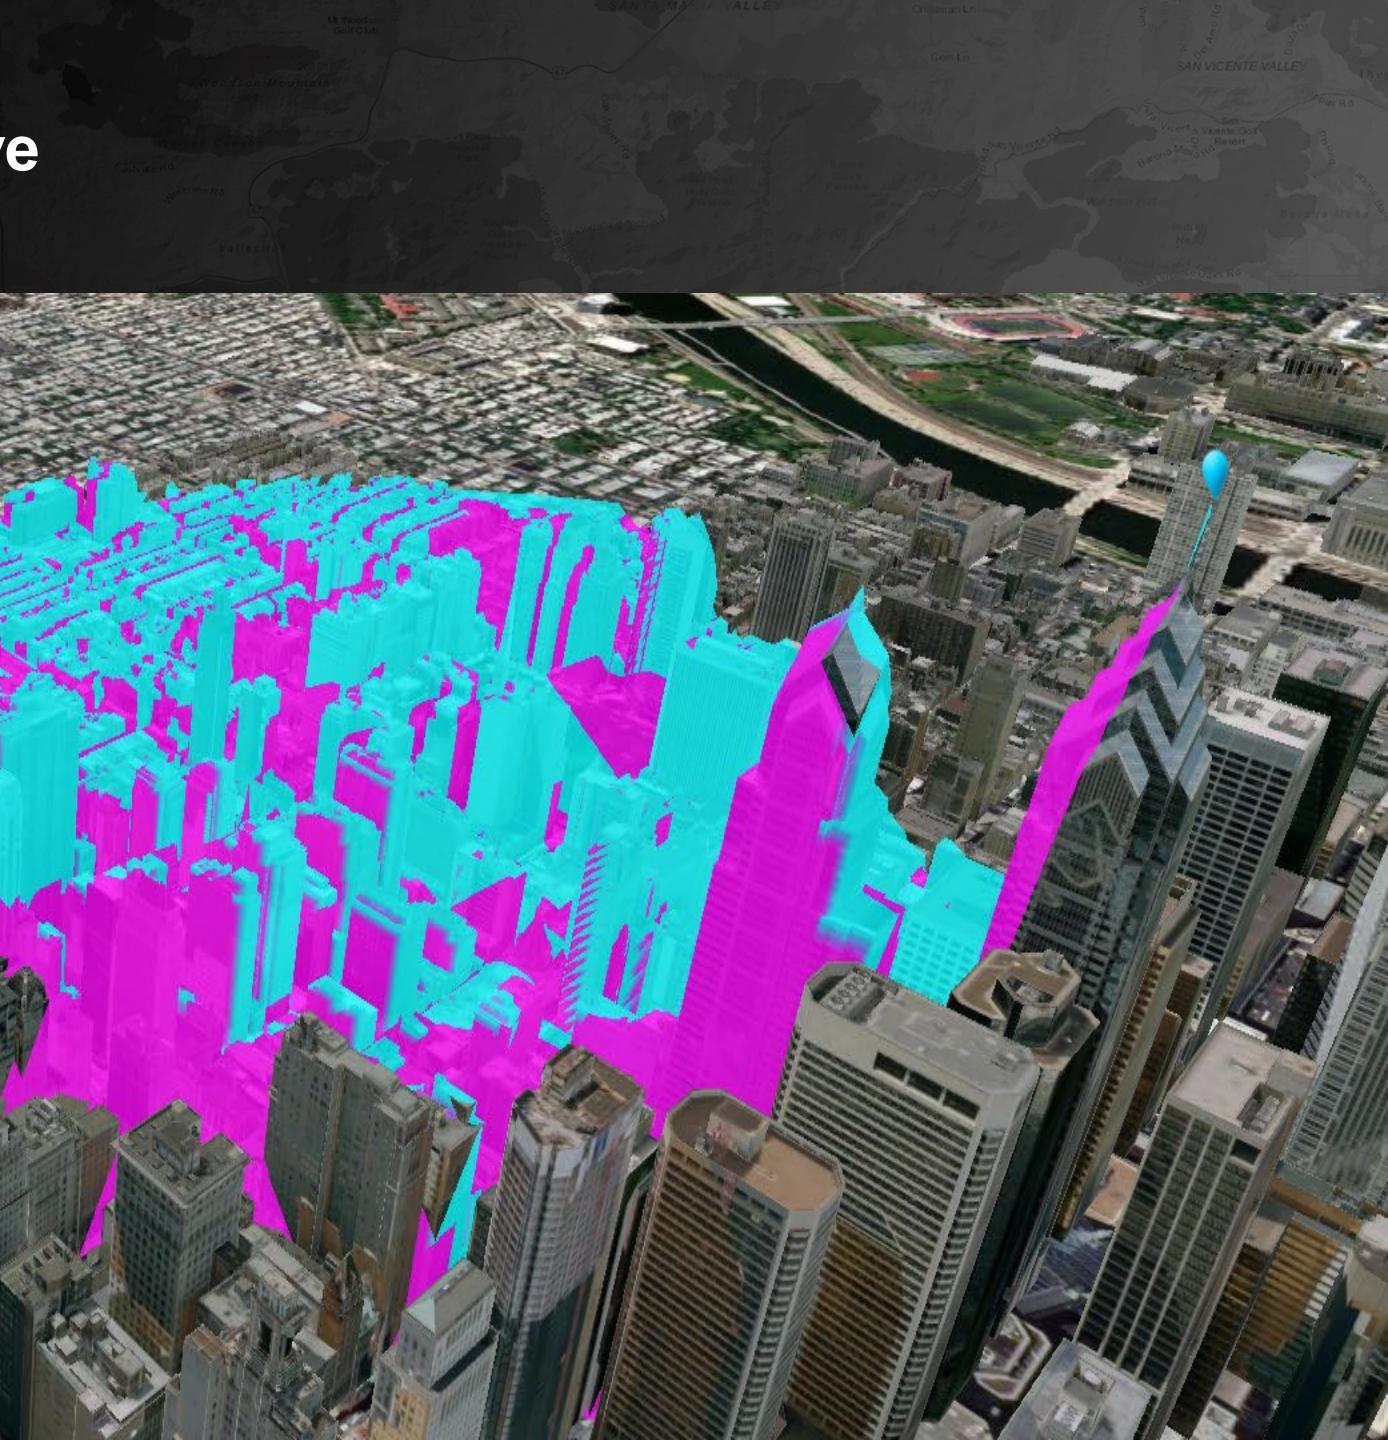


esri | **THE
SCIENCE
OF
WHERE**

Object Extraction from LiDAR with Deep Learning

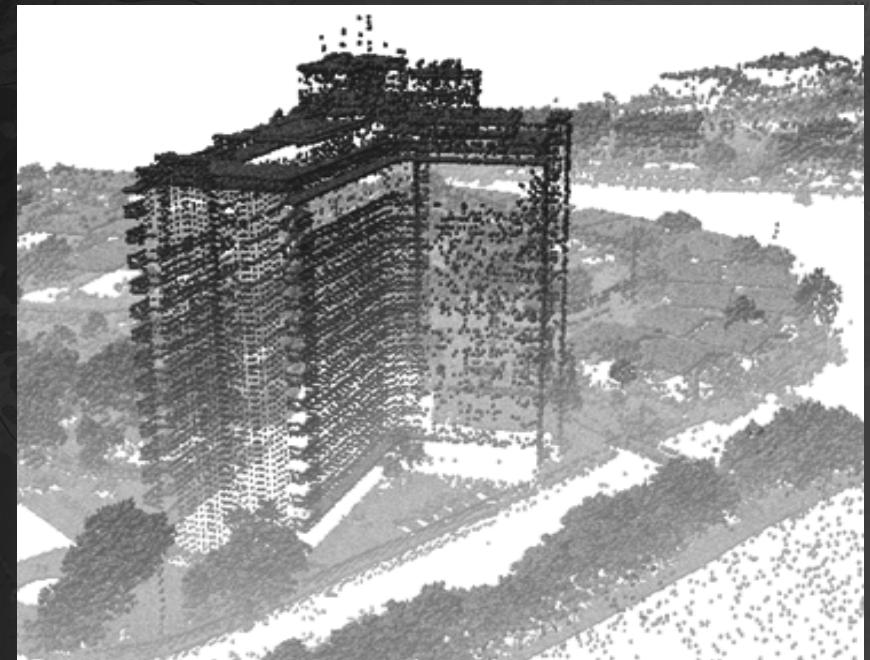
Digital Twin: valuable and expensive

- Digital Twins are **important** for efficient local government operations, urban planning and design, insurance & safety, etc.
- Creation and maintenance of Digital Twin is **expensive** and requires constant inflow of human resources and funds to maintain.
- Traditional raw data sources:
 - LiDAR Point clouds
 - Photogrammetry and resulting **triangulated meshes**



Unclassified point clouds and continuous triangulated meshes

- Both sources require segmentation before they can be used in analysis
- In its raw form, are easily interpretable by humans, but not by algorithms



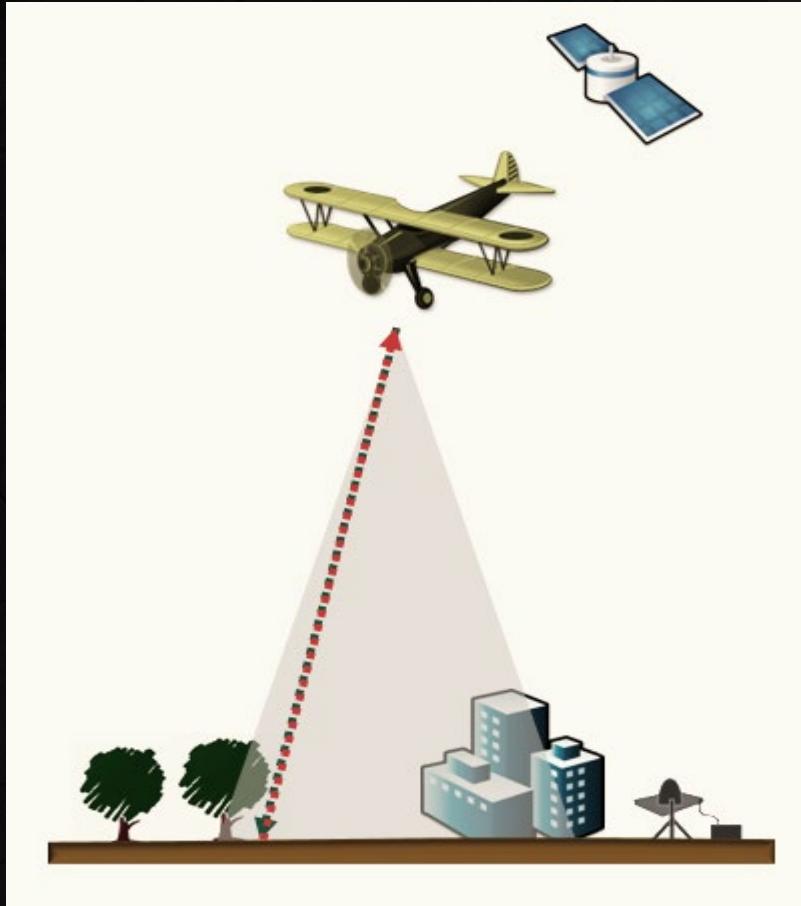
- *How many buildings are in the block?*

- *Estimate # of people in the buildings which are 5+ stores-high on Wednesday from 9-10AM*

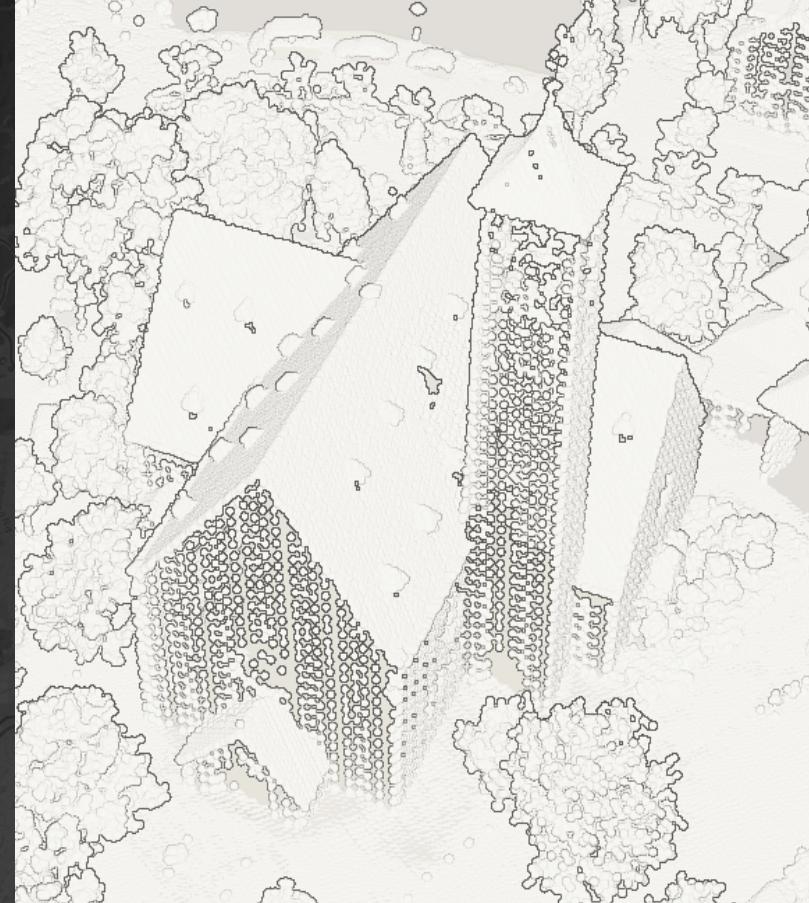
- *What are the windows looking at this part of the park?*

What is LiDAR data?

Massive 3D point collections with additional attributes like:



- Intensity,
- Number of Returns,
- Scan Angle,
- RGB,...

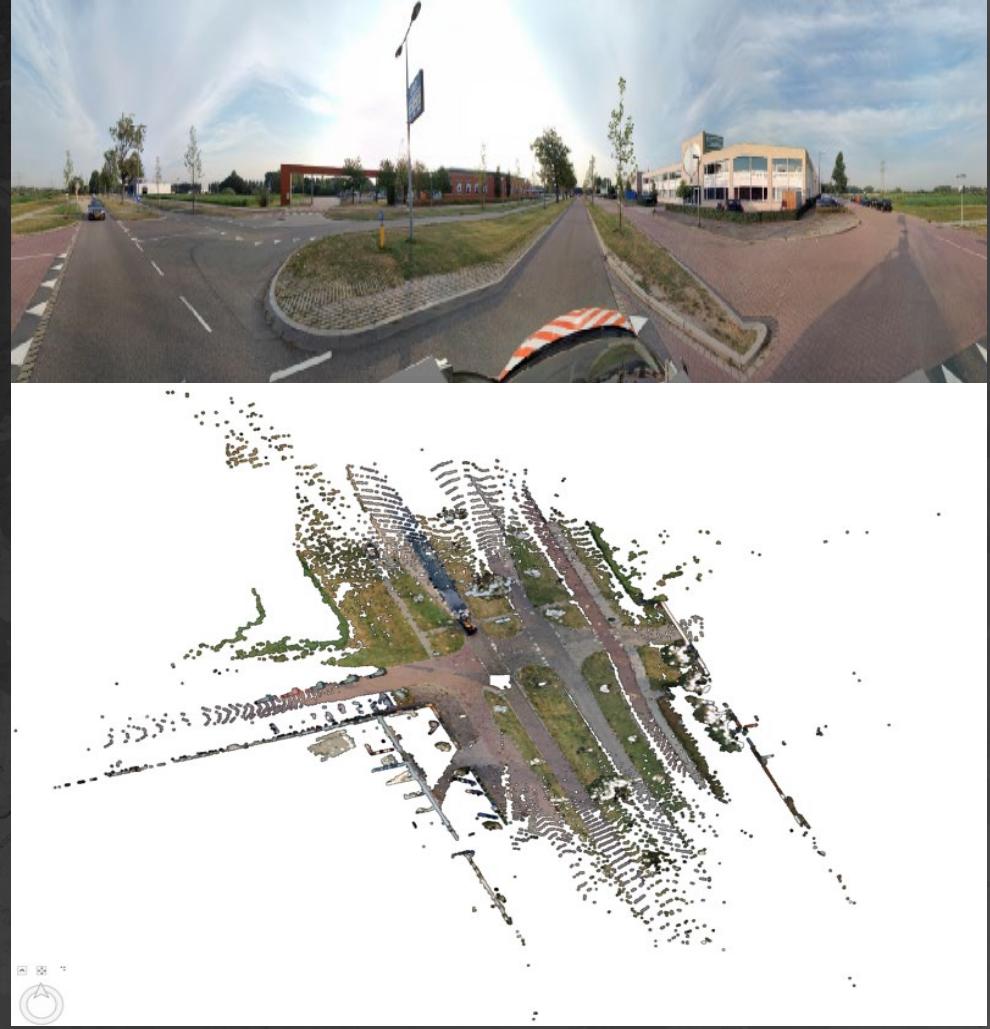


What is LiDAR data?

Massive 3D point collections with additional attributes like:



- Intensity,
- Number of Returns,
- Scan Angle,
- RGB,...

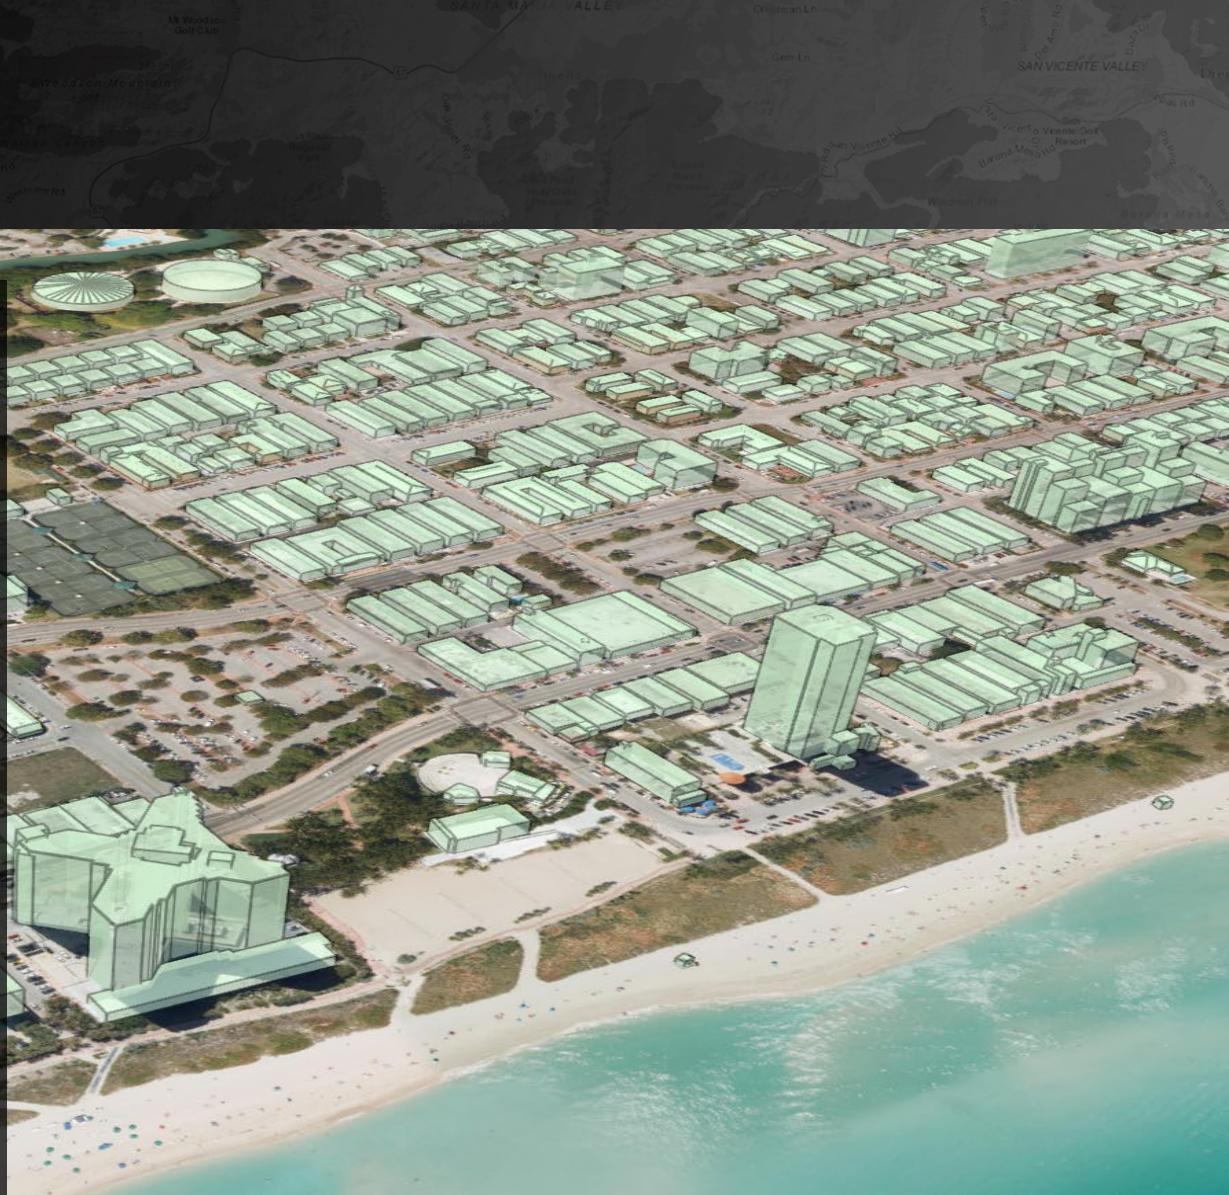


What can be extracted from point clouds?

- 3D building models
 - schematic, or
 - detailed (up to RGB+D textures)
- Electric grid
 - overhead conductors, poles, stay wires, transformers, etc.
- Trees
 - canopy polygons, precise height
- Street furniture
 - Signs, traffic lights, dividers, bus stops, fire hydrants, etc.



3D Building
models
from airborne LiDAR



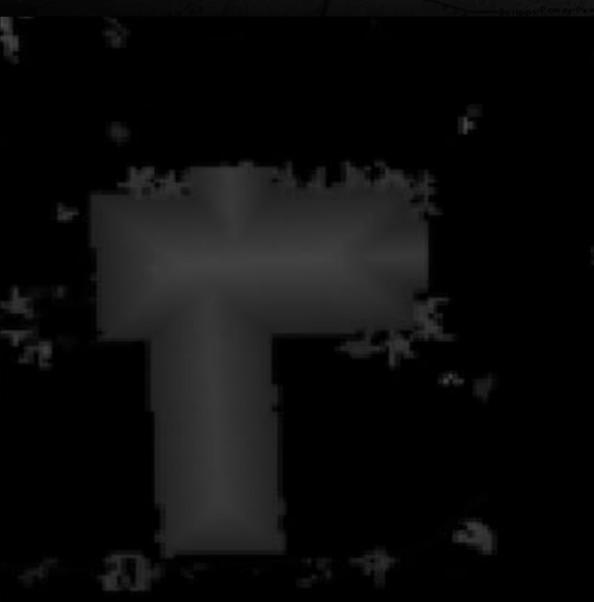
THE
SCIENCE
OF
WHERE

Digital Twins & 3D Building Models: Realism vs Cubism

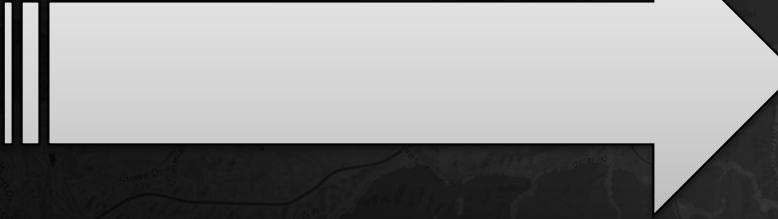
Two directions / types of models:

1. High-fidelity models of historical or culturally-significant buildings, which look is frozen or changes very little in time:
 - Manual model construction
 - Significant initial investments, require little to maintain
2. Schematic models for commercial, industrial, residential zones – zones which are subject to constant changes:
 - Largest in area, number of buildings, number of people hosted
 - Periodic updates are required
 - Data collection and models update should be quick, accurate, cost-efficient

PoC: 3D Building model reconstruction from aerial LiDAR



Rasterized Aerial LiDAR



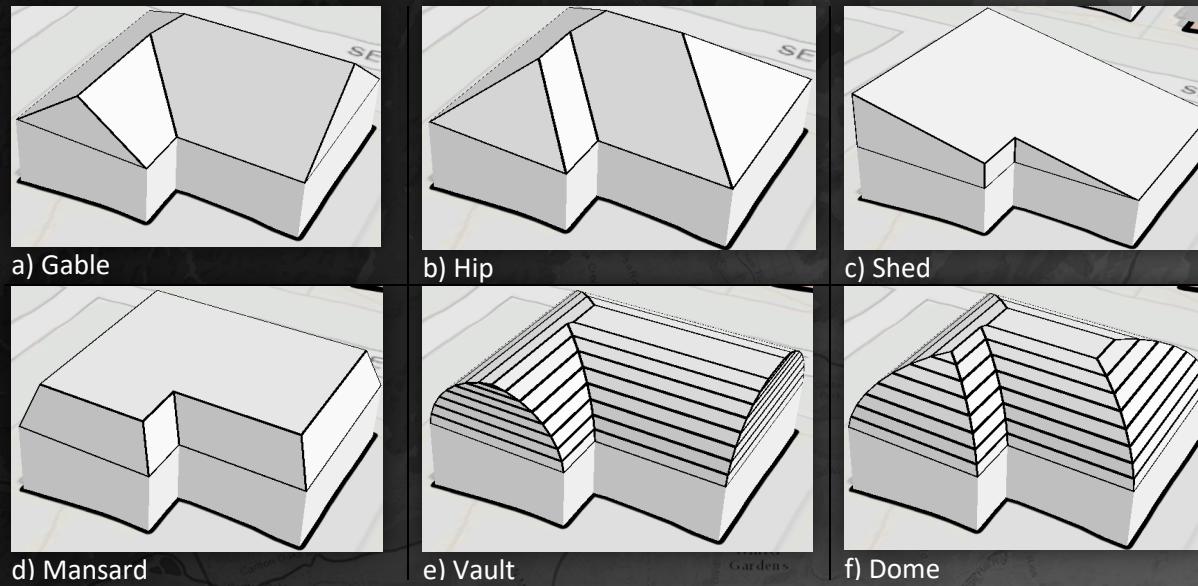
Manually digitized Hip
(purple) and Gable
(orange) segments



3D reconstruction of
building using manually
digitized segments

PoC: 3D Building model reconstruction from aerial LiDAR

- Manually digitizing roof segments:
 - Over 3,000 man hours were spent on digitizing about 213,000 polygons covering the area of 200 square miles.
 - ~71 polygons / man hour.

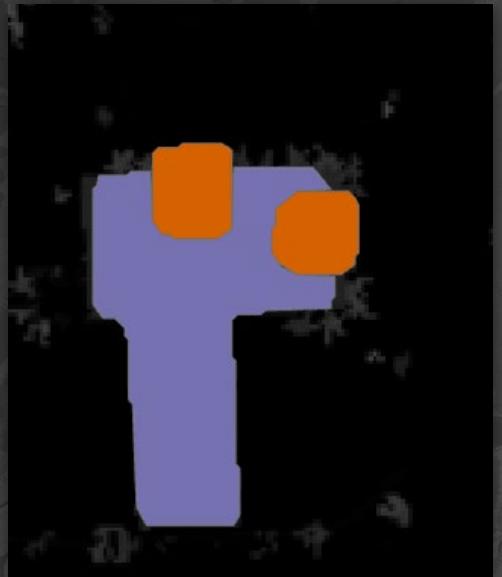


PoC: 3D Building model reconstruction from aerial LiDAR

- Using Mask R-CNN to digitize roof segments
- Not as accurate as humans,
but much faster: 60,000 polygons / hour.
- Regularize Building Footprints helps with
accuracy.



Manually digitized “ground truth” data from the Test set



Prediction produced by the neural network



But can we work in true 3D (without rasterization)?

There is a traditional workflow with deterministic GP tools:

1. **ClassifyLASGround**
2. **ClassifyLASBuilding**

To get building footprints:

3. **LASPointStatisticsAsRaster**
 - with LAS layer filtered on class 6 (building)
 - using the 'Most Frequent Class Code' option
4. **RasterToPolygon**
 - Turn off the Simplify polygons option

5. **EliminatePolygonPart** to remove small holes (could alternately have performed some manipulation on the raster side for this)

6. **RegularizeBuildingFootprint** to straighten things out.

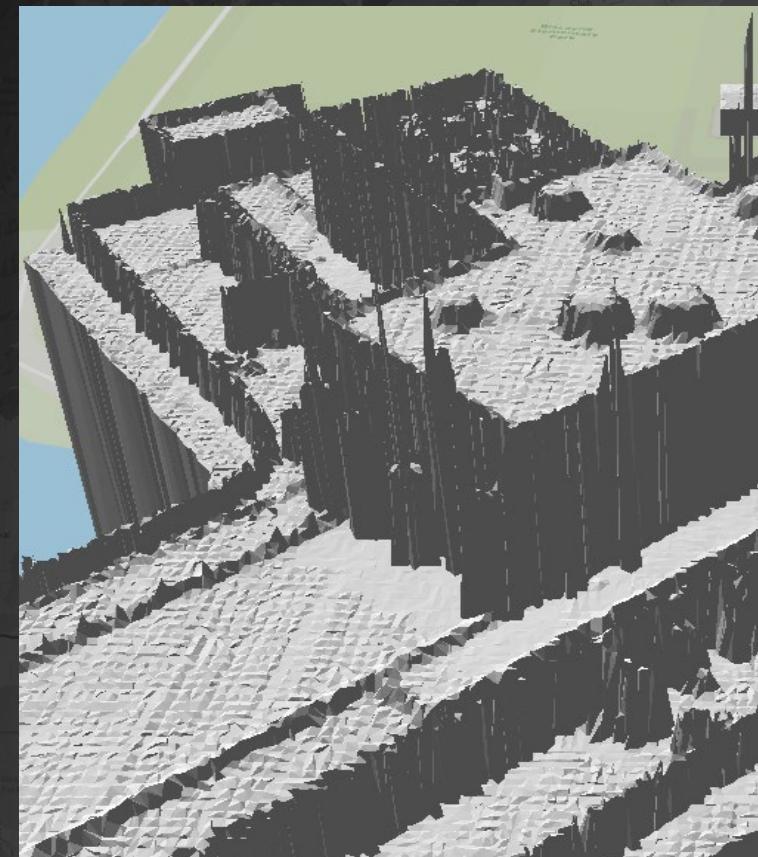
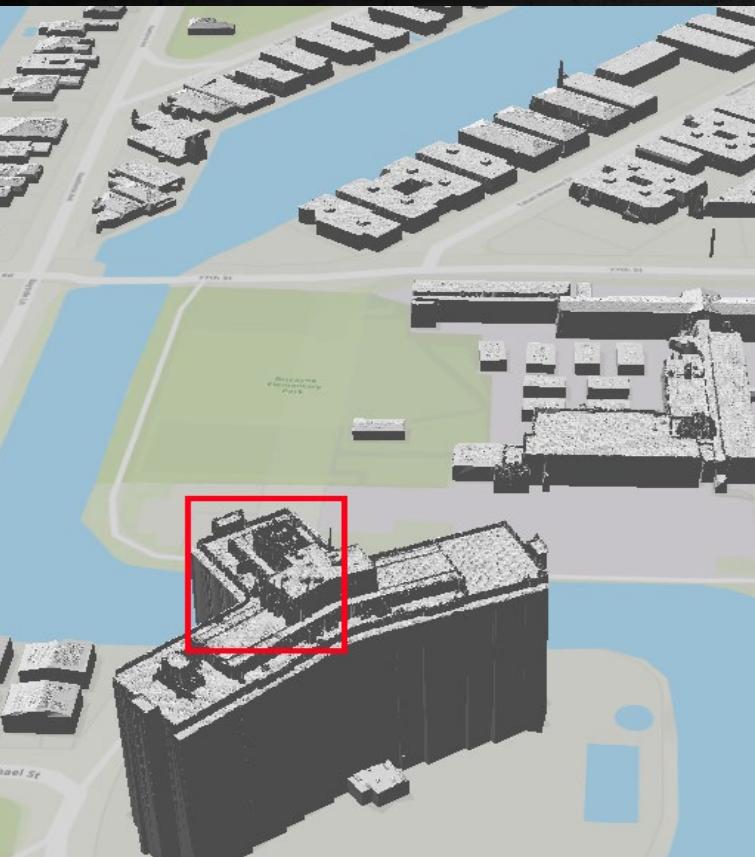
To extract shells:

7. **LASDatasetToRaster** with input LAS layer filtered on class 2 points to make DEM
8. **LASBuildingMultipatch**



Working in true 3D point cloud

- Resulting models are **not** suitable for manual editing because of largen ## of vertices
- Important to **keep the noise level down**



Working in true 3D point cloud

Main sources of noise in resulting models come from misclassification of

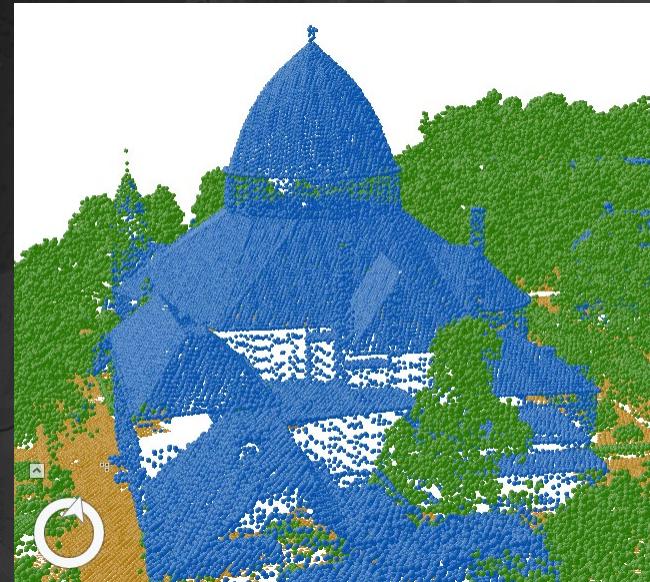
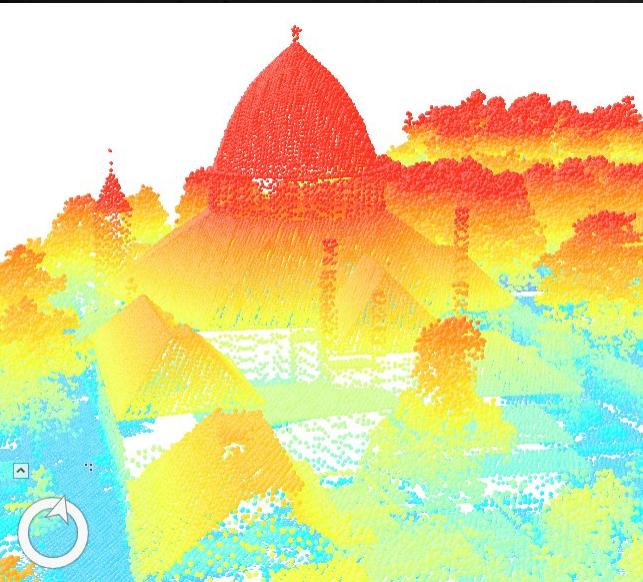
- Ground
- Buildings
- **ClassifyLASBuilding**

does not produce reliable results
when there is a vegetation in direct
proximity to the building



Can we use DL to classify Point Clouds?

- Point clouds are unordered, sparse data collections where direct application of convolutional kernels does not work efficiently.
- Good news in recent years: PointNet, PointNet++, Graph Convolutional networks, Deep Sets, PointCNN, etc.



PointCNN

- Trained on 1.8B X-Y-Z points of airborne point cloud collected over Amsterdam, Netherlands.
- 0.97 accuracy just after 6.5h of training on Nvidia QUADRO V100.
- Testing on city of Utrecht, Netherlands



- Classify Building Points

• PointCNN

PointCNN

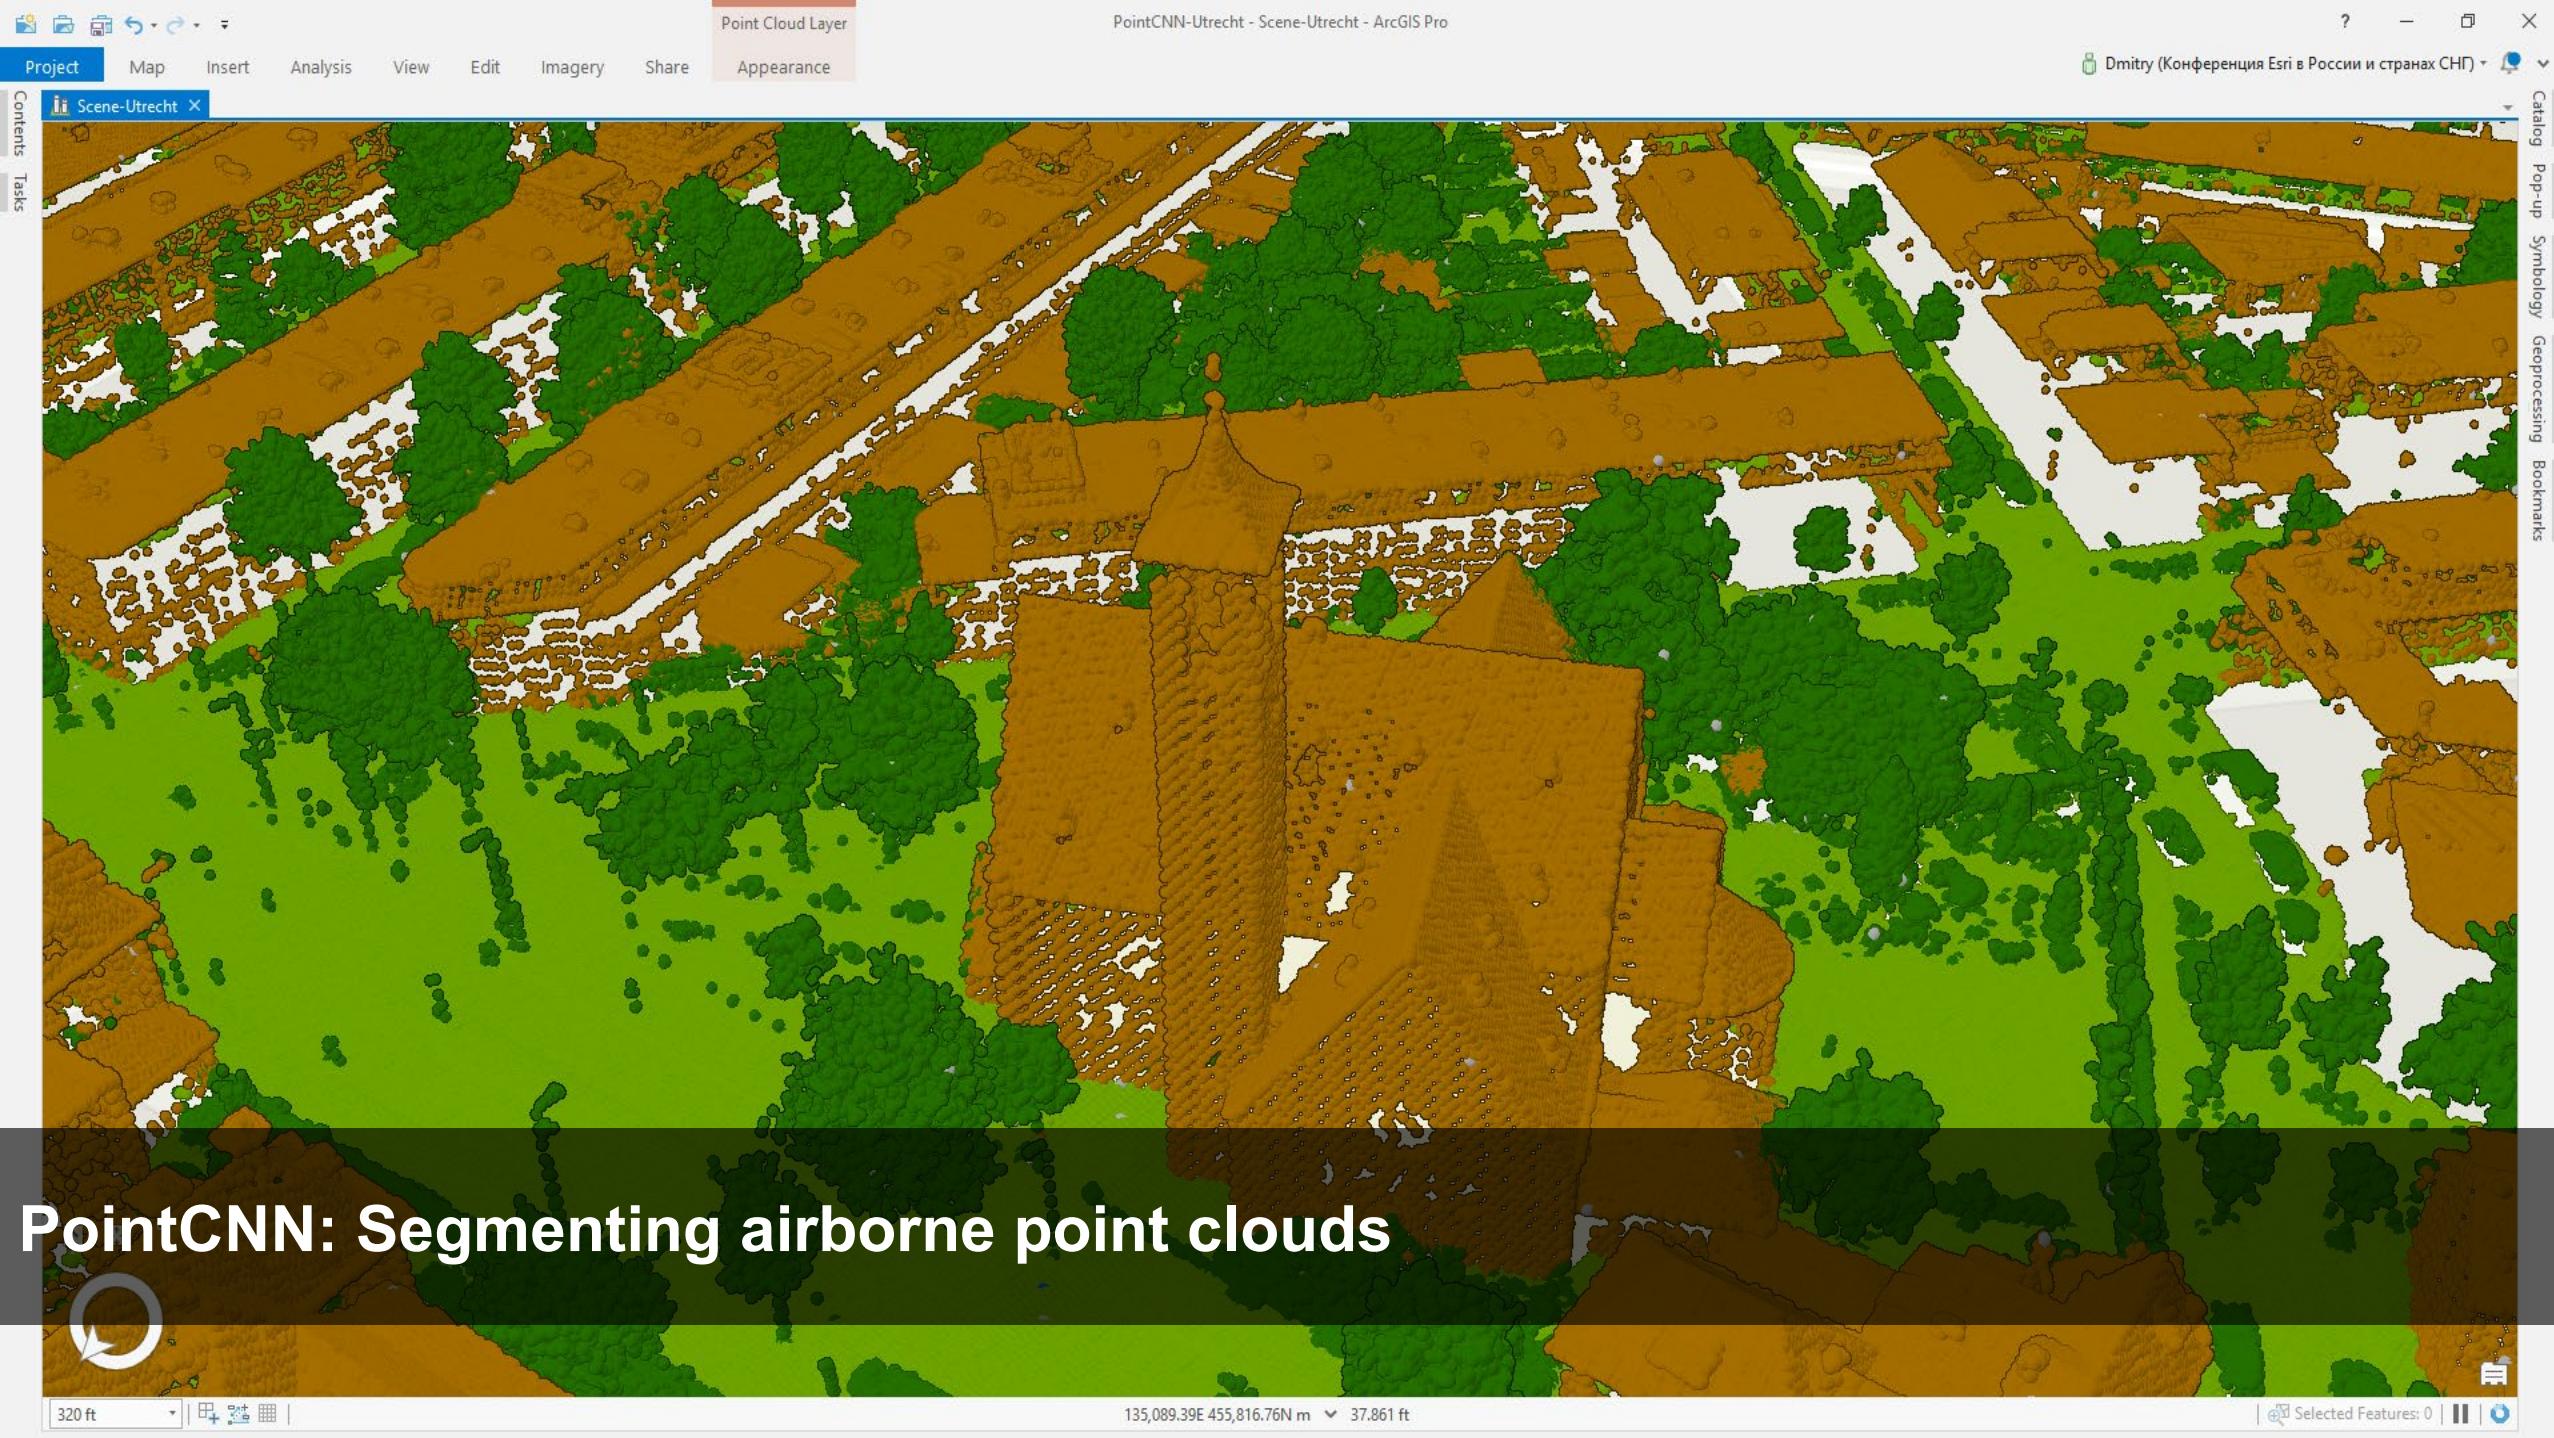
- Trained on 1.8B X-Y-Z points of airborne point cloud collected over Amsterdam, Netherlands.
- 0.97 accuracy just after 6.5h of training on Nvidia QUADRO V100.
- Testing on city of Utrecht, Netherlands



PointCNN vs ClassifyBuildingPoints

- Significantly lower noise level in the models reconstructed from point cloud classified by a PointCNN model.





PointCNN: Wires & Poles

from airborne LiDAR



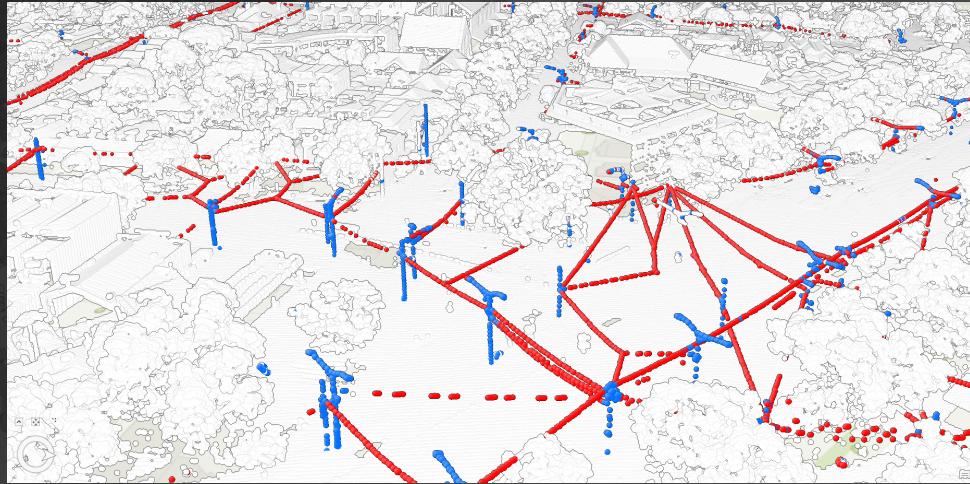
THE
SCIENCE
OF
WHERE

PoC: Wires detection in airborne LiDAR

AAM Group, Australia: collecting airborne LiDAR point clouds to detect the power lines and any easement encroachment.



50,000 man
hours per year of
manual labeling



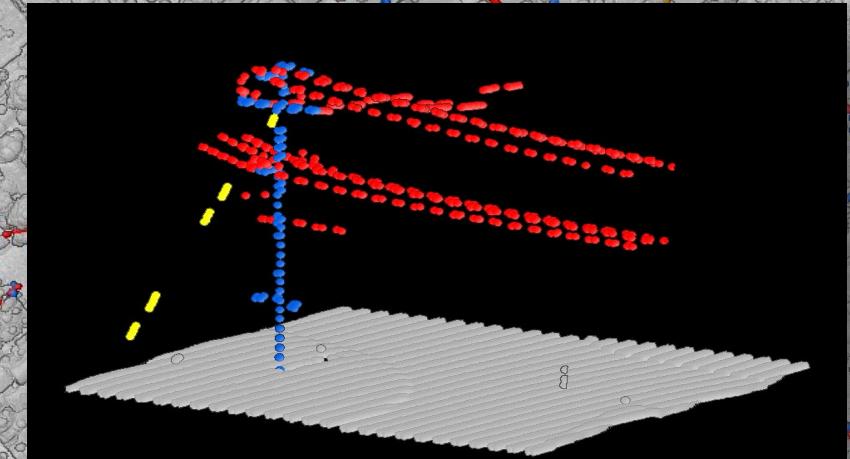
PoC: Wires detection in airborne LiDAR

- AAM Group shared 10B+ of manually labeled points to train a neural network.
- Thousands of miles of transmission lines.



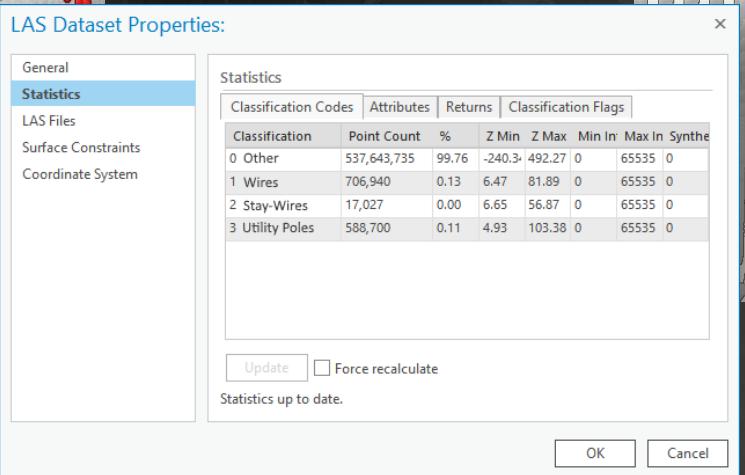
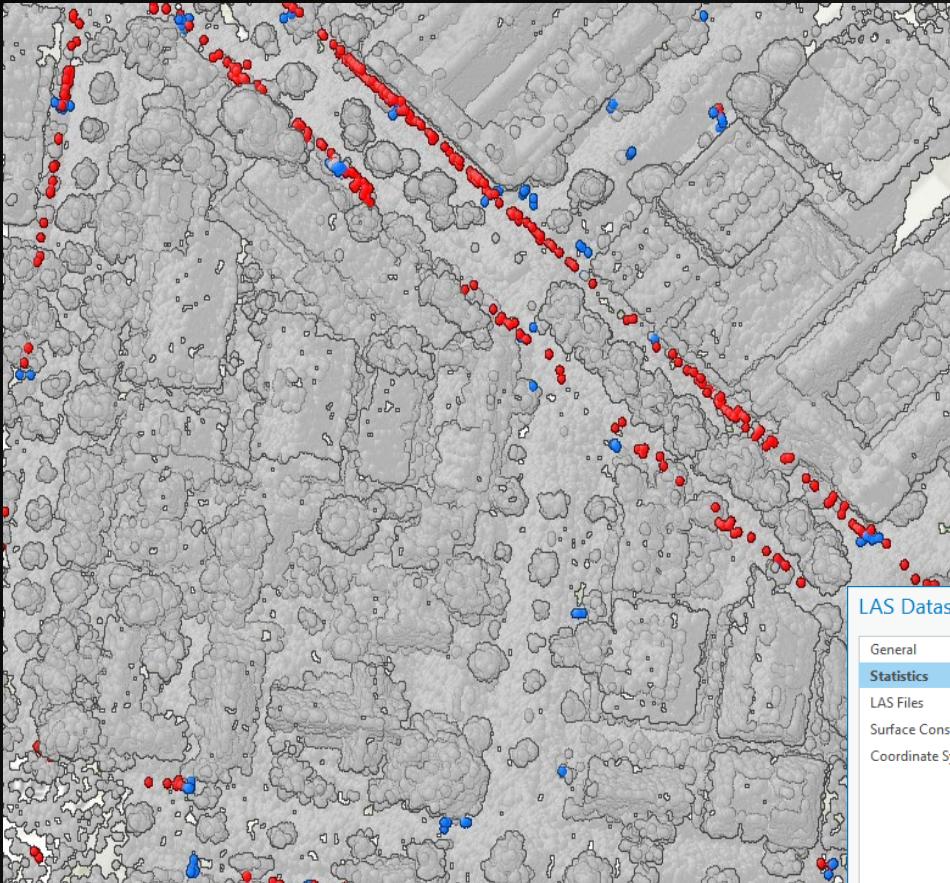
Four object classes:

1. Wire
2. Stay Wire
3. Pole
4. Other



PoC: Wires detection in airborne LiDAR

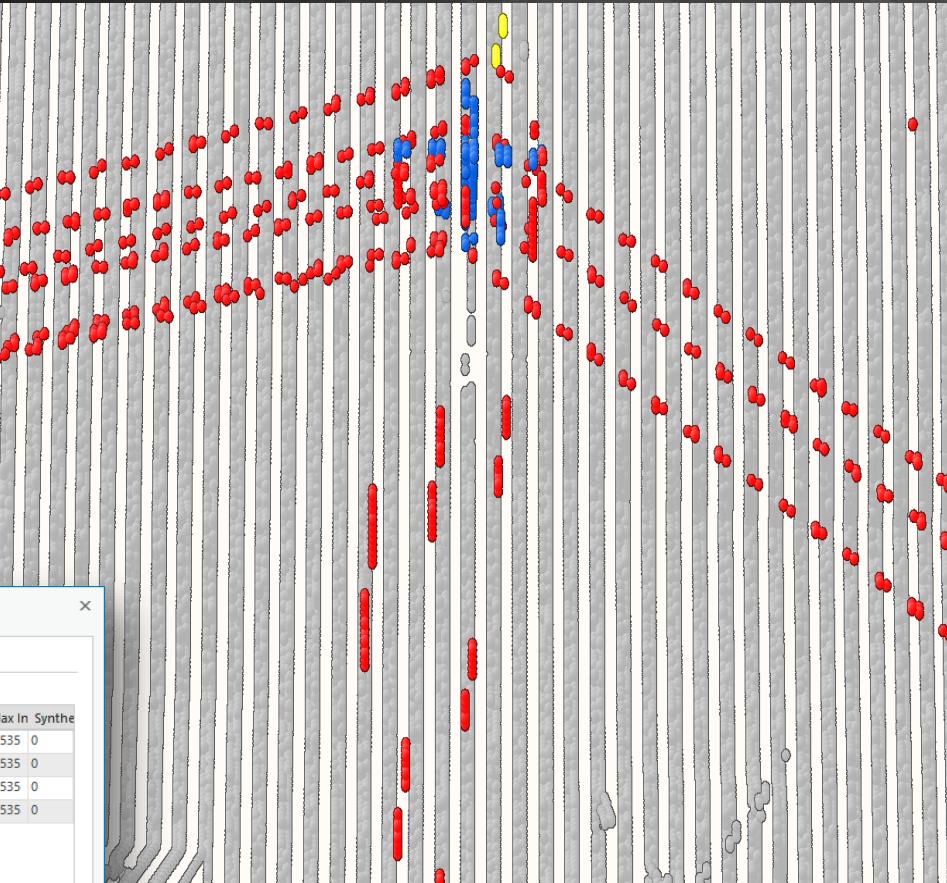
With wires, there challenges though -



Significant class imbalance:

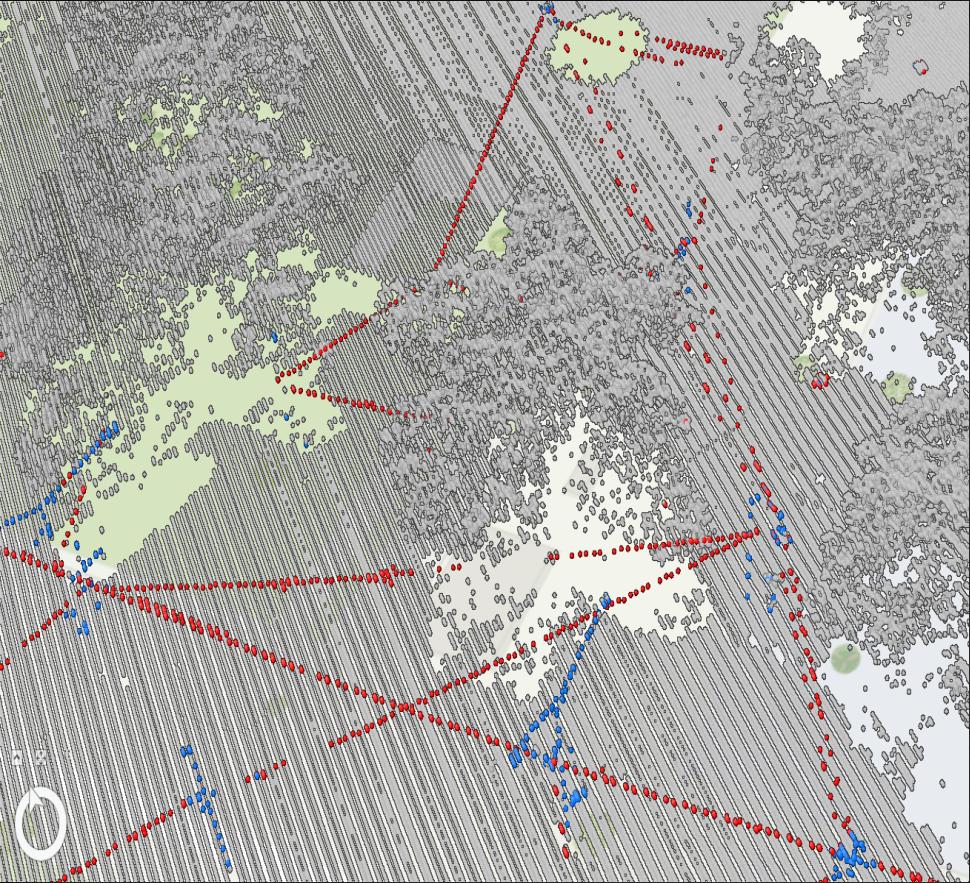
99.76% points are of the Other class.

Scan lines angle causes the Wire point density fluctuation

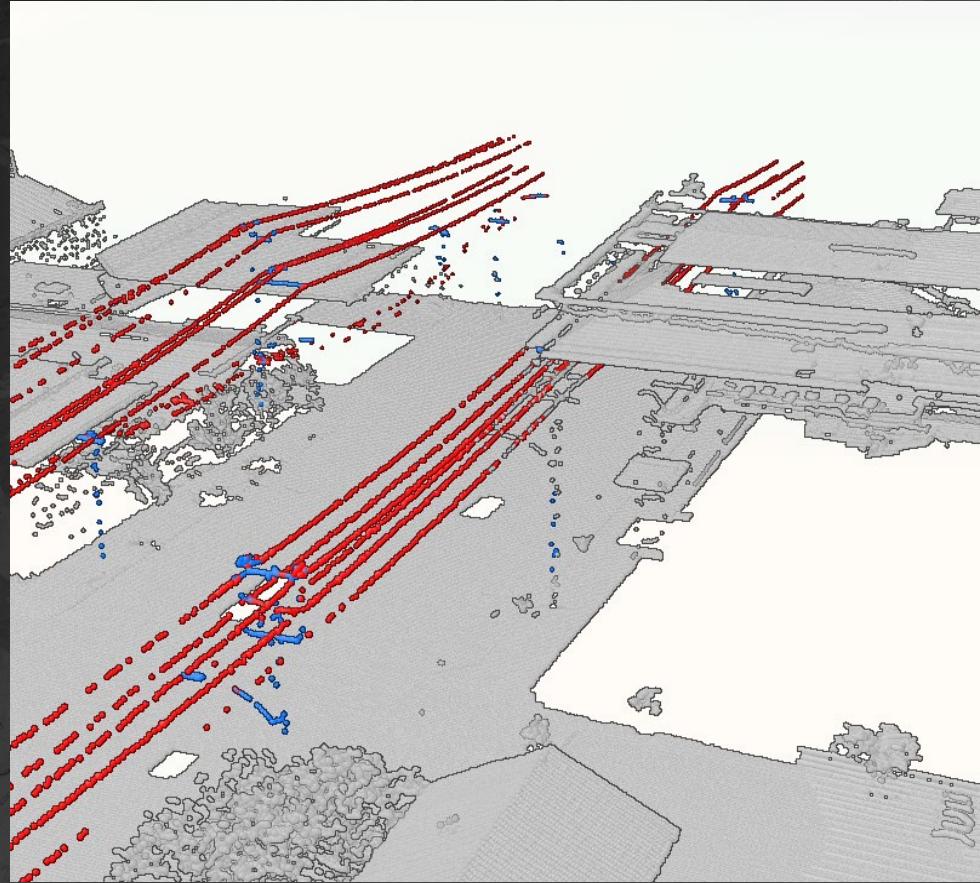


PoC: Overhead conductor inspection

With wires, there challenges though -



Wires are zero-area
point neighborhoods
which easily blend
into tree canopies
and buildings



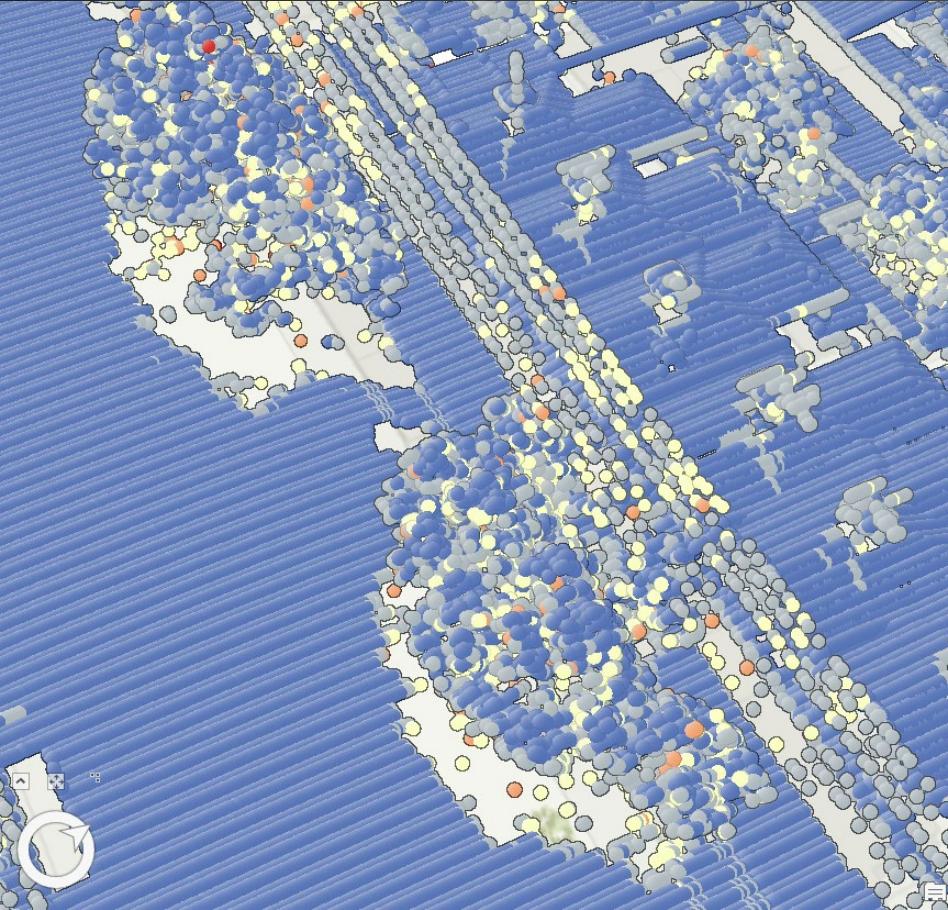
PoC: Overhead conductor inspection

...but there are some good news too -



Intensity and the Number of Returns on the Wire points is often different than of the surroundings

This allows for training PointCNN on
XYZ +
Intensity +
Number of Returns



PoC: Overhead conductor inspection

Training & Results:

after training on a single GV100 for ~20 hours

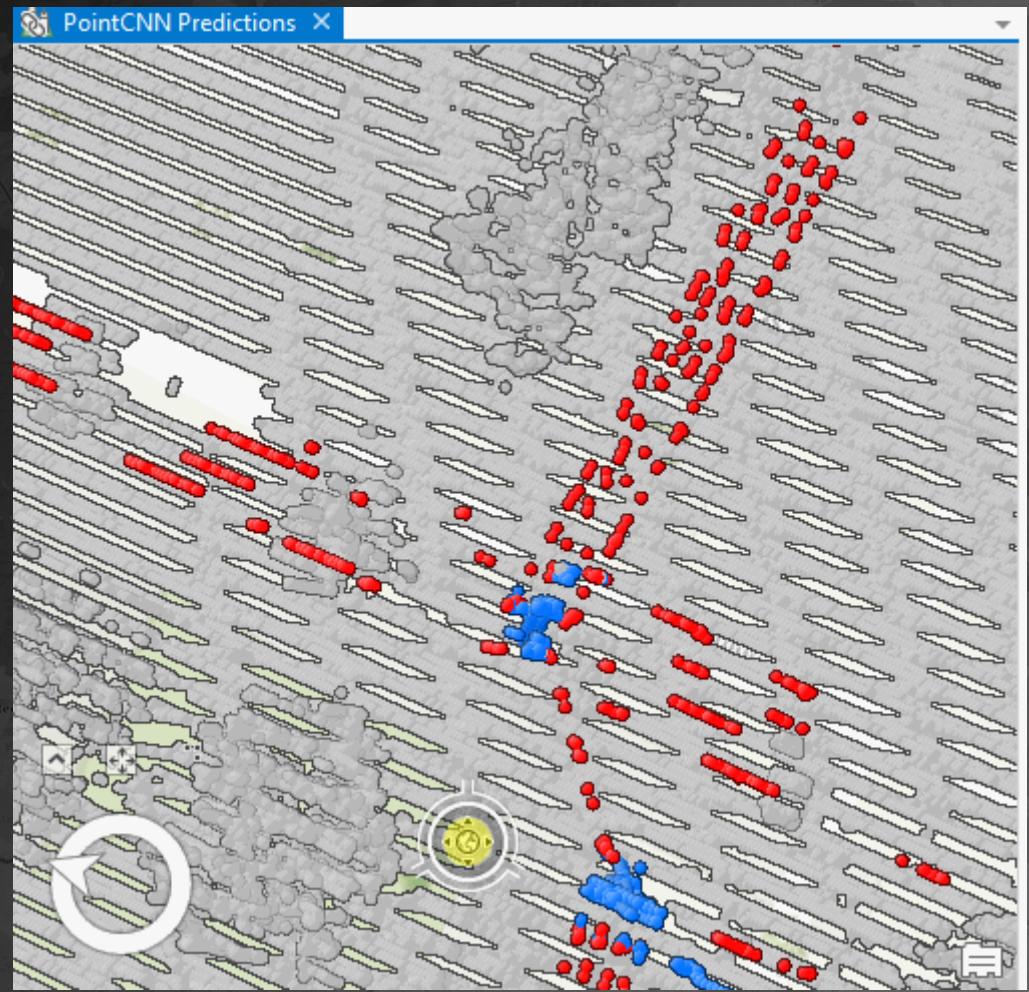
@E232K best RECALL

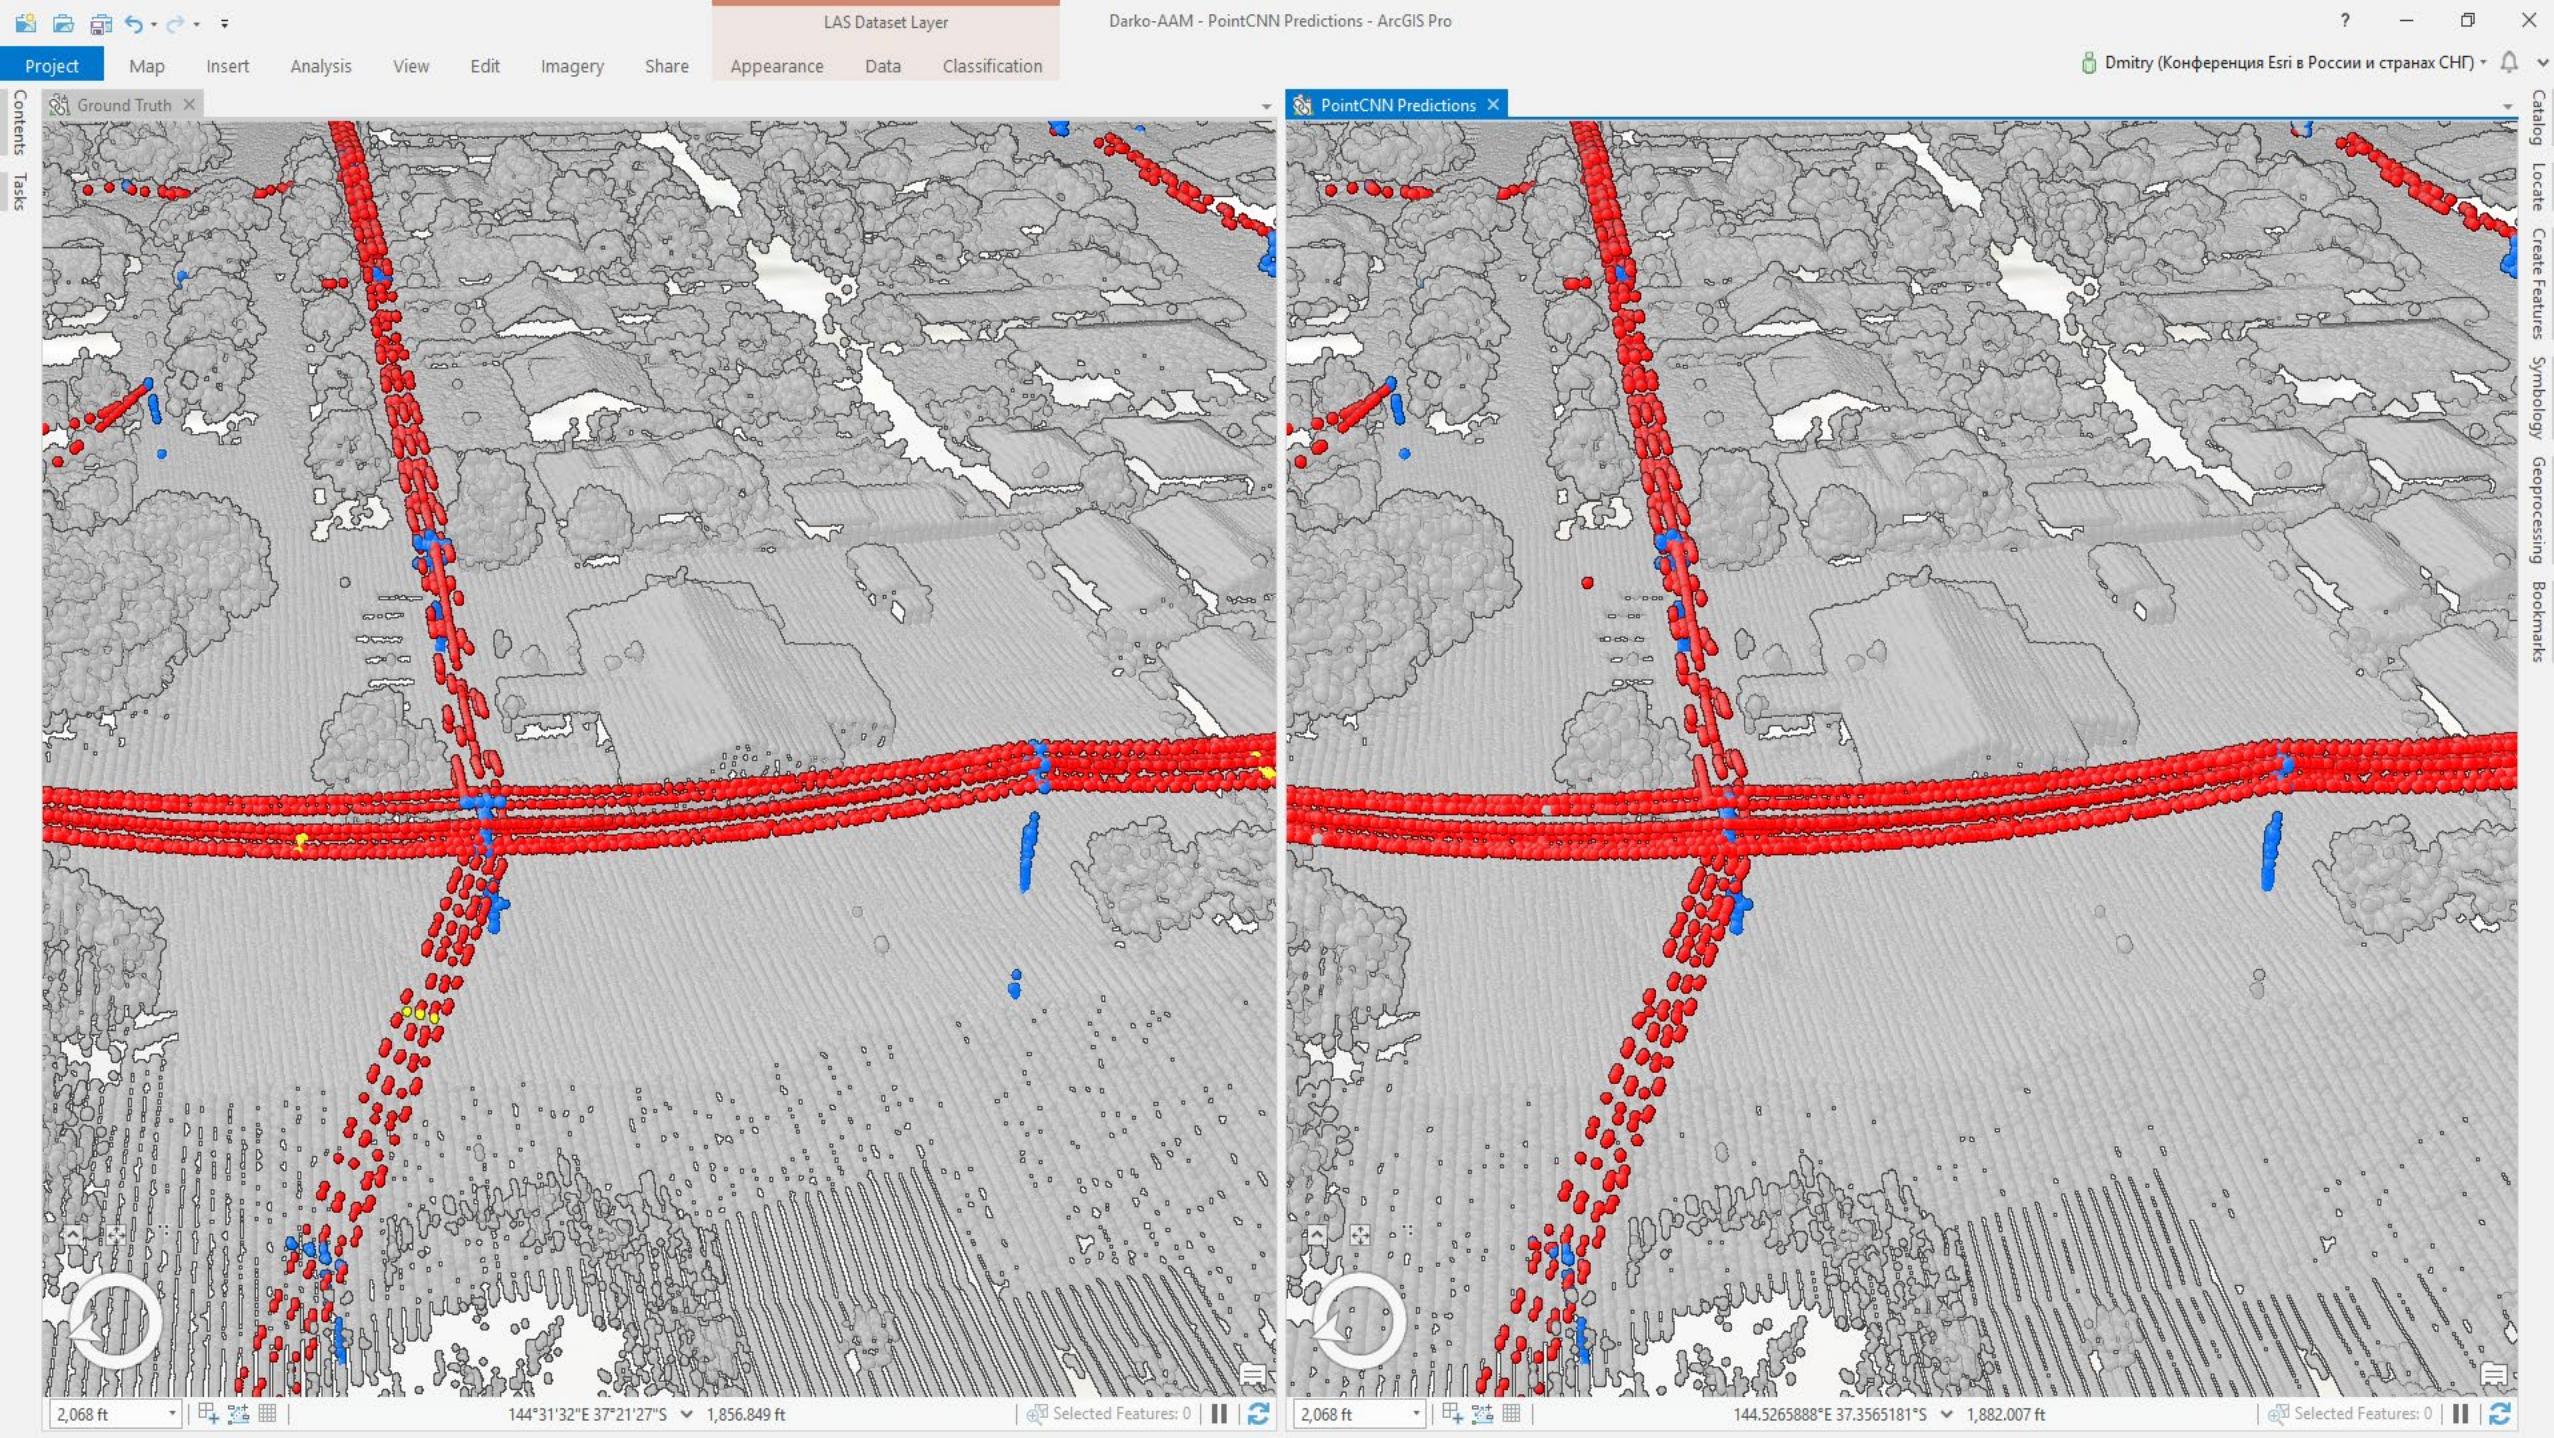
OTHER	WIRE	STAY-WIRE	POLE
-------	------	-----------	------

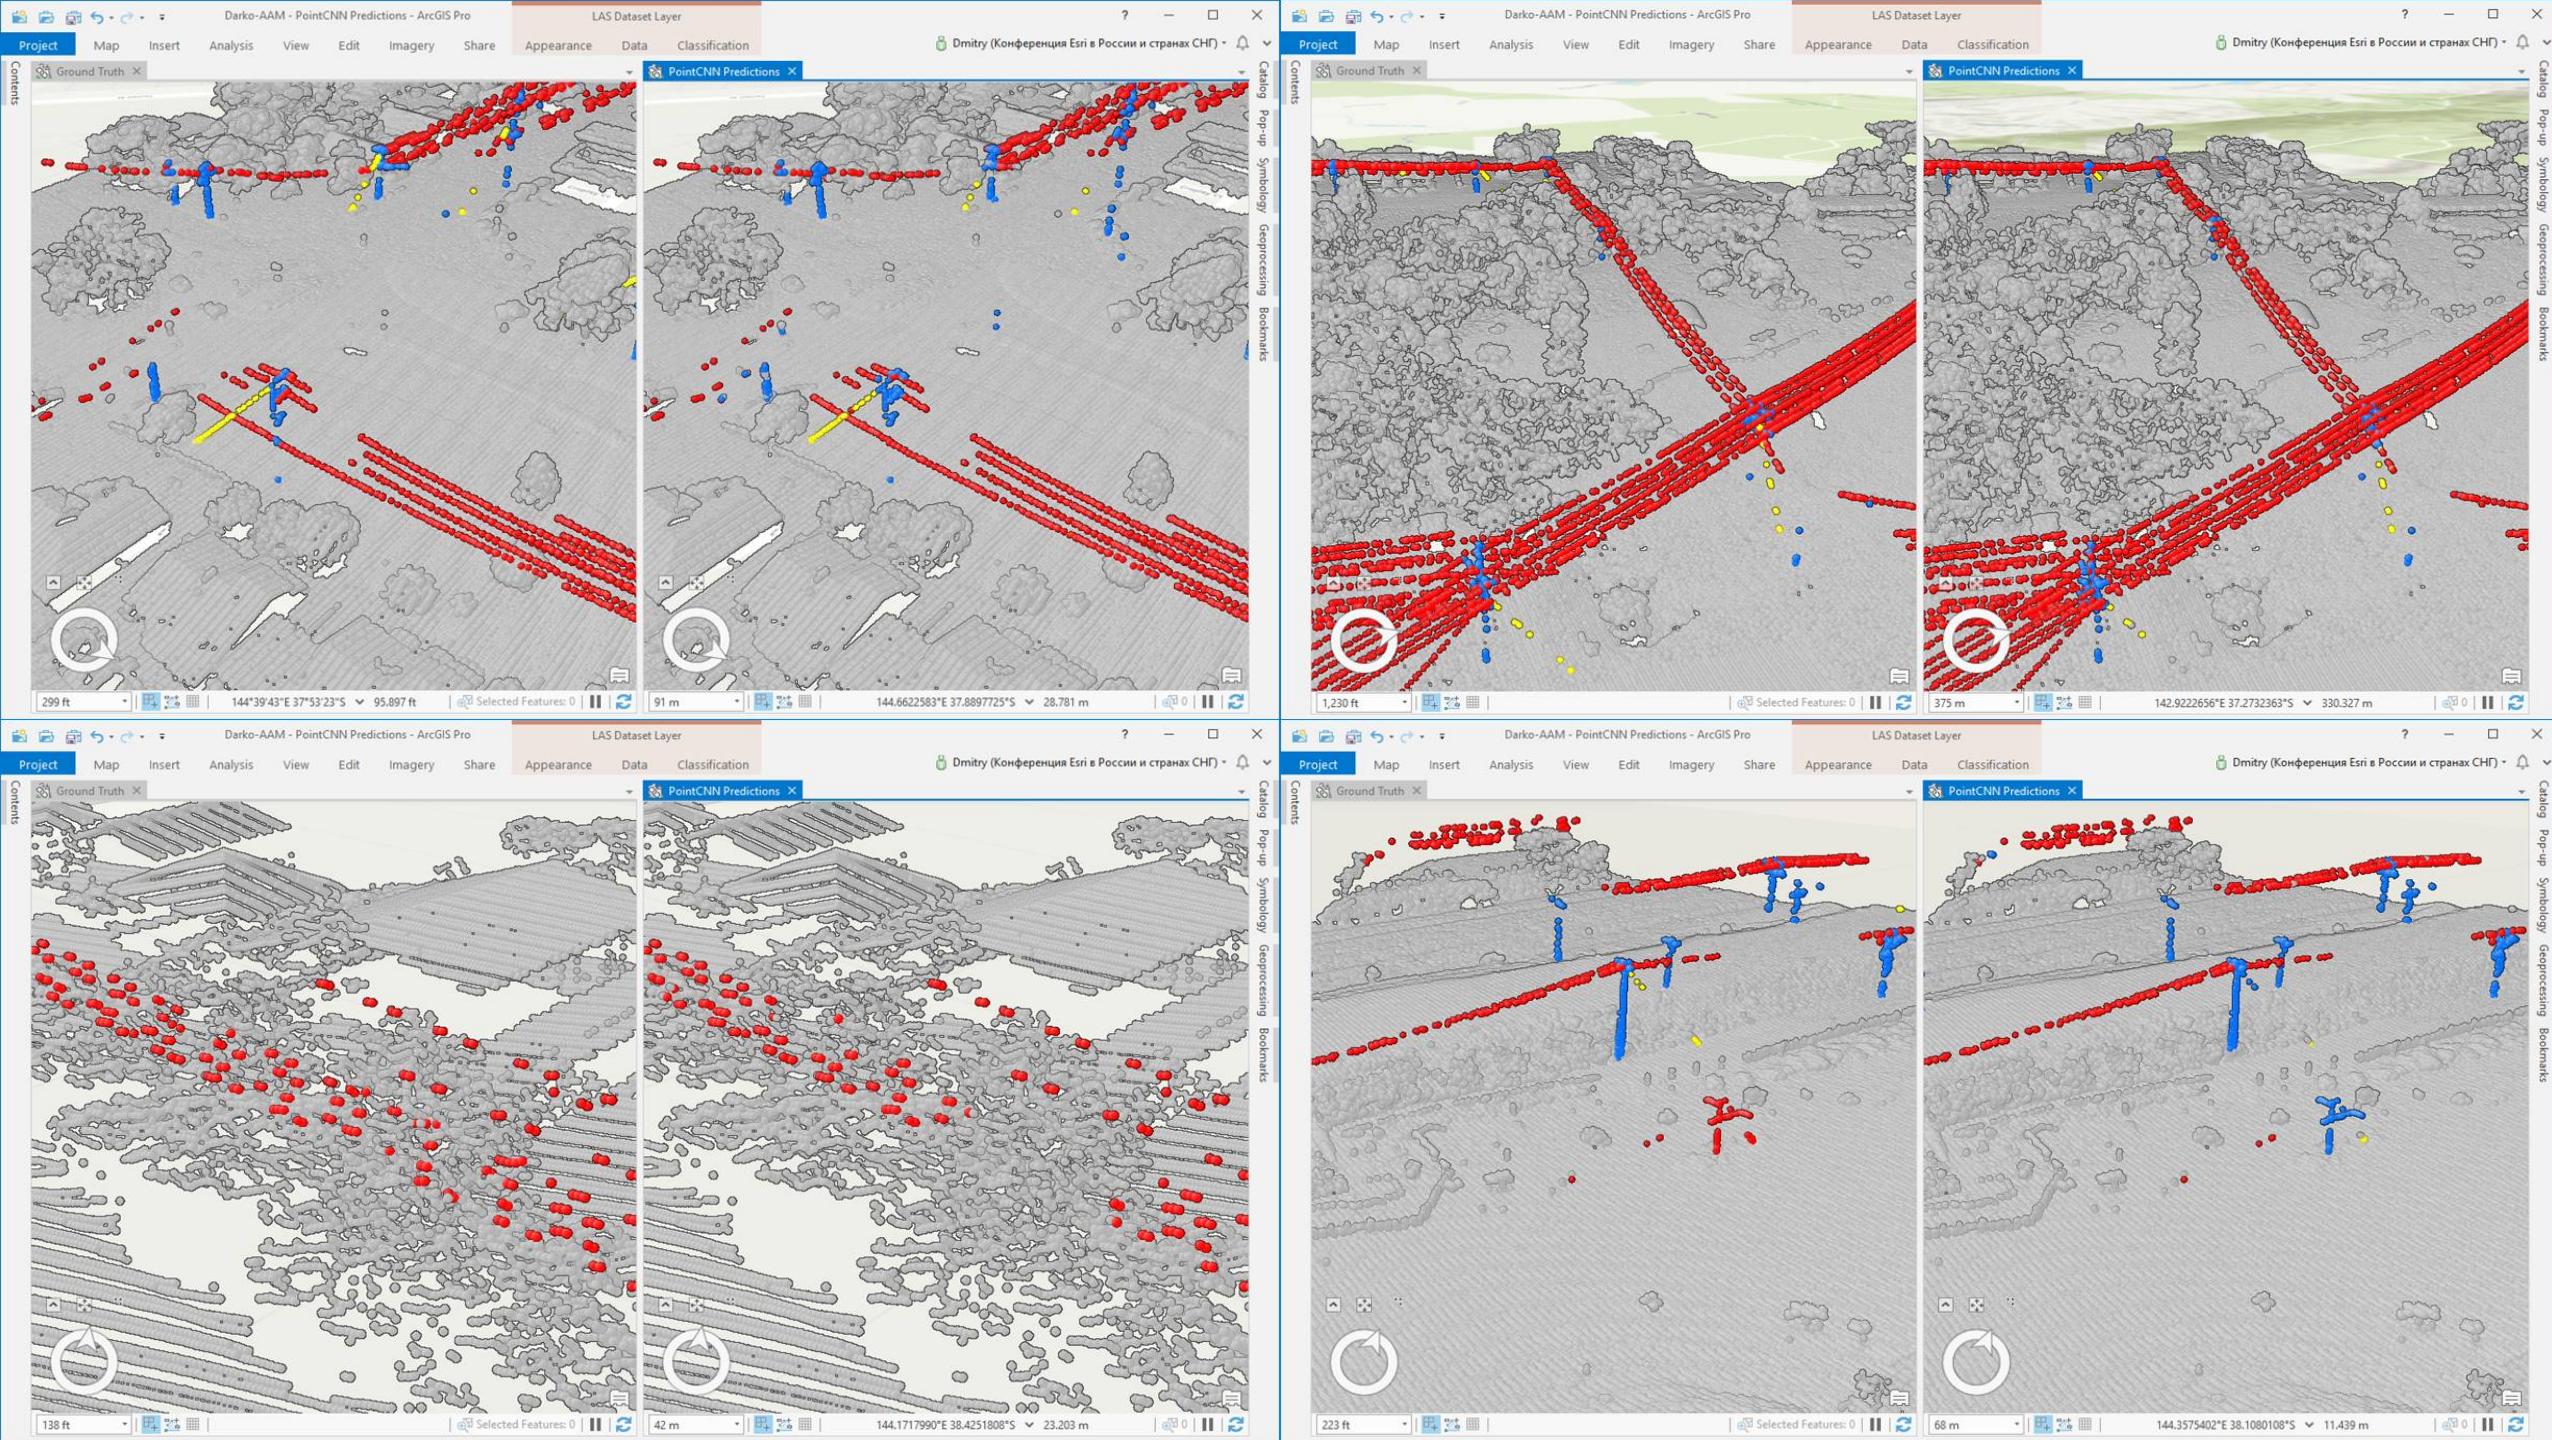
Precision: [0.99988538 0.96672749 0.83674406 0.80313546],

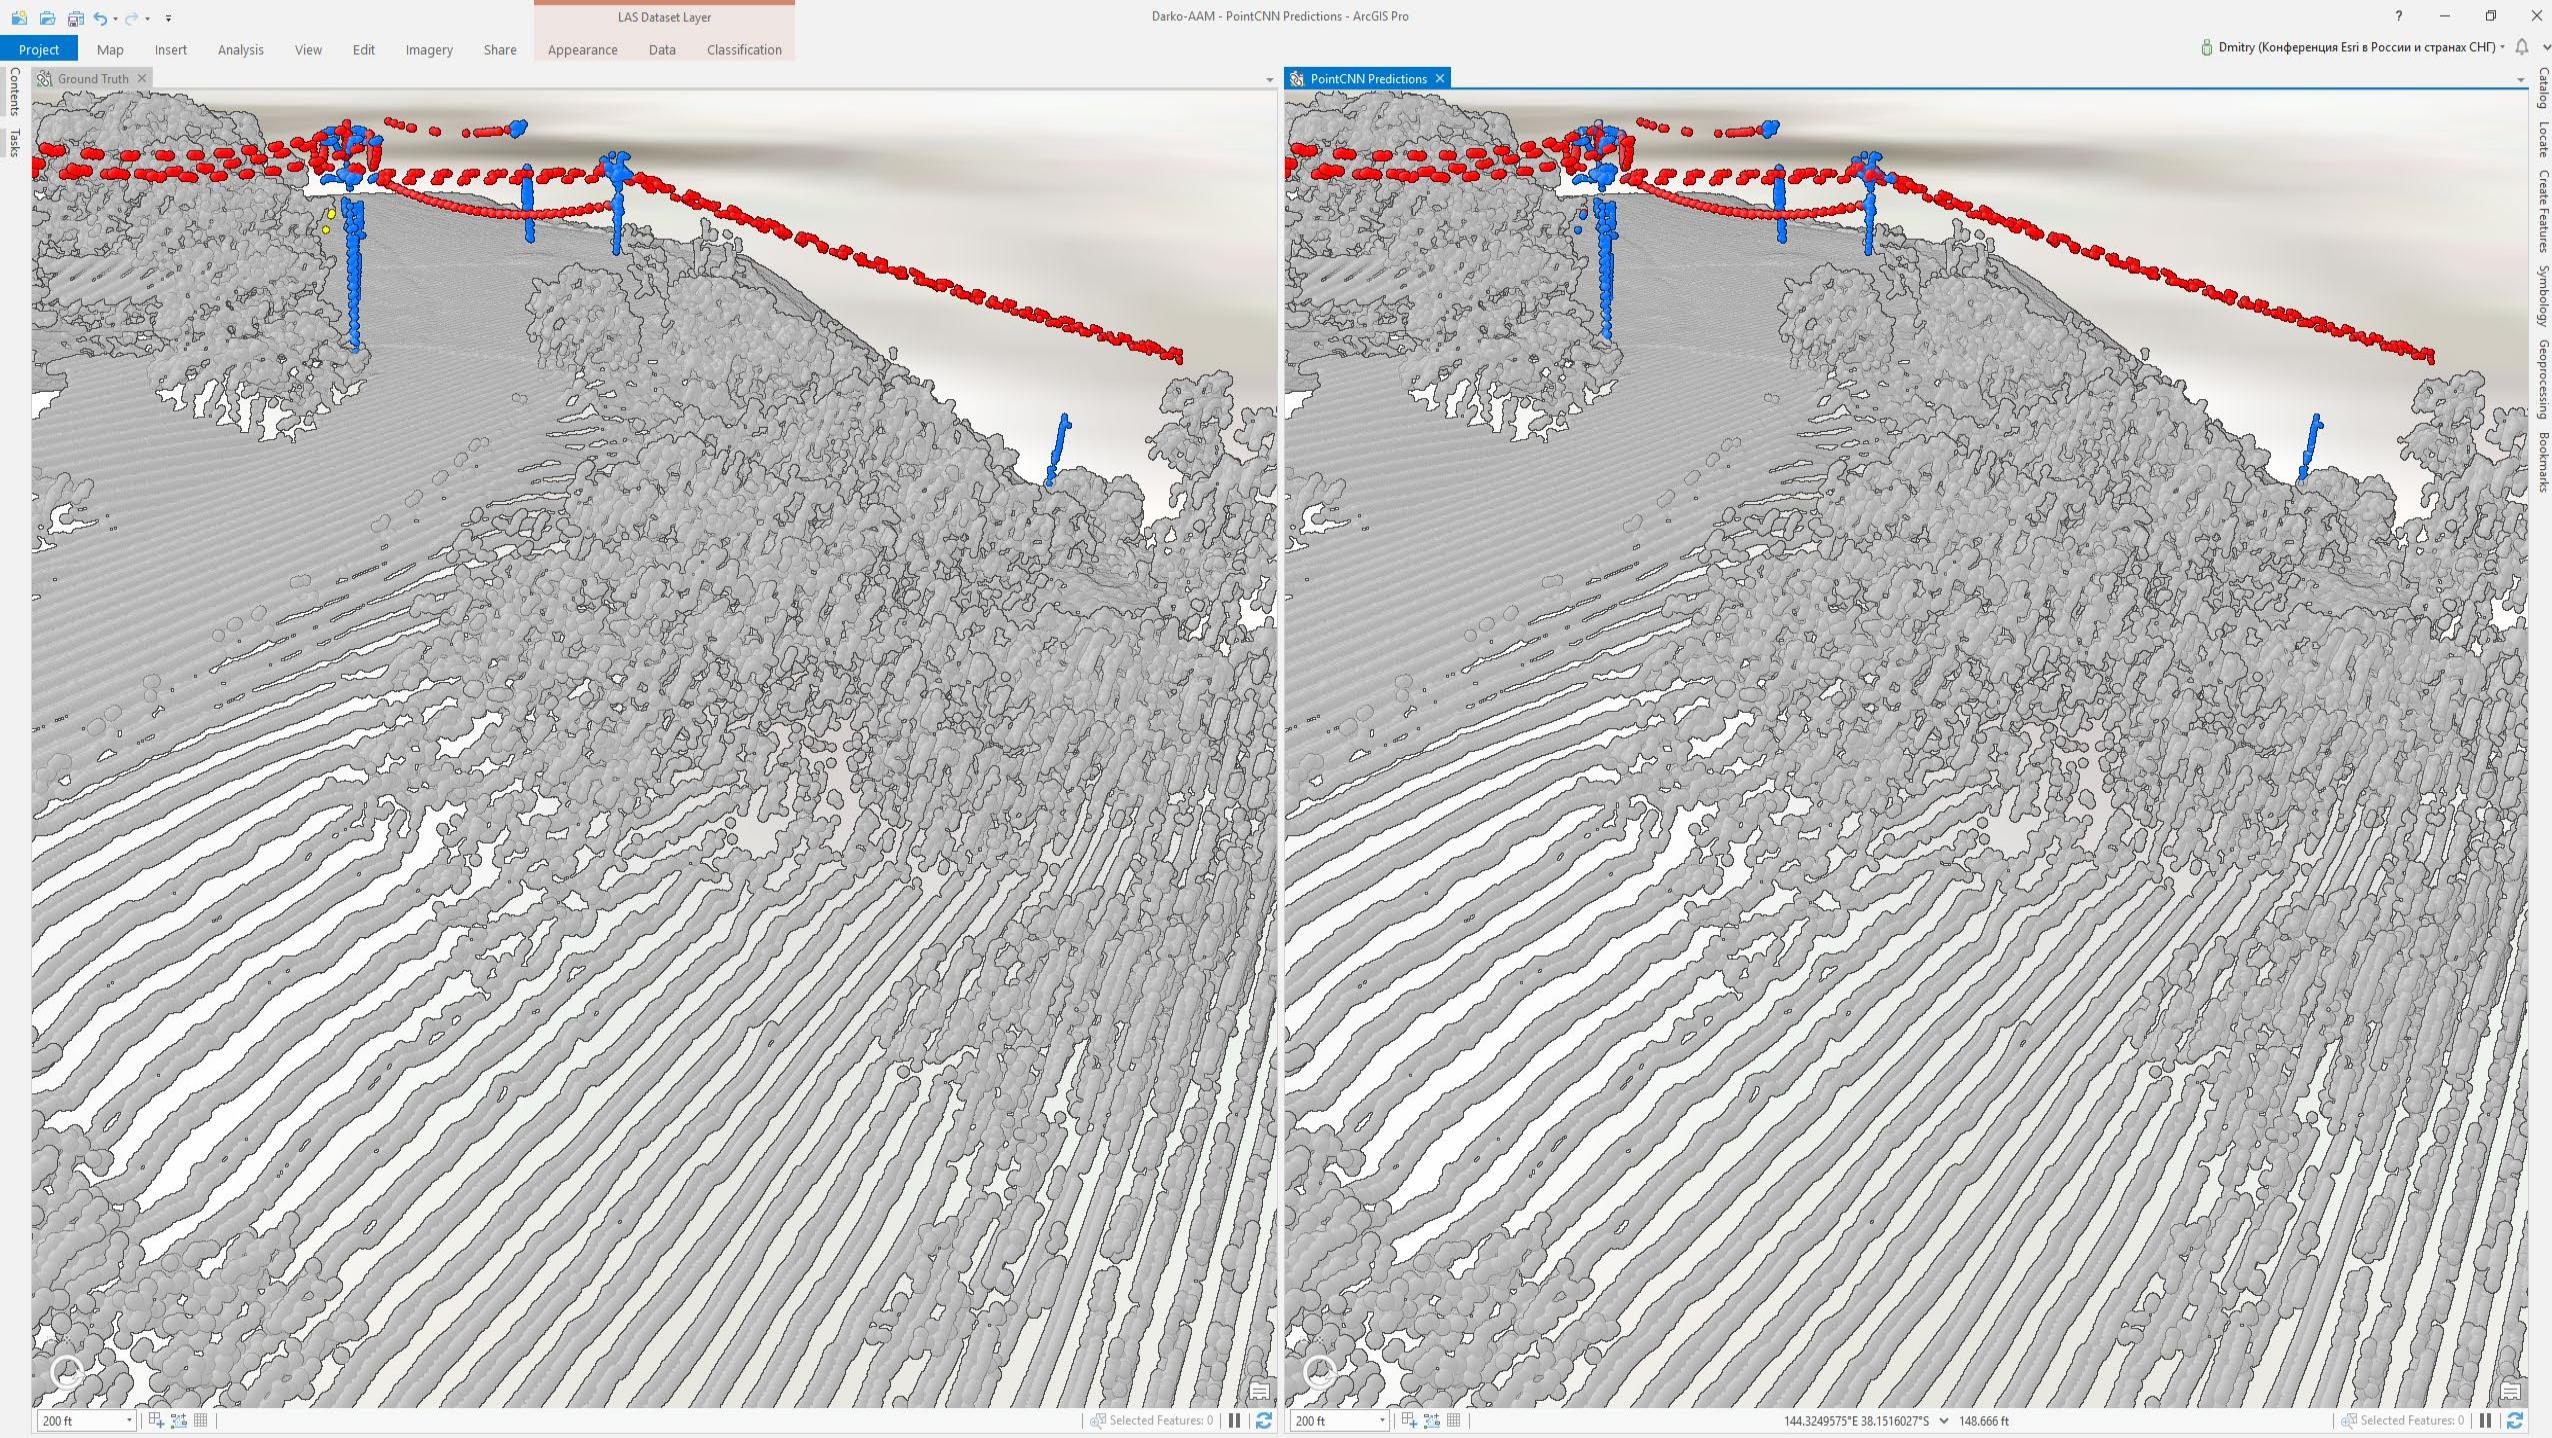
Recall: [0.99987221 0.98060249 0.21455632 0.77497643],

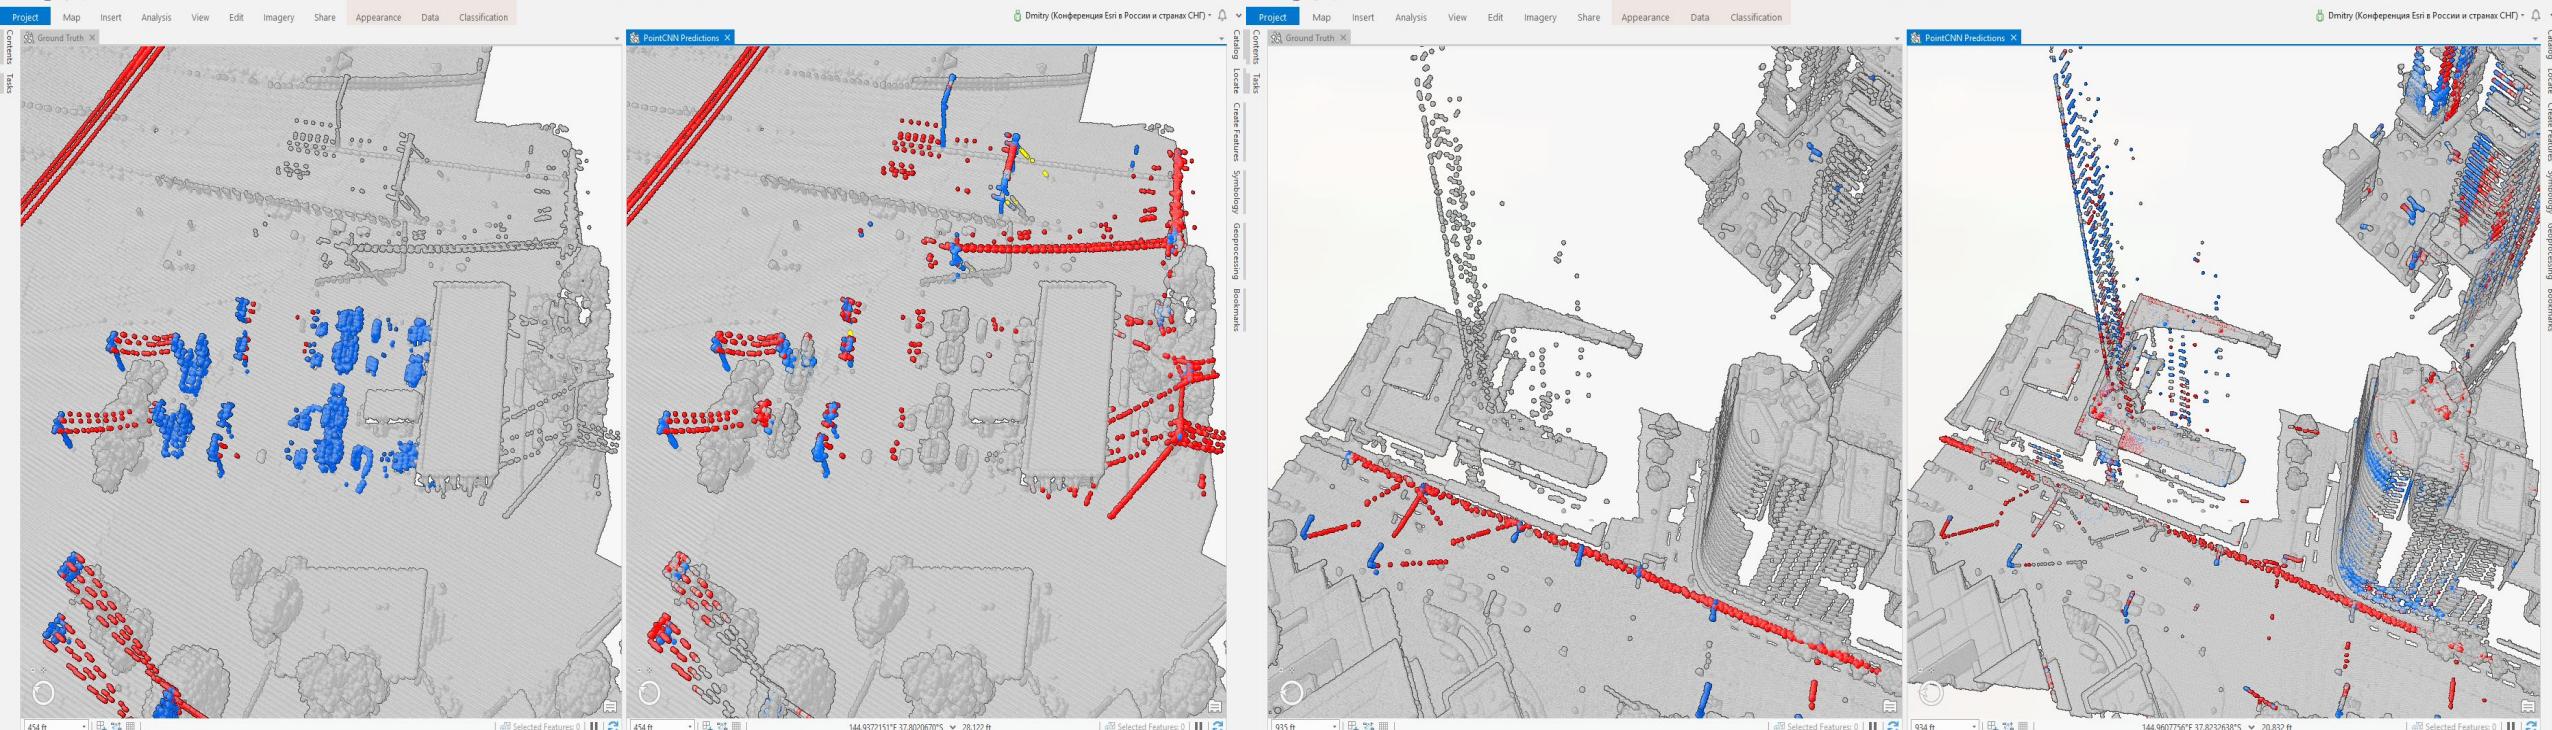
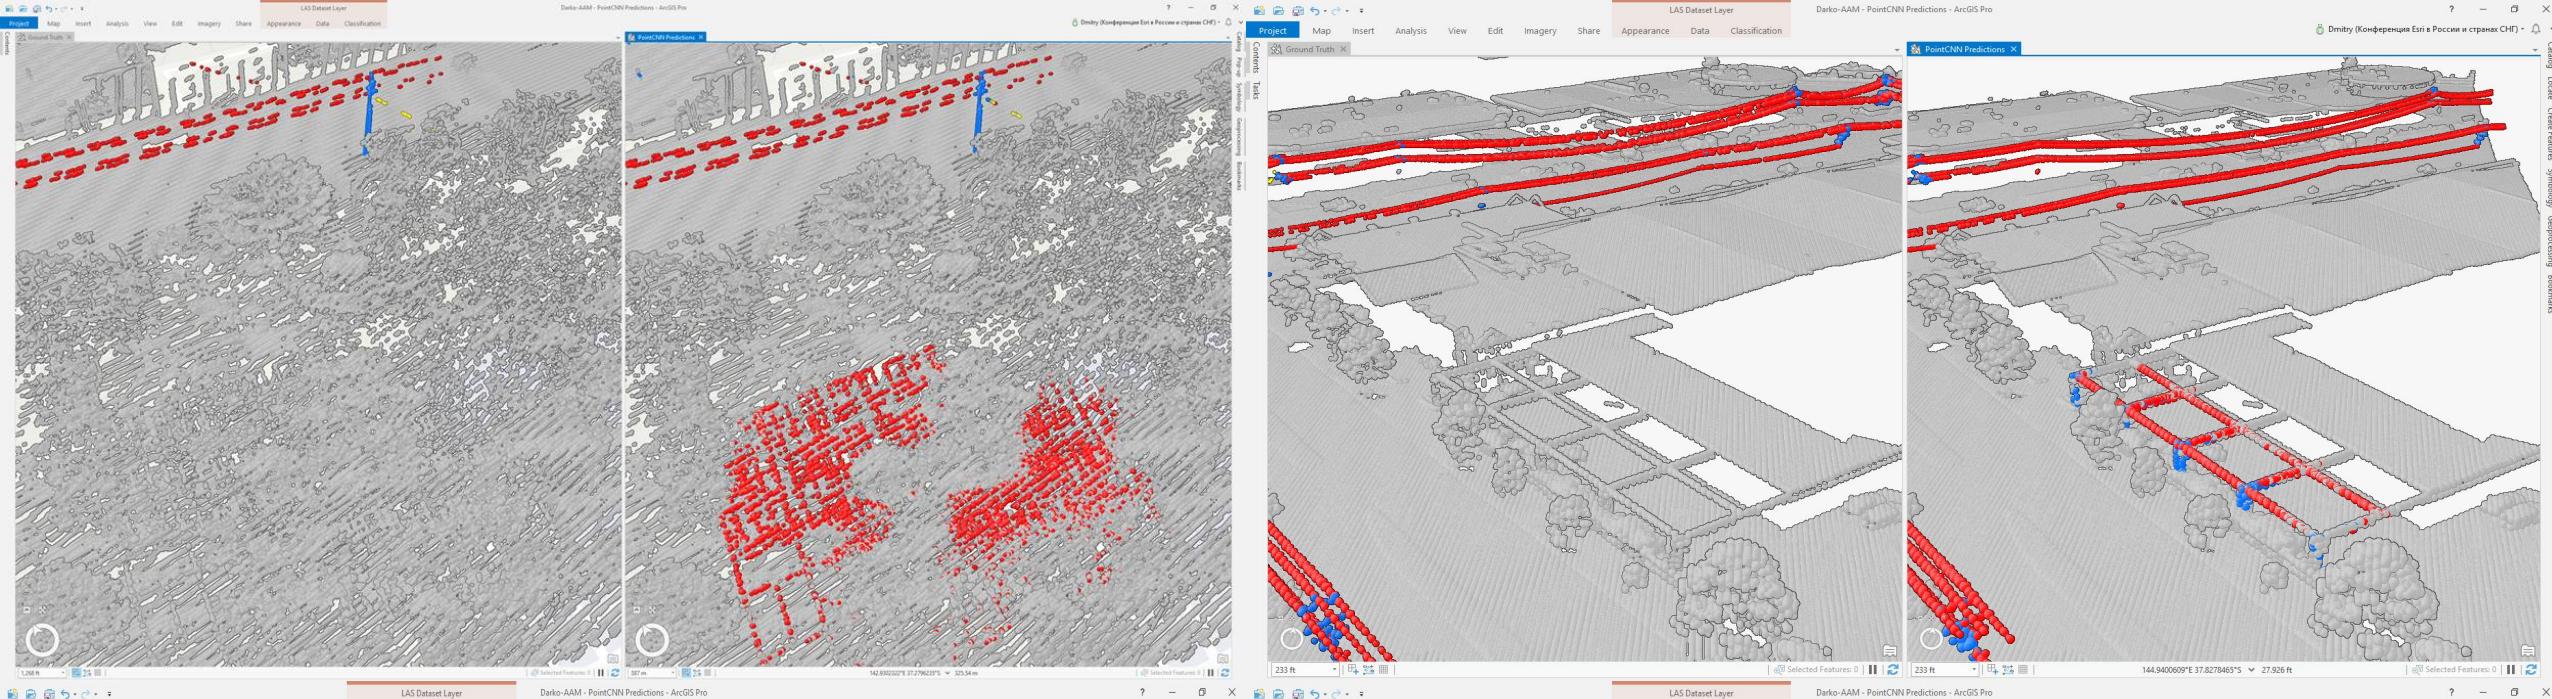
F1 score: [0.9998788 0.97361556 0.3415365 0.78880471]

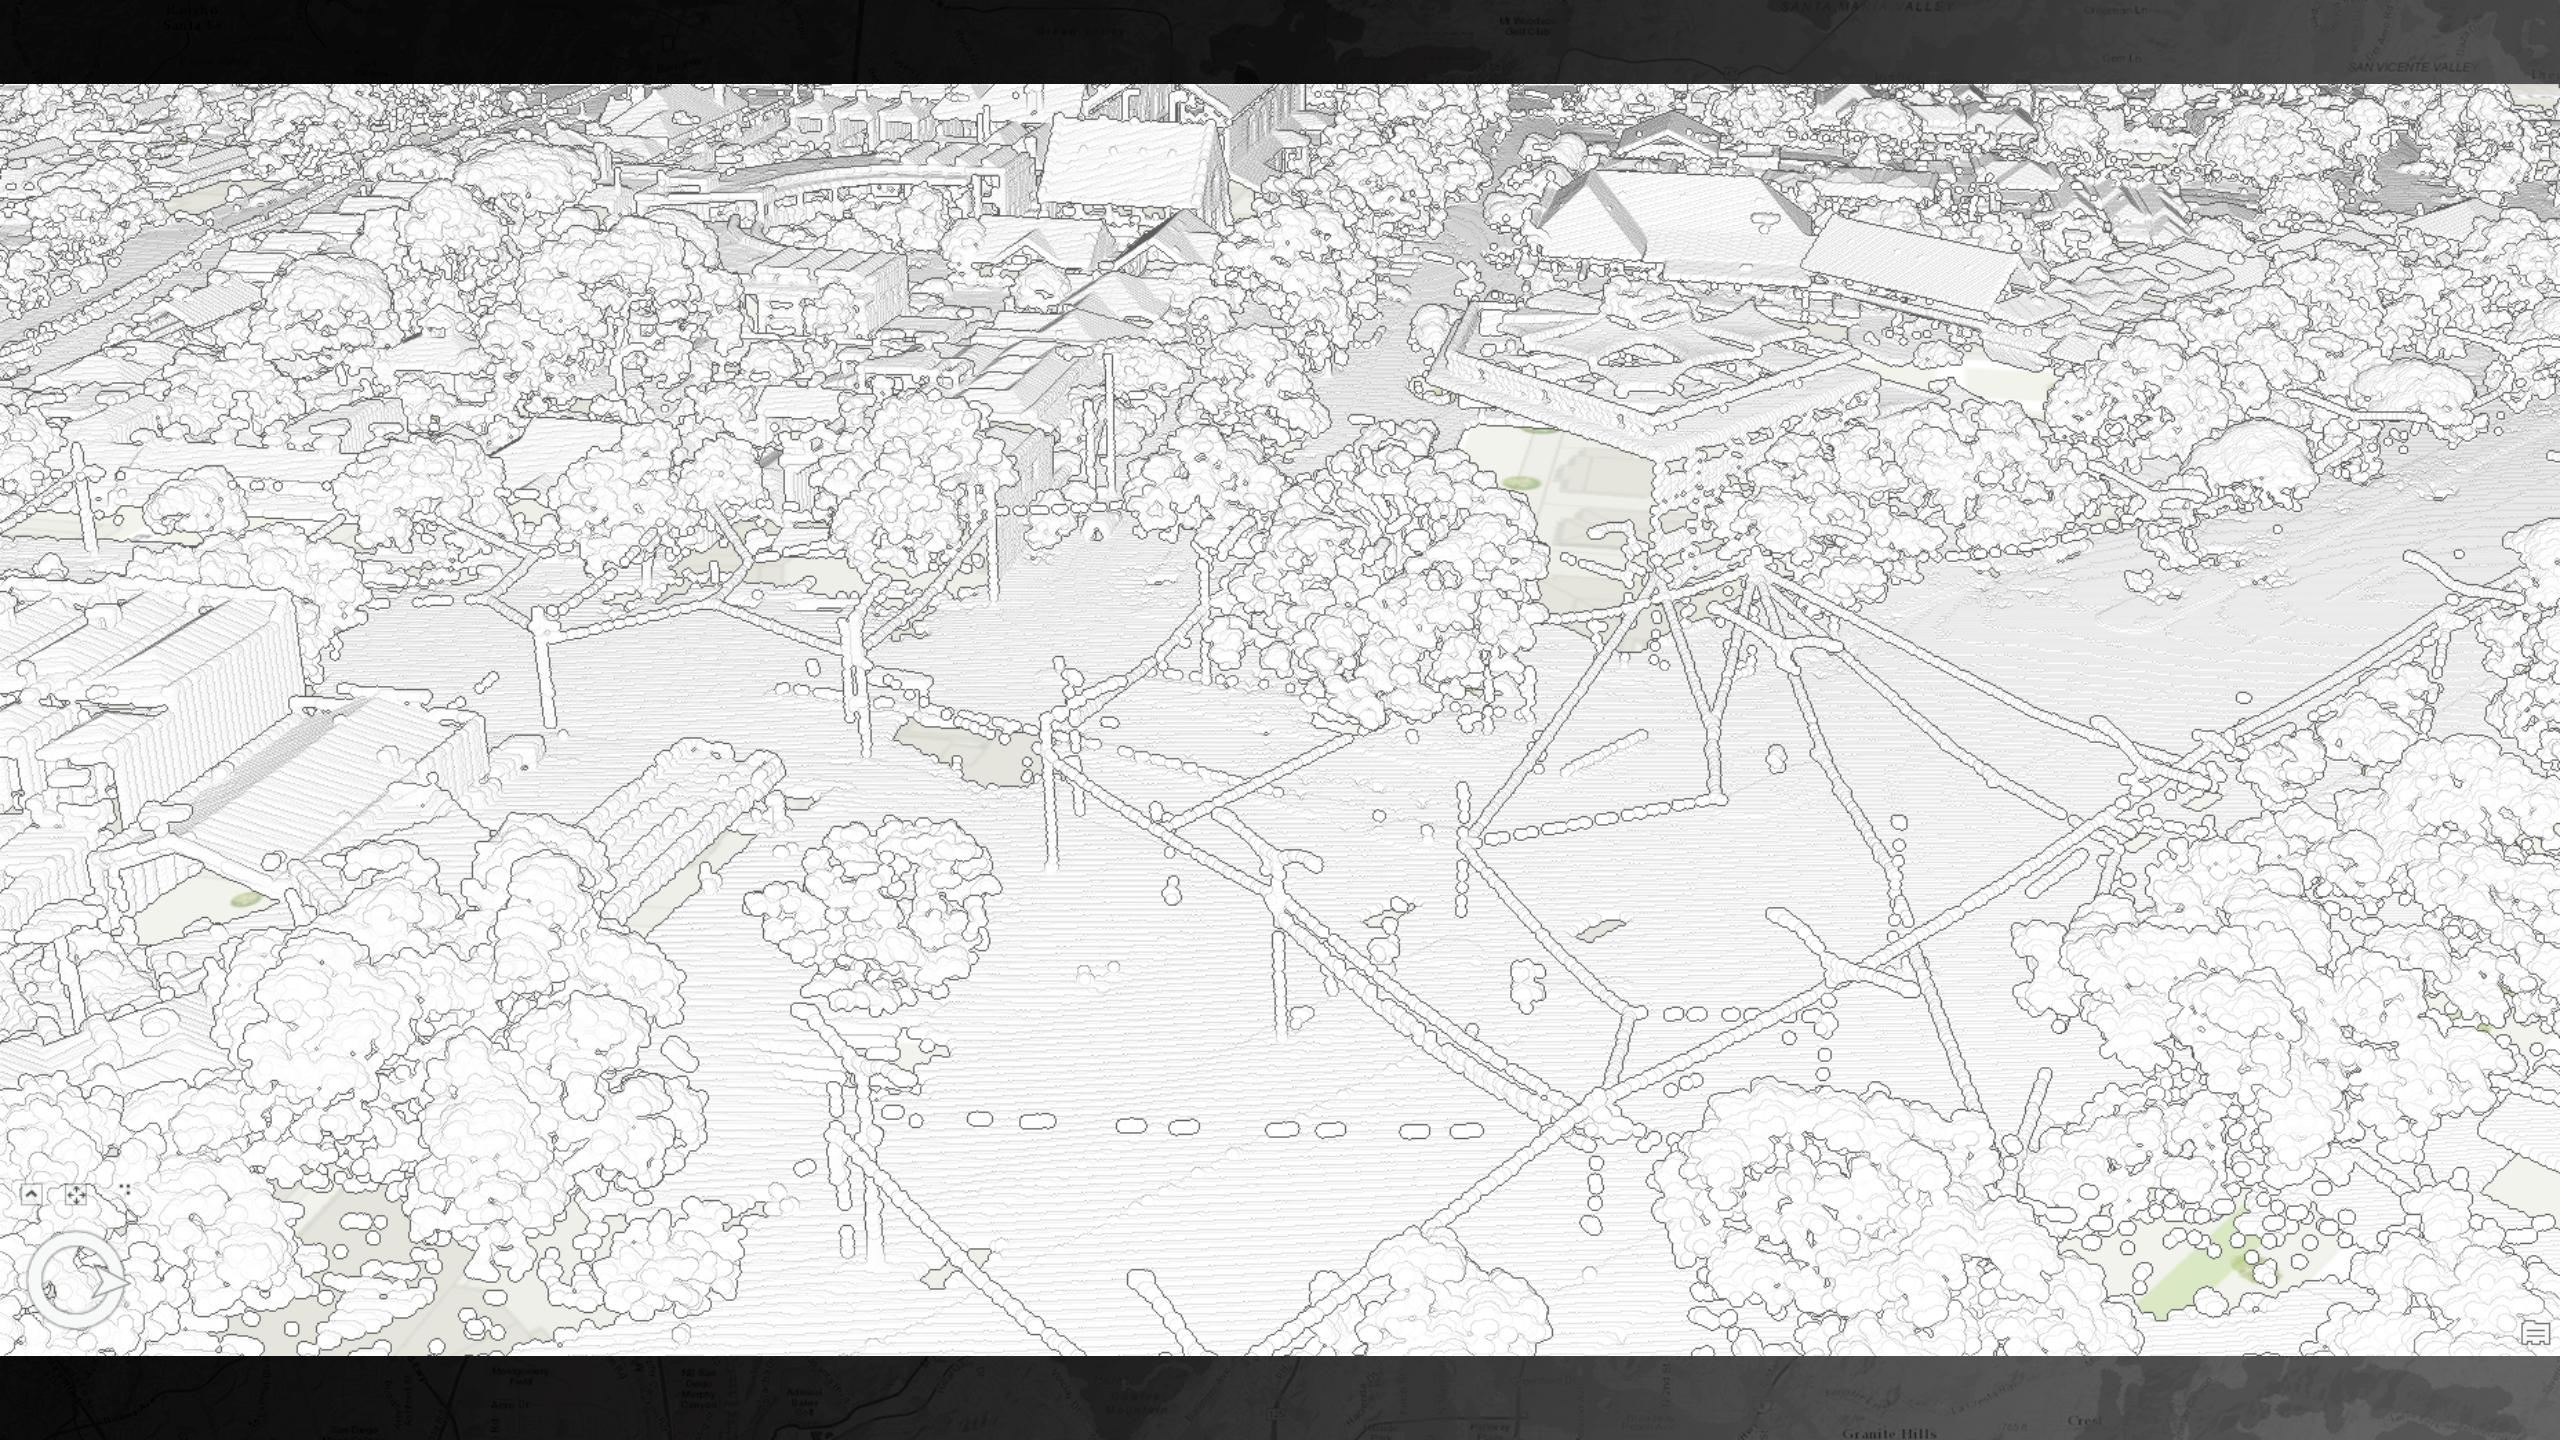


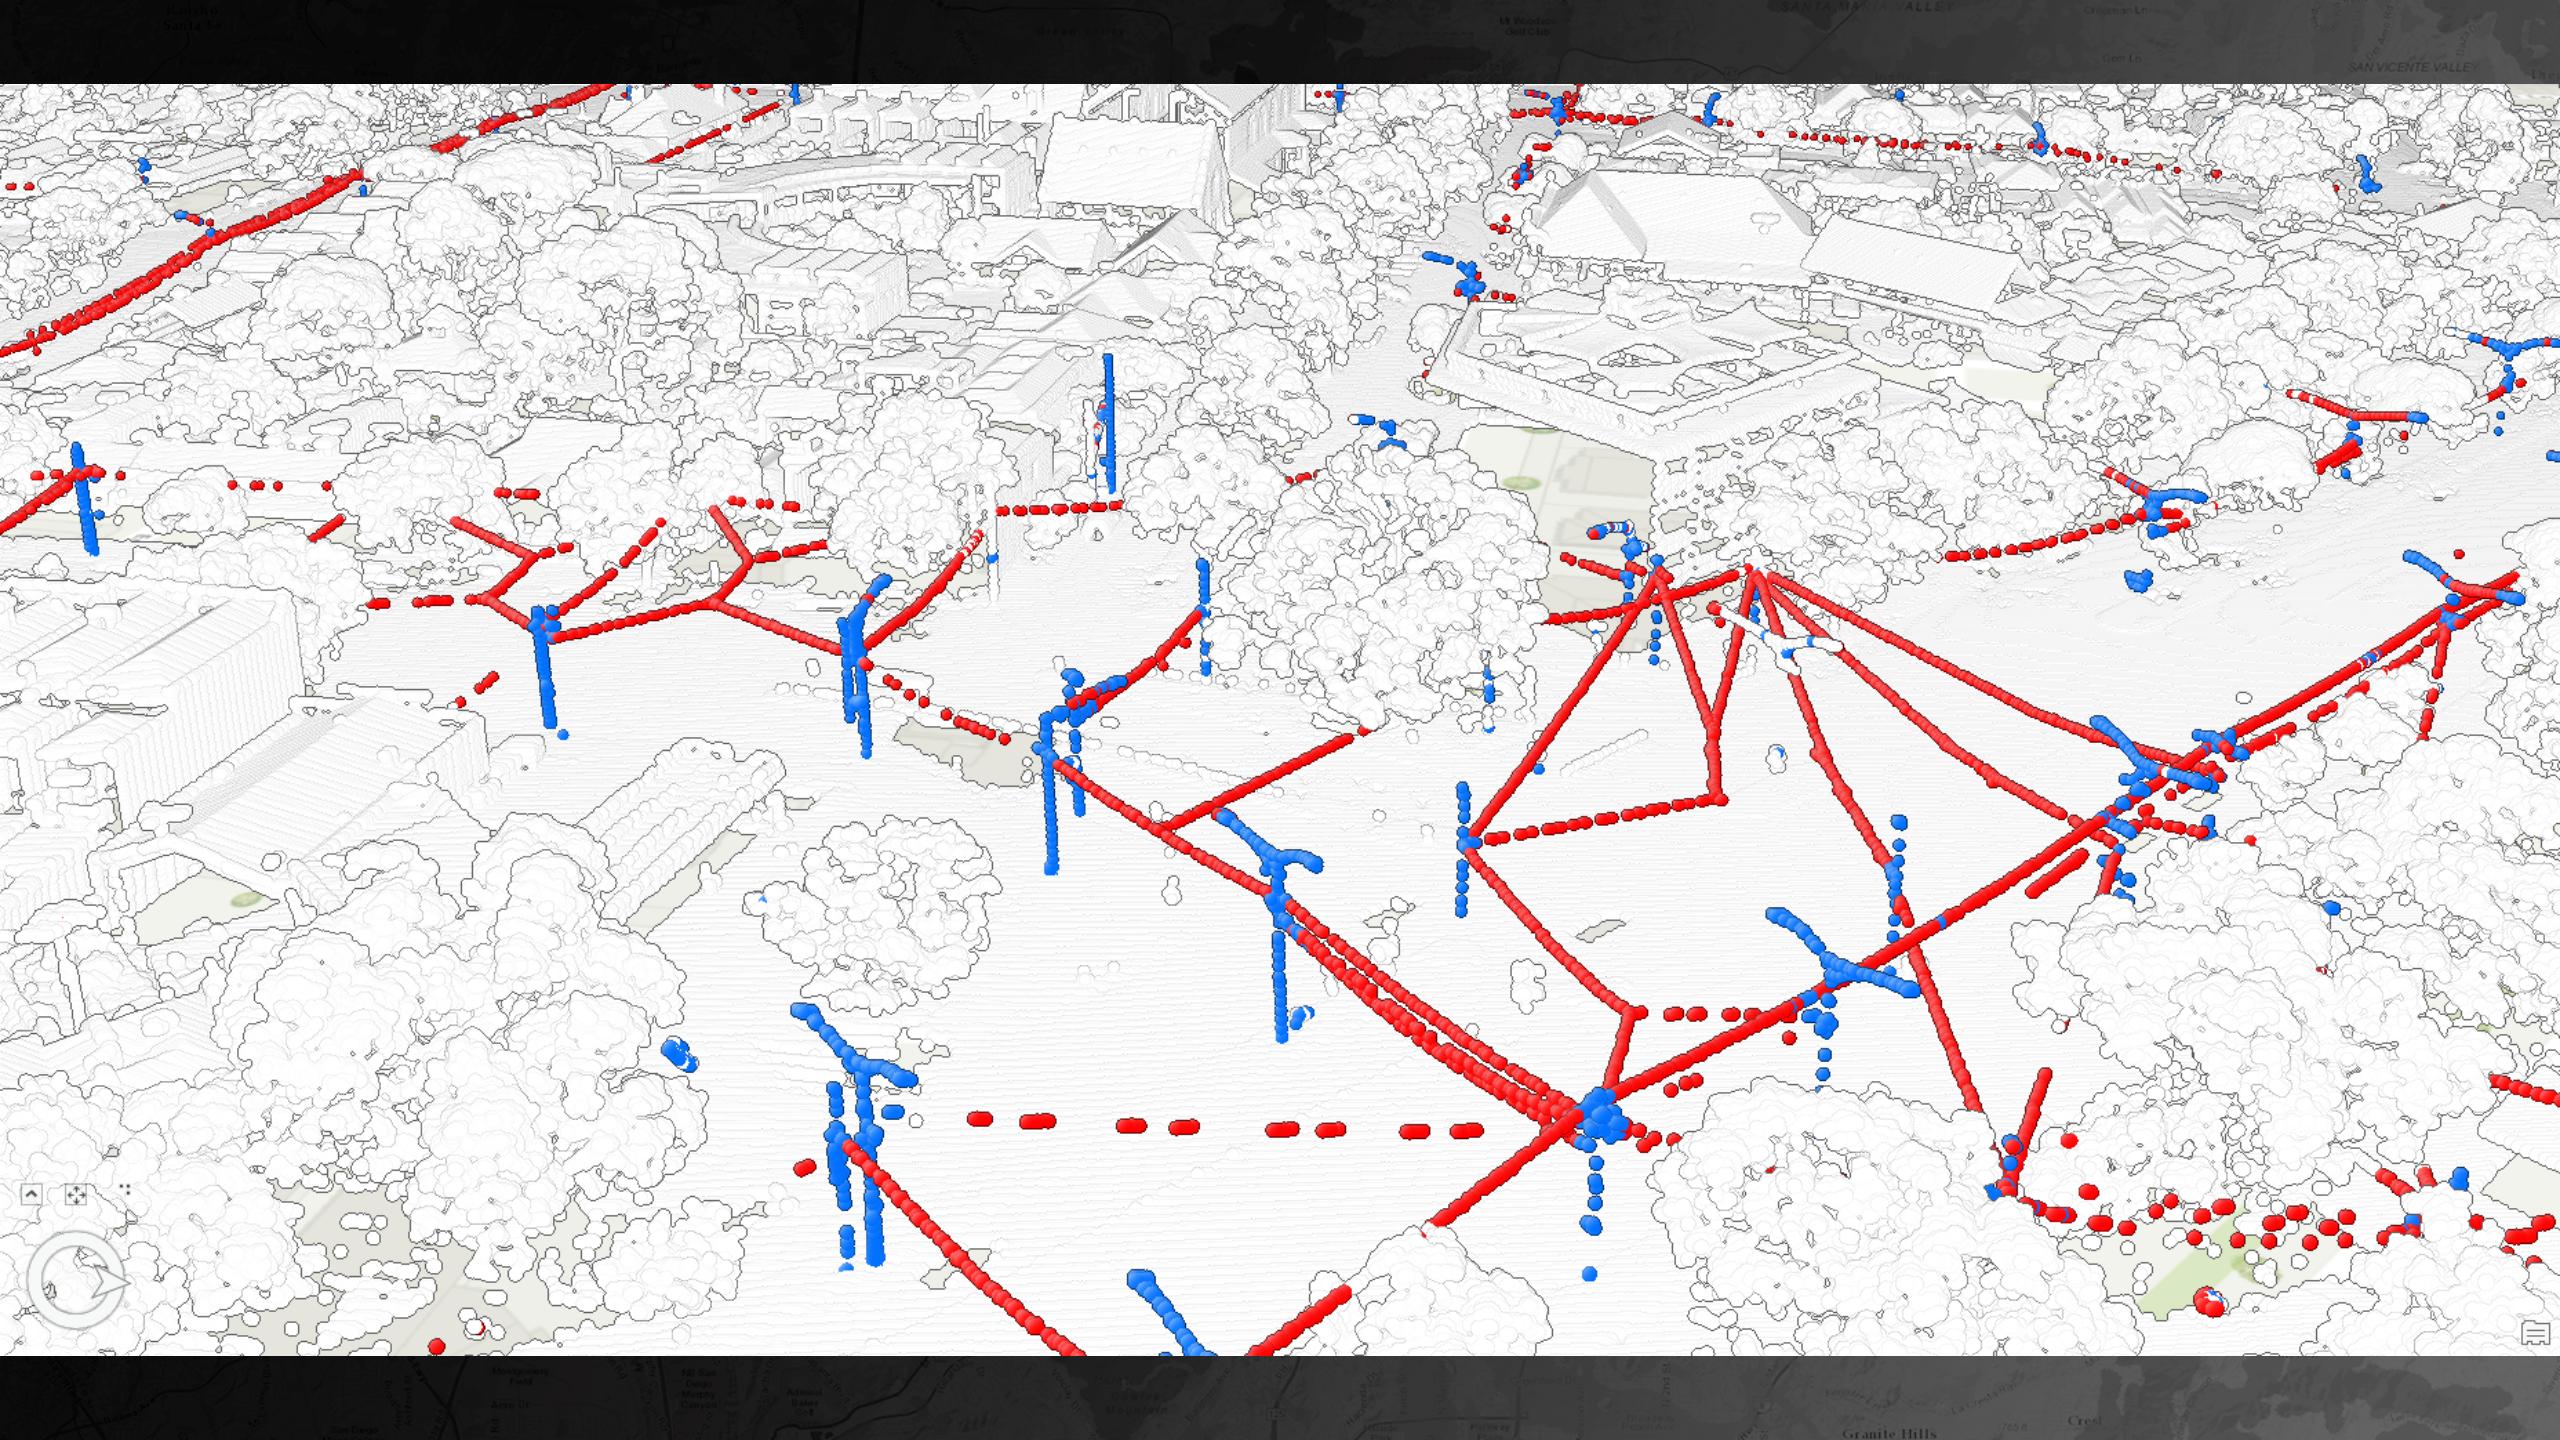


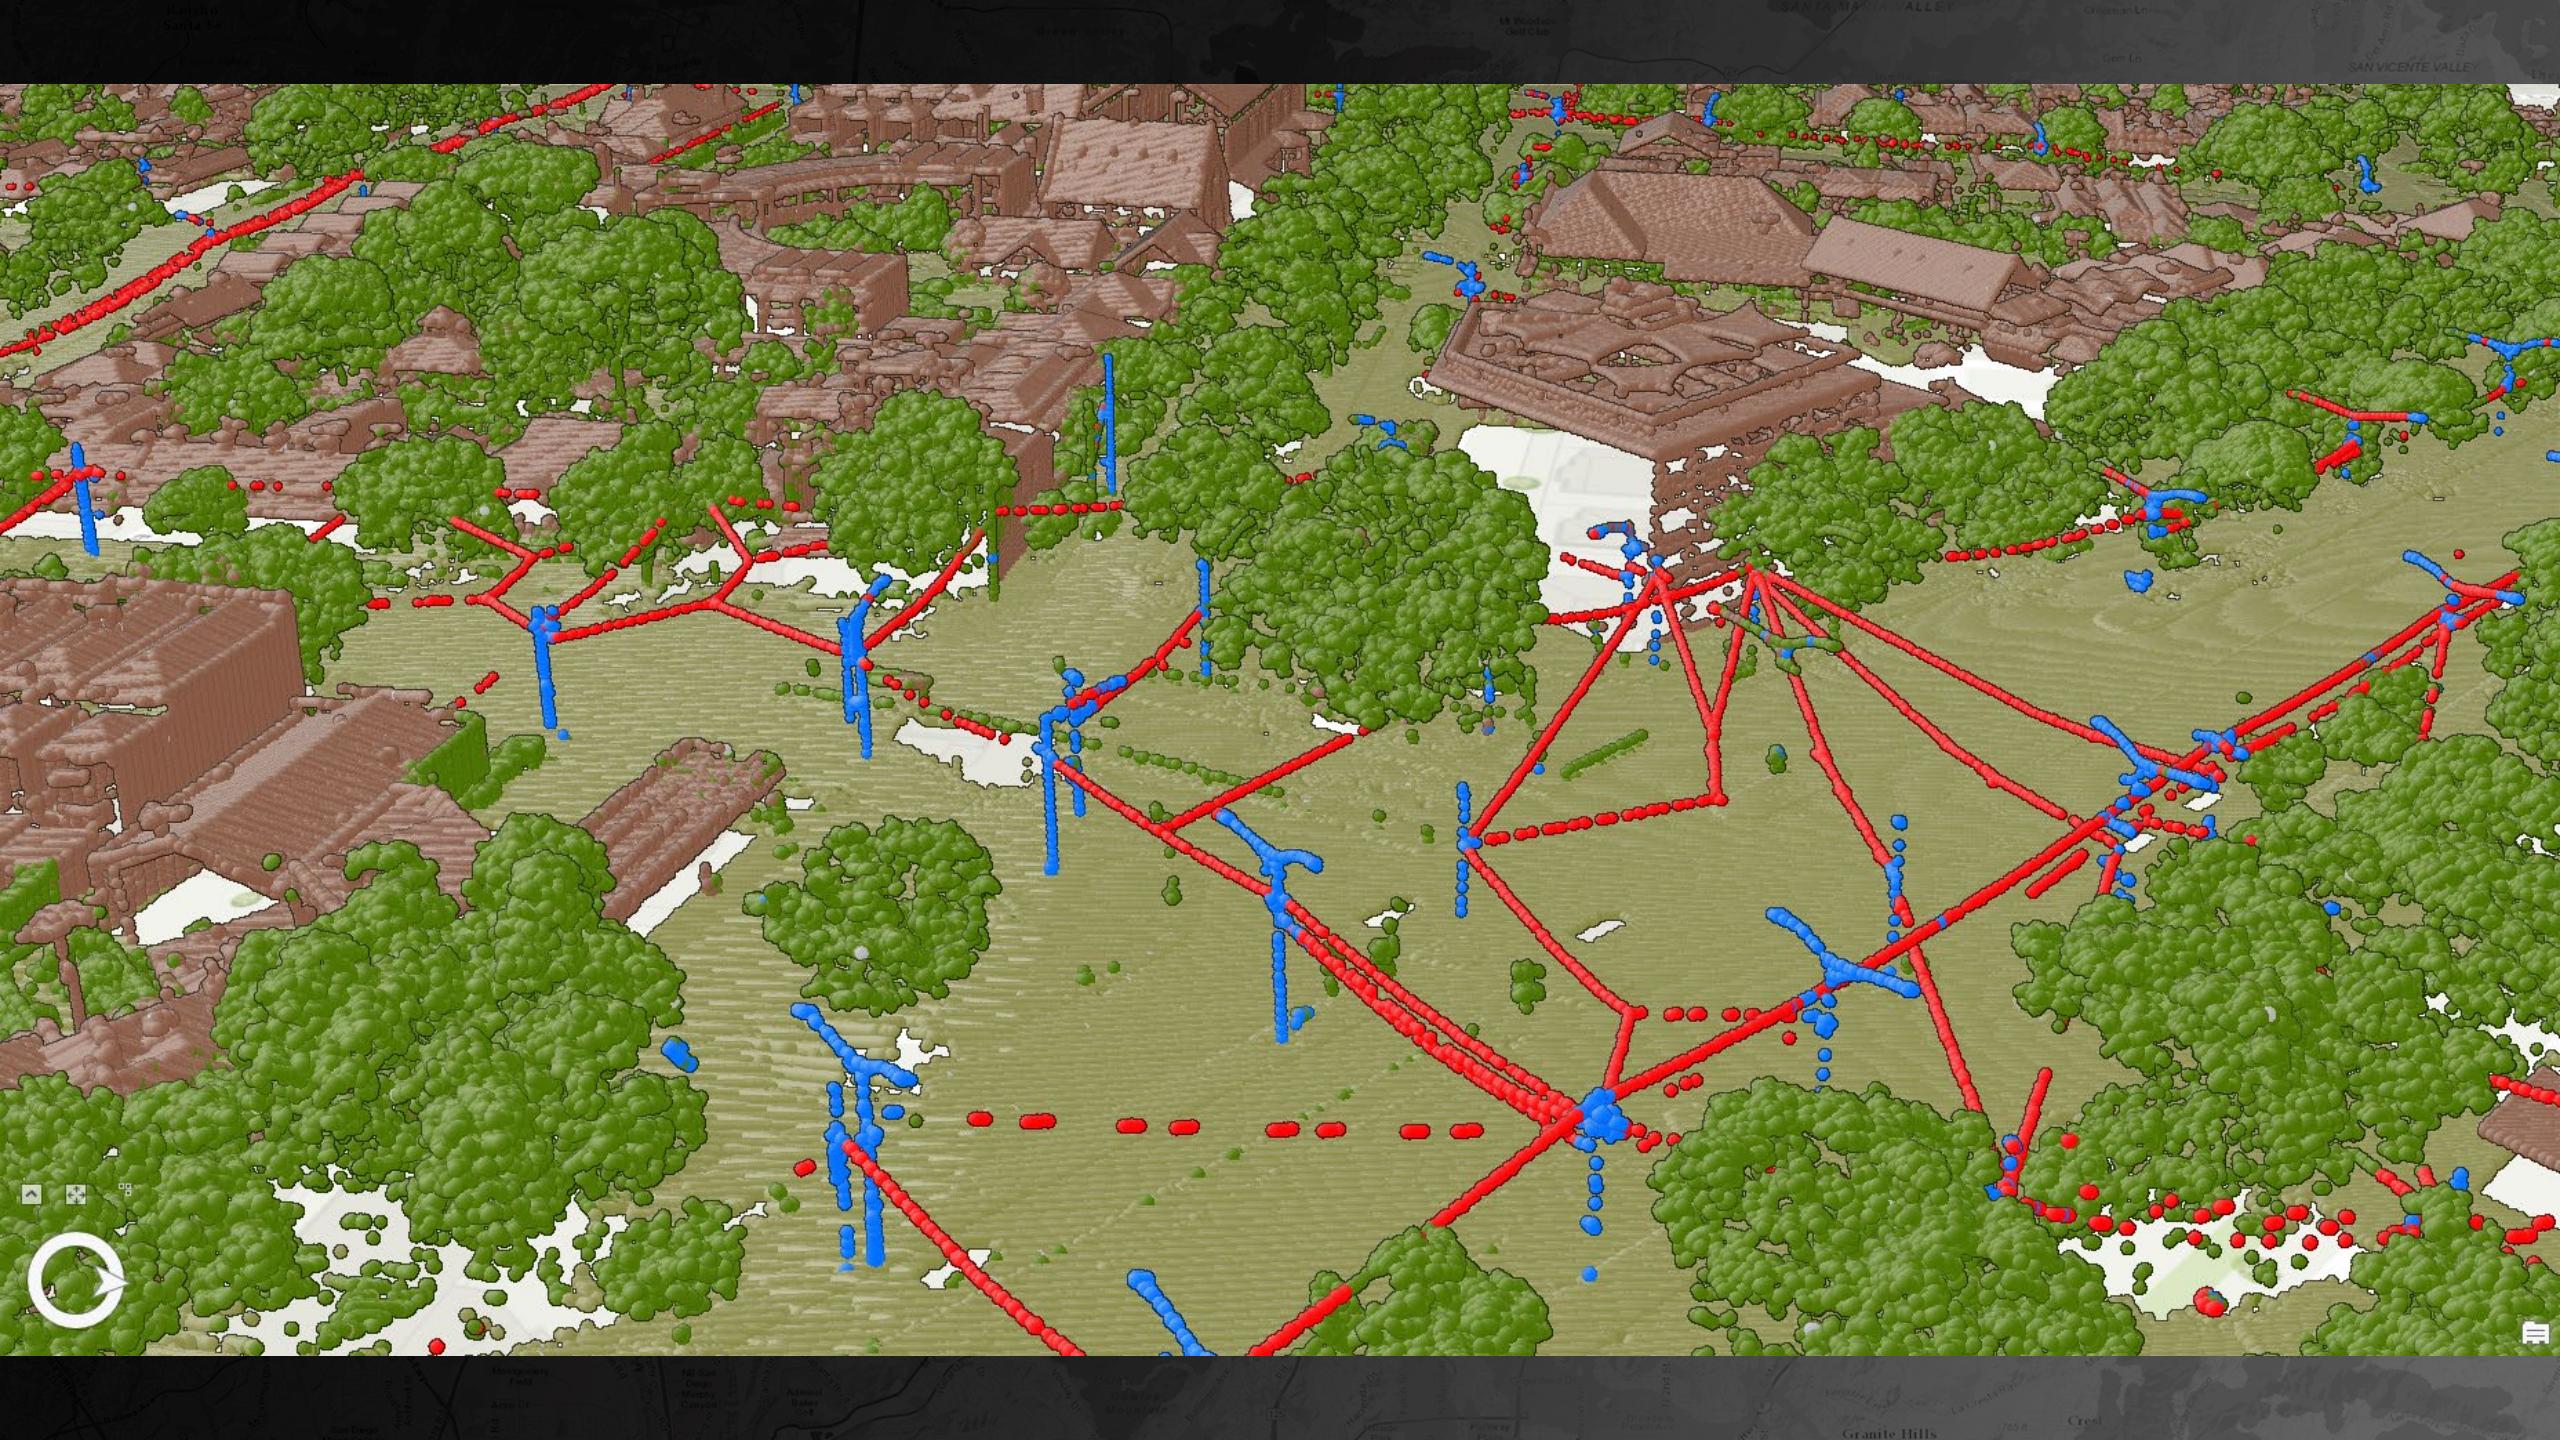


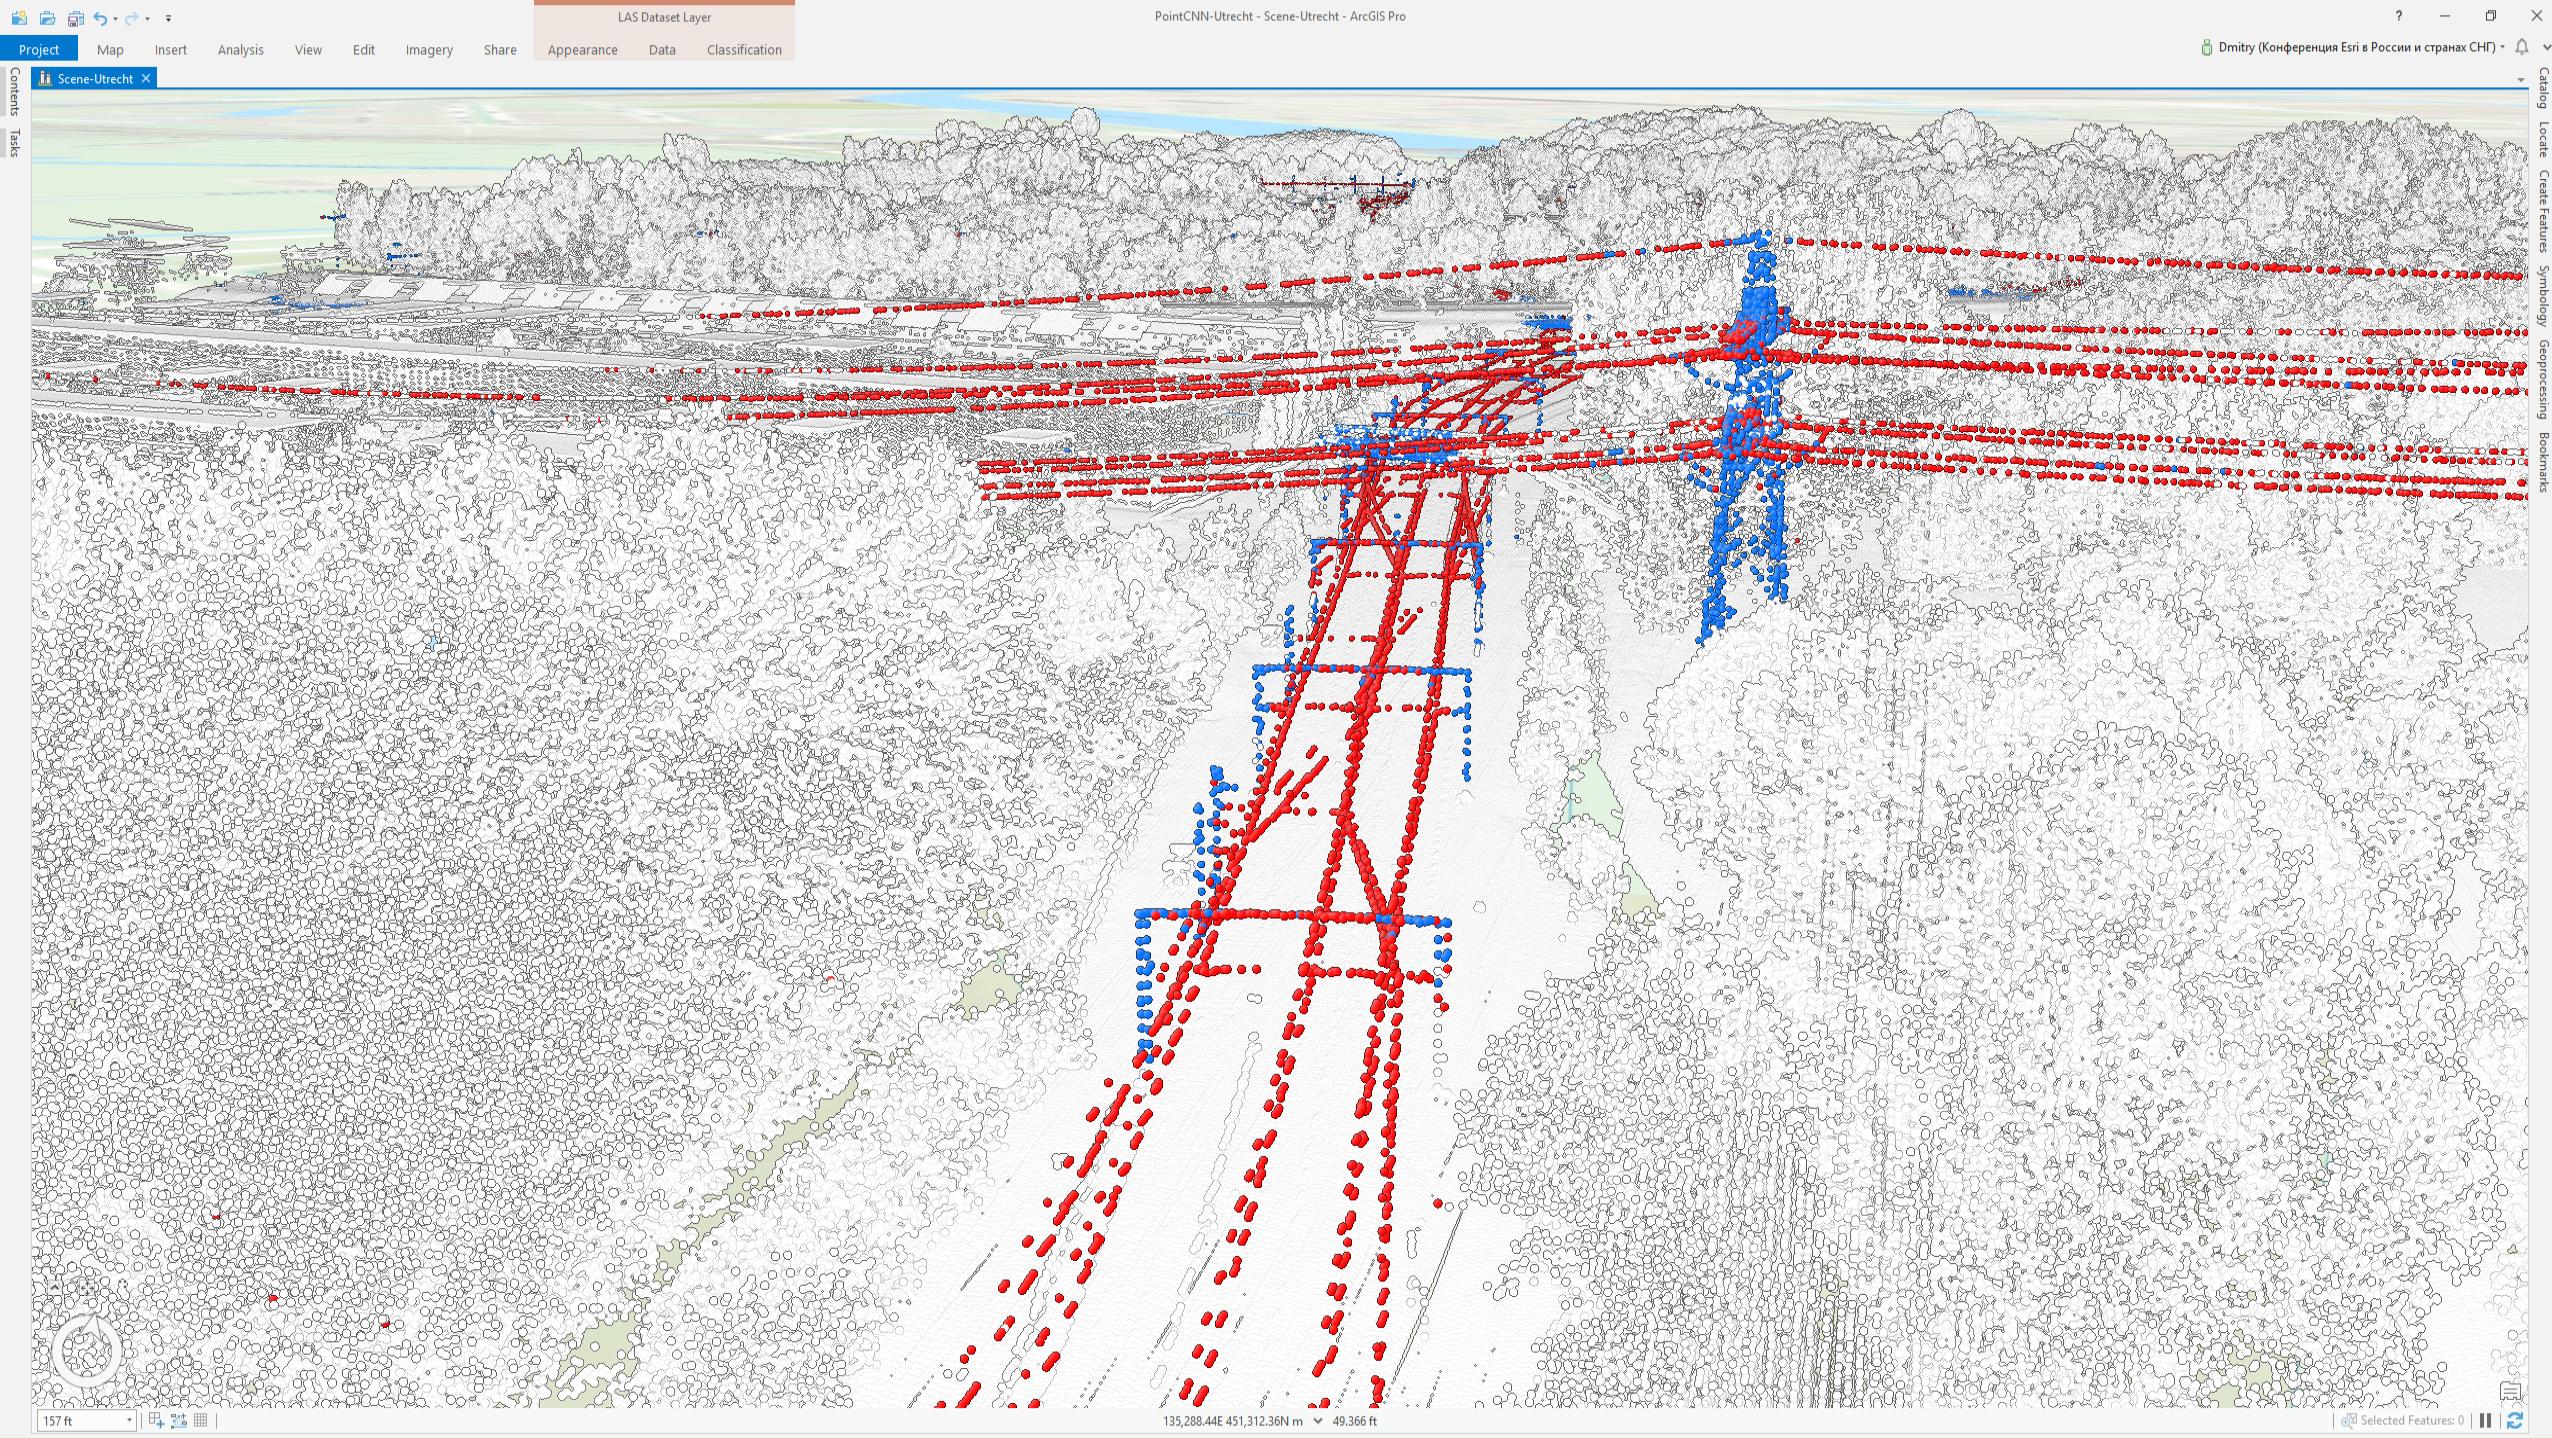


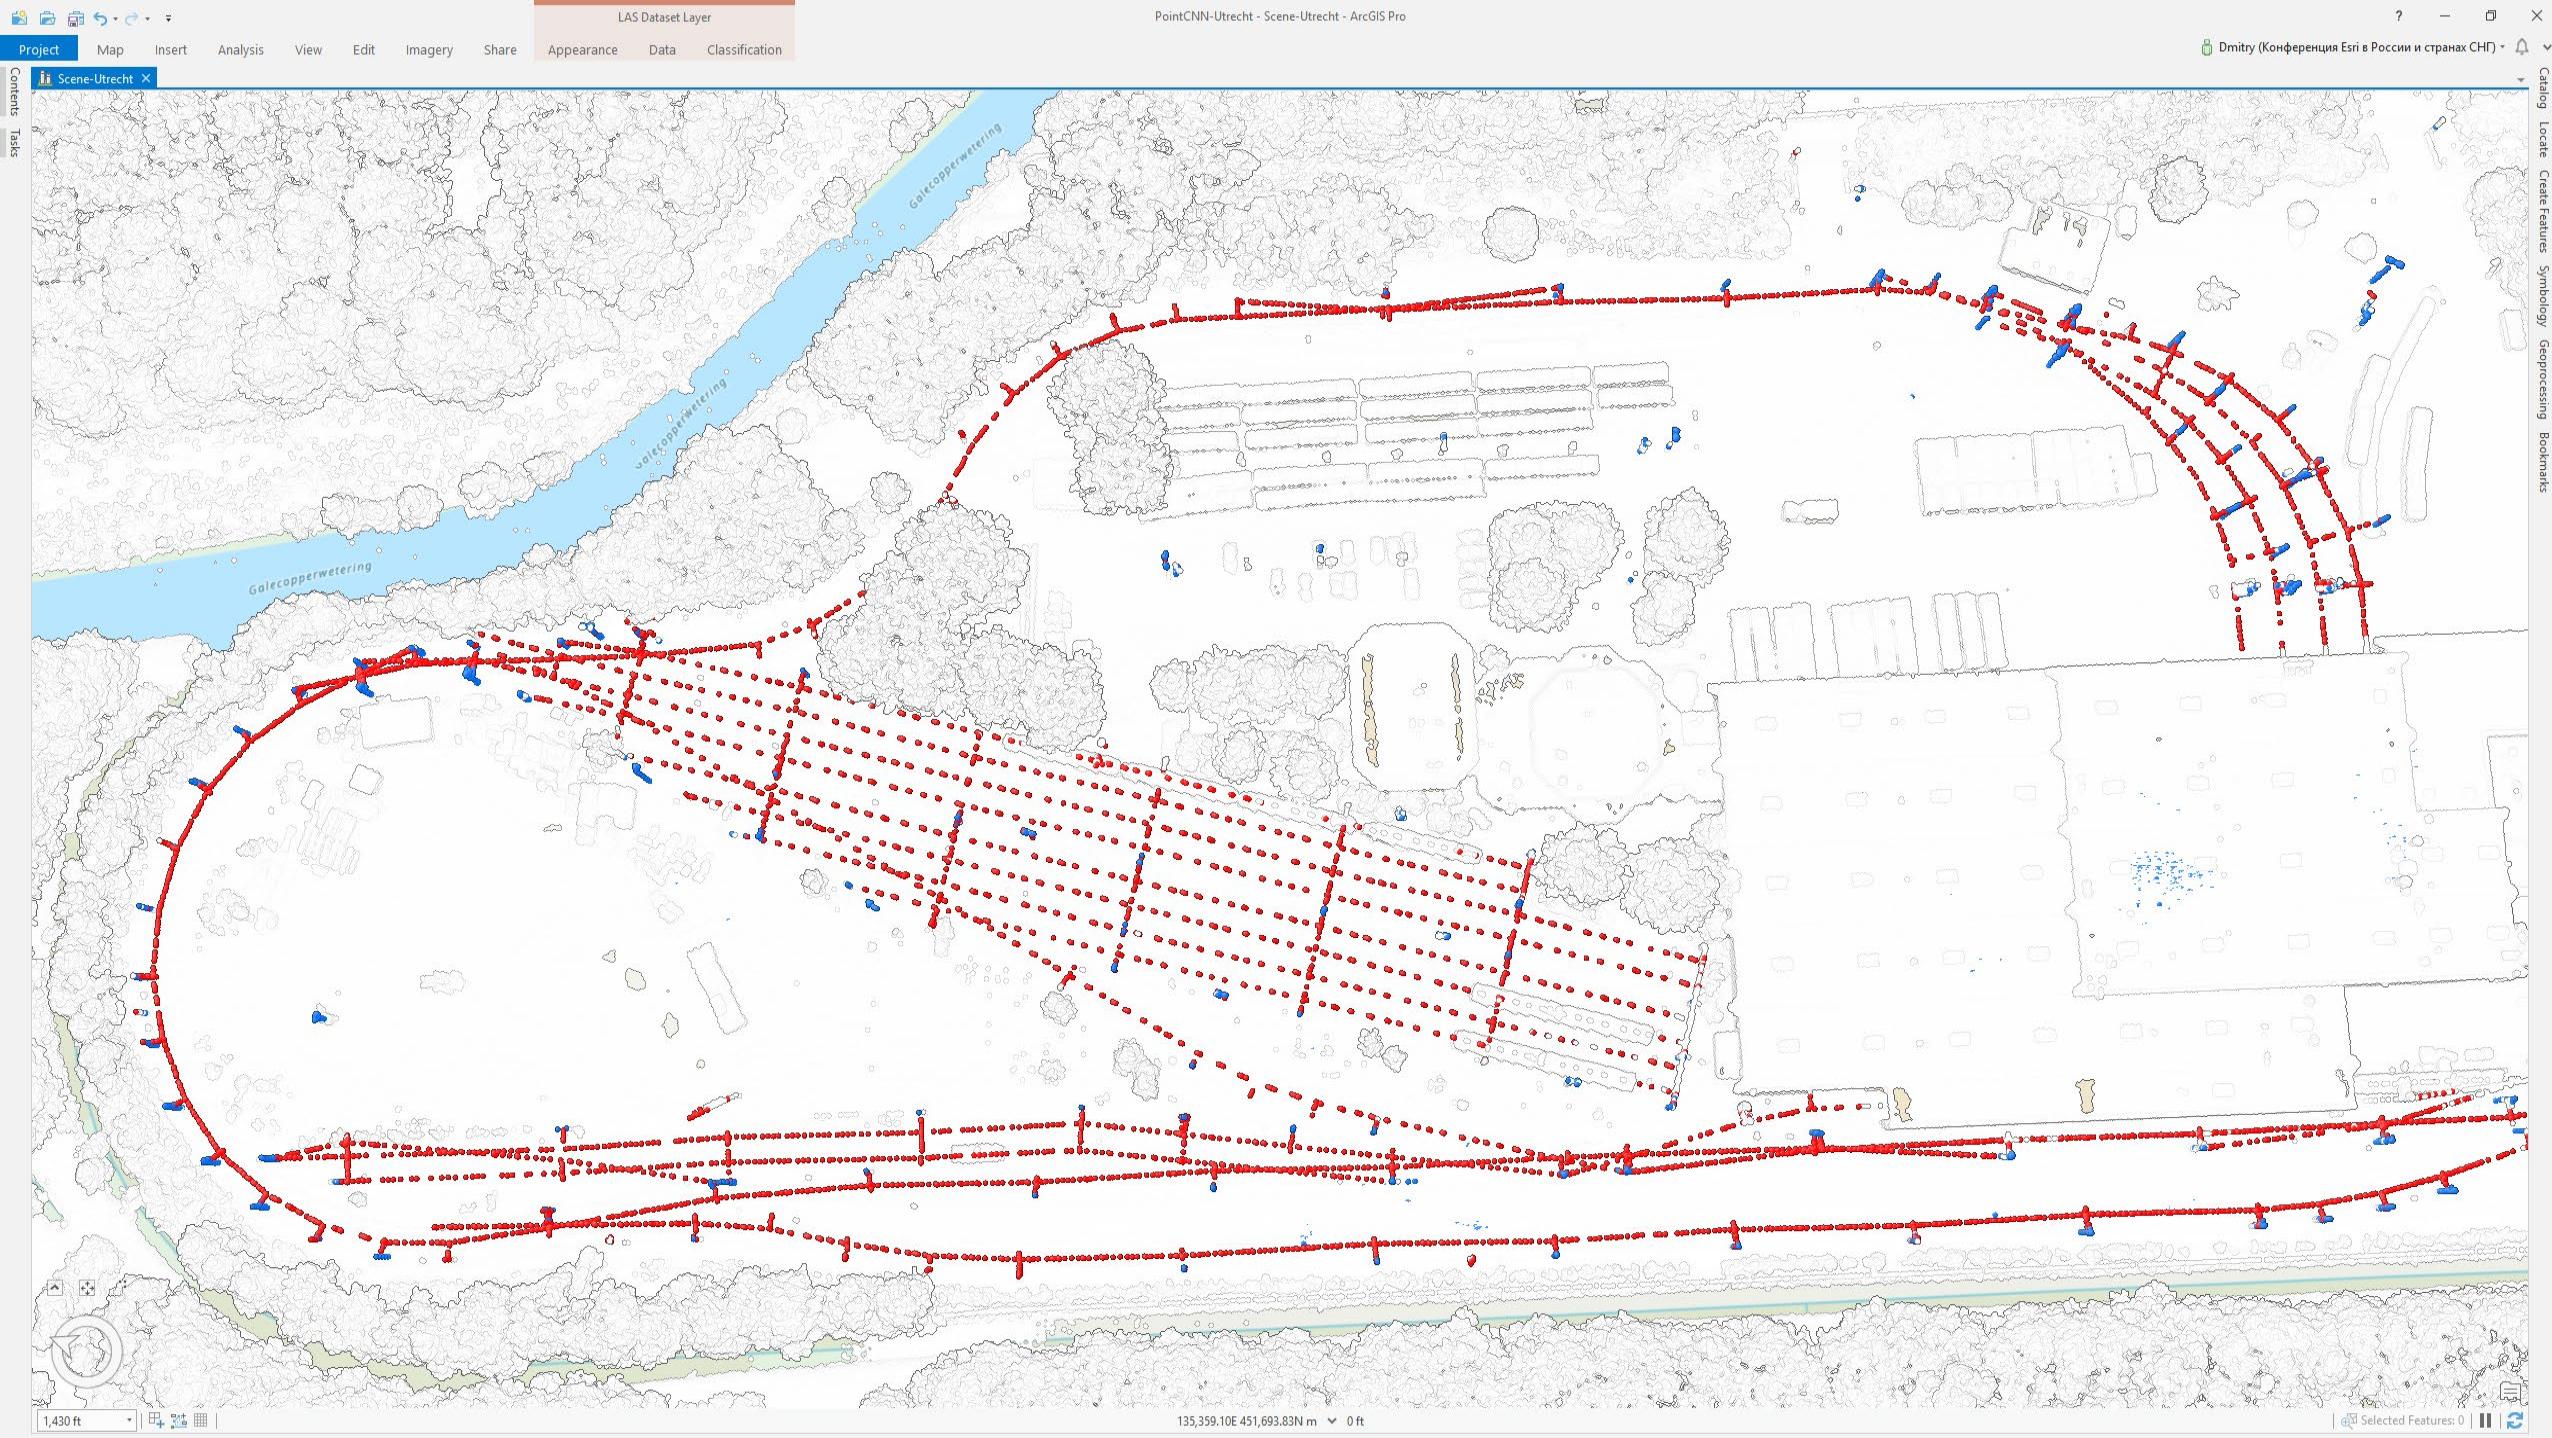












Railway Asset Detection

David Yu



THE
SCIENCE
OF
WHERE

Automating Rail Asset Detection

Using Deep Learning x GIS
David Yu



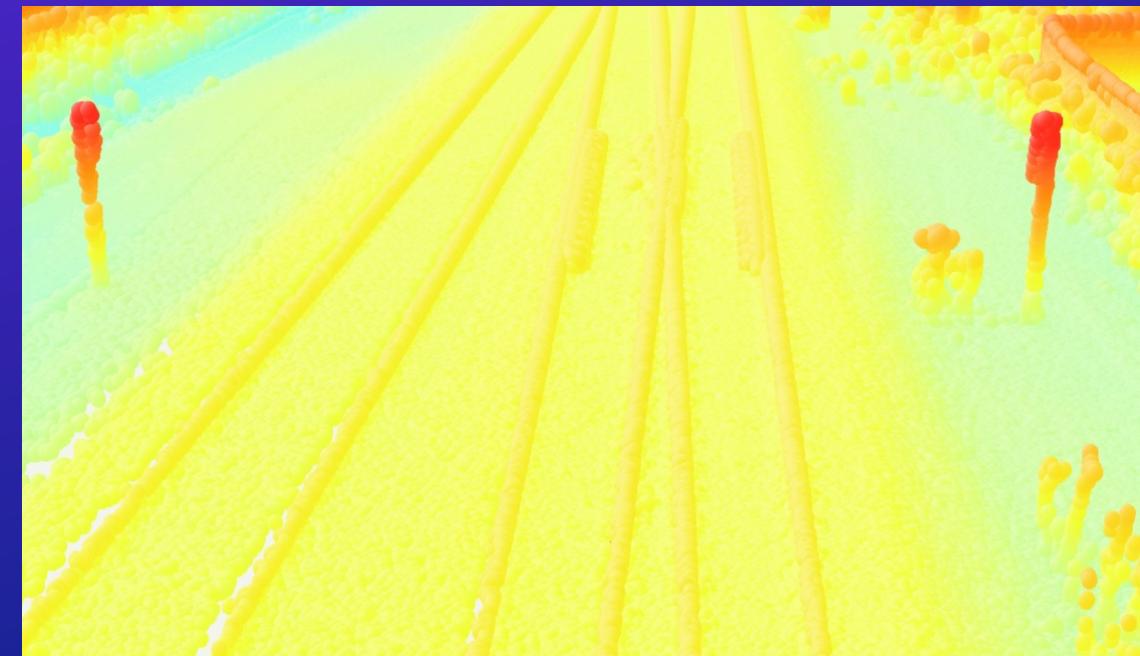
Challenge / Opportunity / Value

Challenge:

- Auditing, inventorying, managing change - at scale
- Asset collection & change detection

Value:

- PTC compliance
- Decreased cost
- Increased Accuracy
- Increased Frequency



Machine Learning / Input / Outcome

Conventional approach:

- Convert to image and apply CNN for feature extraction

Problem statement:

- Big data (e.g., over billion points)
- Train a deep learning model to detect objects directly from 3D point clouds?

Inputs:

- 3 subdivisions worth of labeled asset locations
 - Crossings, Mileposts, Switches, Signals
- 3D point clouds for study areas
- Track centerline

Outcomes:

- Geospatially enabled and trained deep learning model
- Found locations areas with previously unknown assets
- Identified mislabeled data

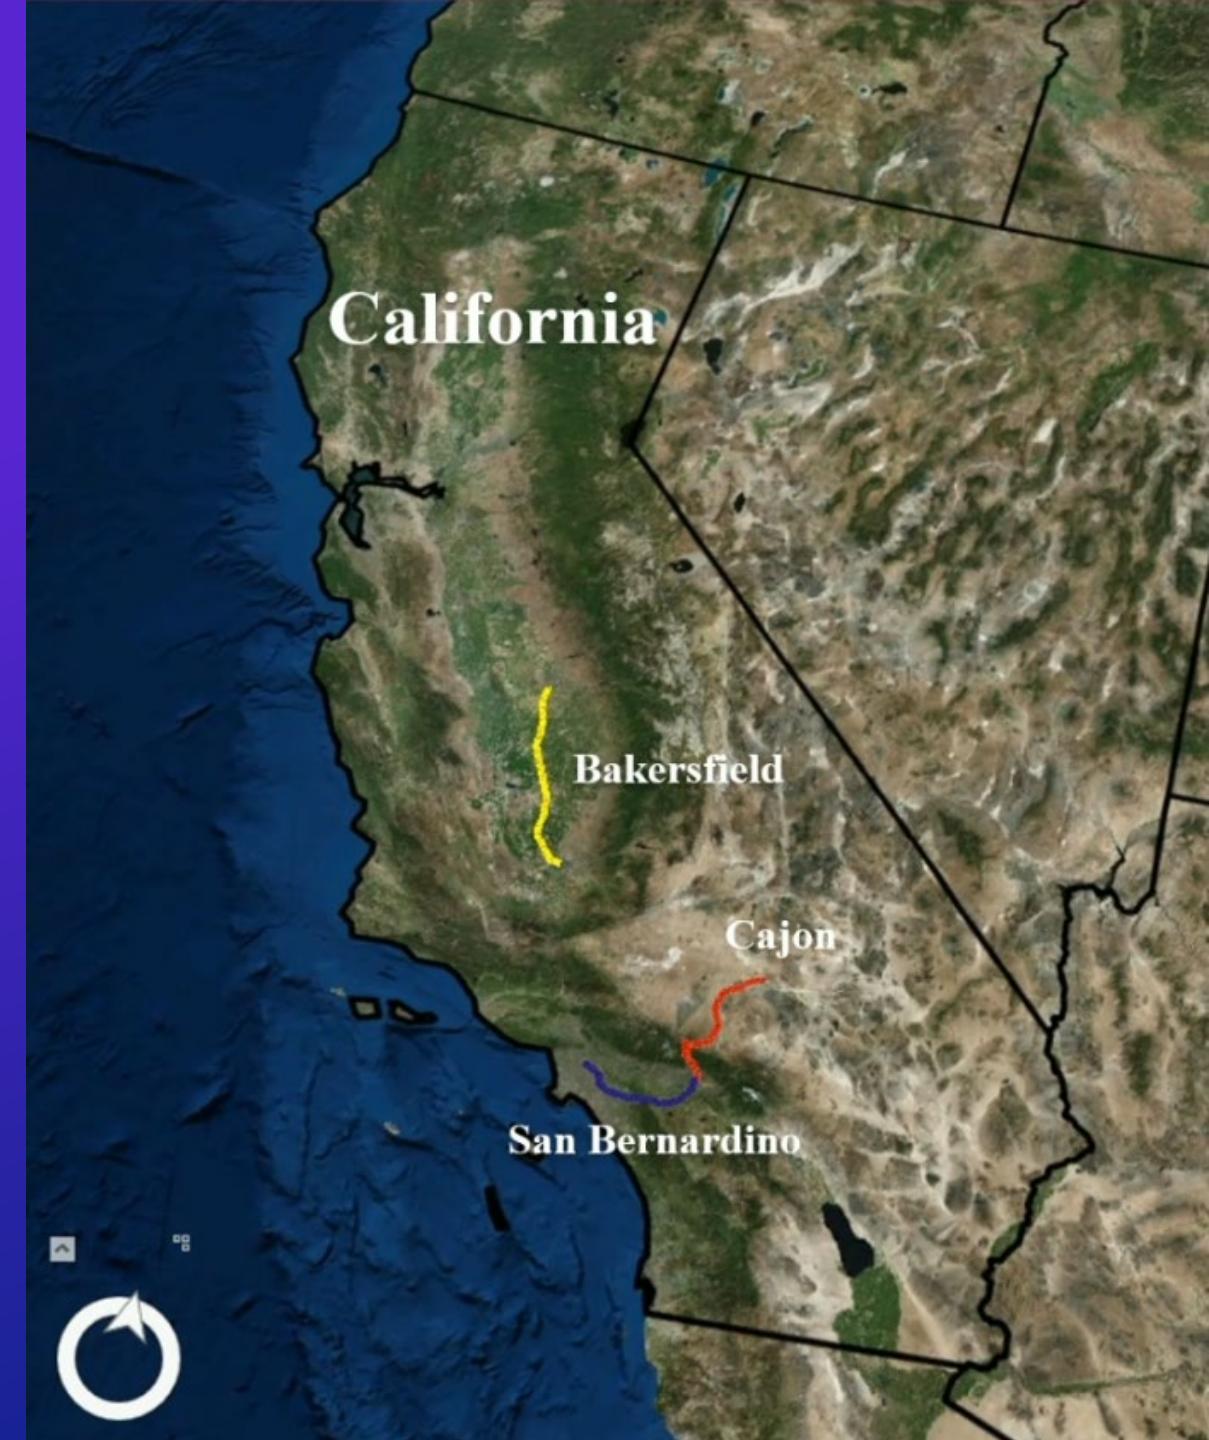
Study Area

Three regions in SoCal:

- Bakersfield
- Cajon
- San Bernardino

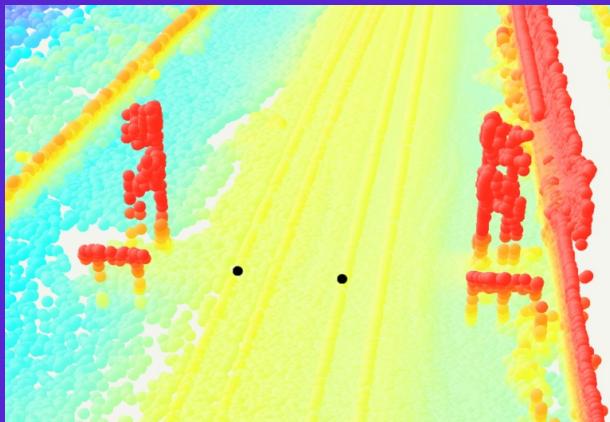
Data

- 3D point clouds for study areas
- Ortho images in TIFF format
- Labelled asset data (point feature class)
 - 150 labels per class per region
 - ~1,800 labels ($150 \times 4 \times 3$)

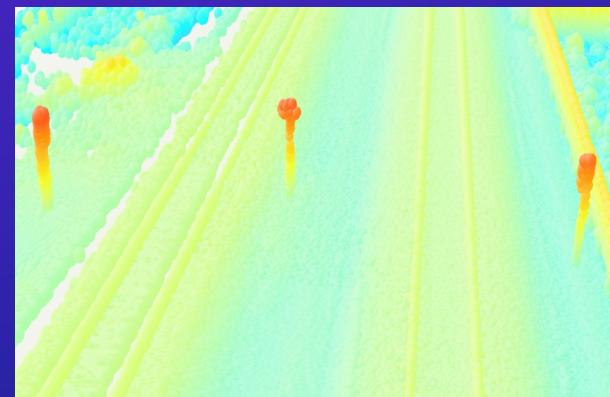
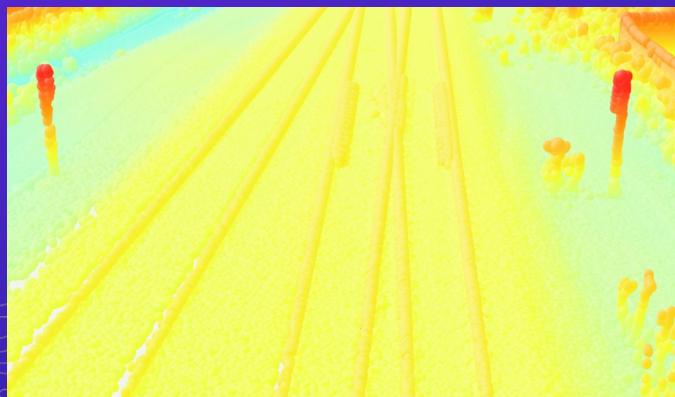


Sample of Objects

➤ Signal



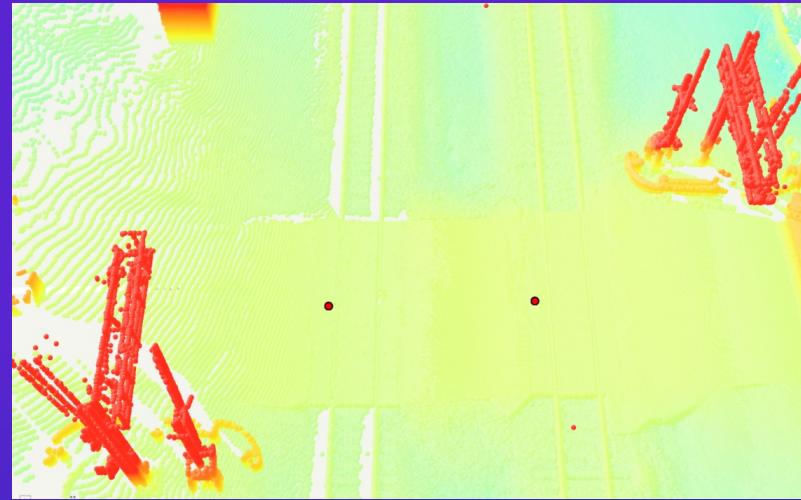
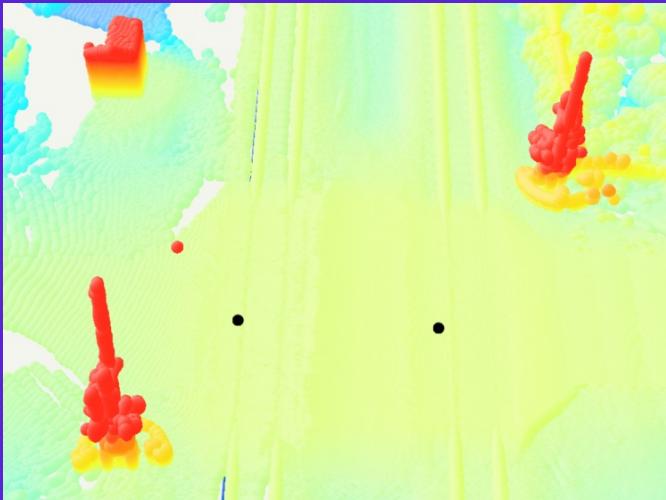
➤ Mile Post



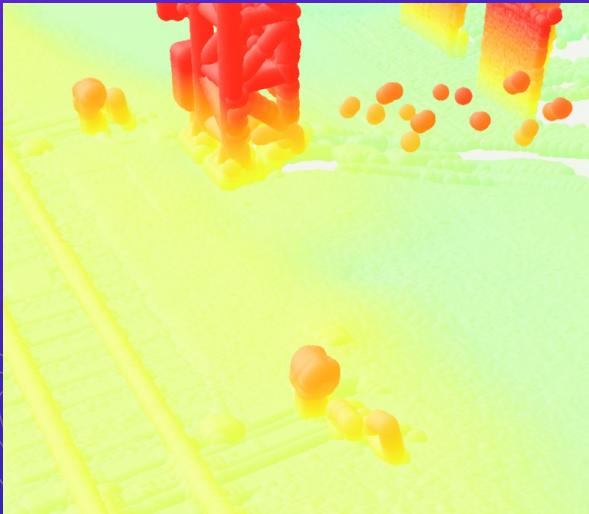
JIMBAUX.com

Sample of Objects

➤ Crossing



➤ Switch

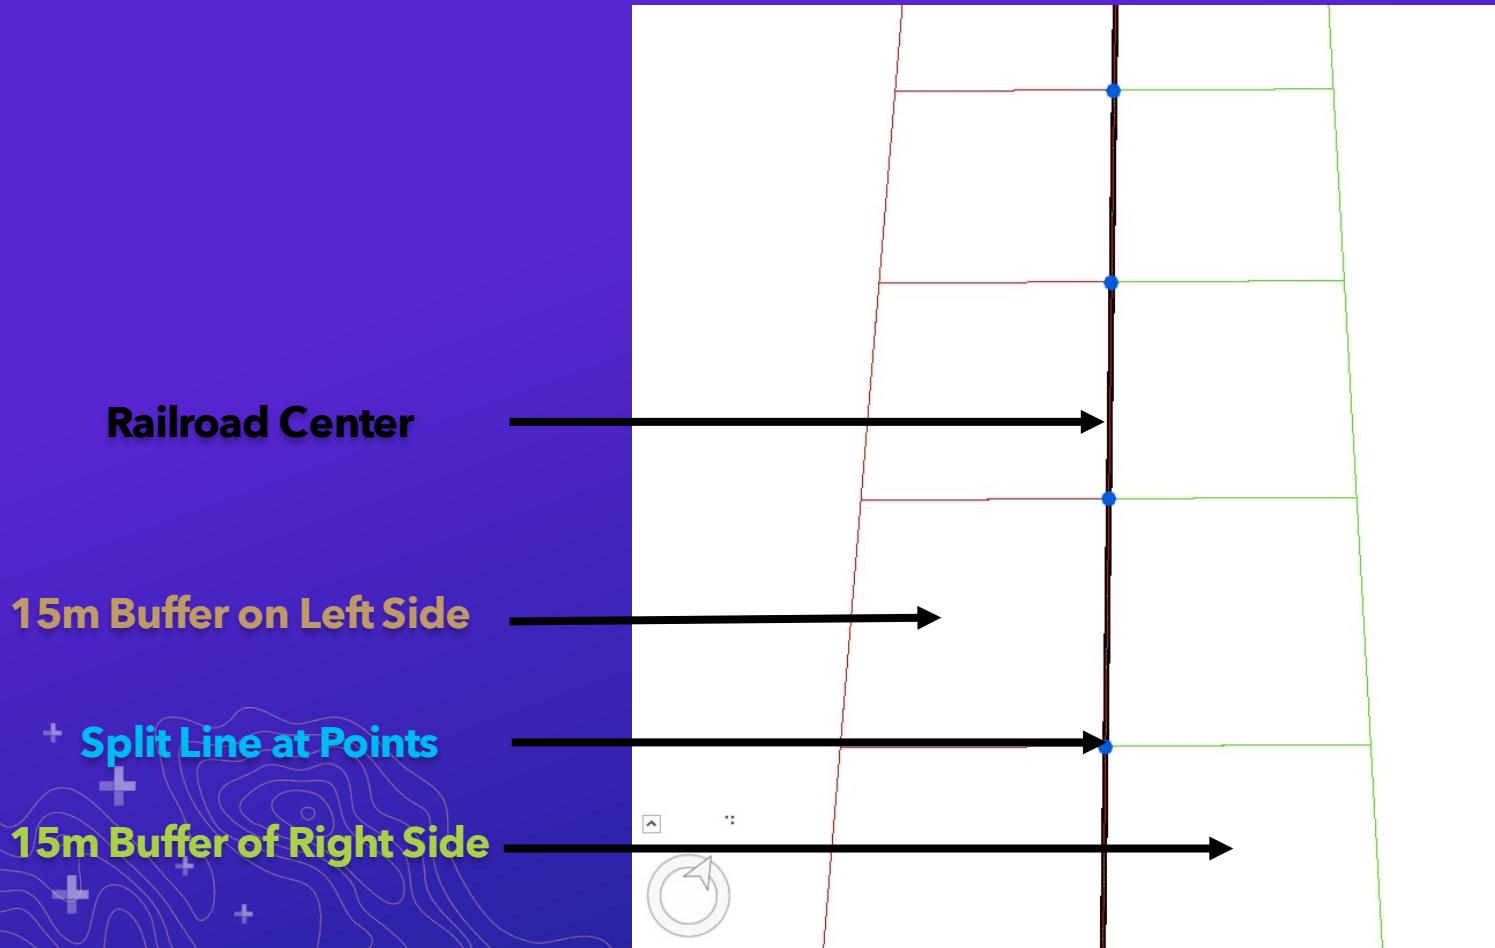


Data Processing, Modeling & Accuracy Assessment



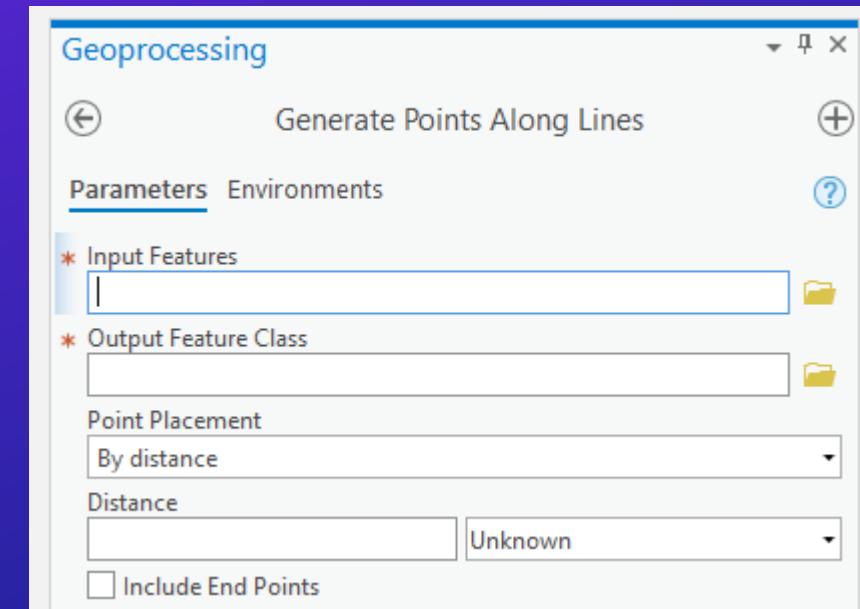
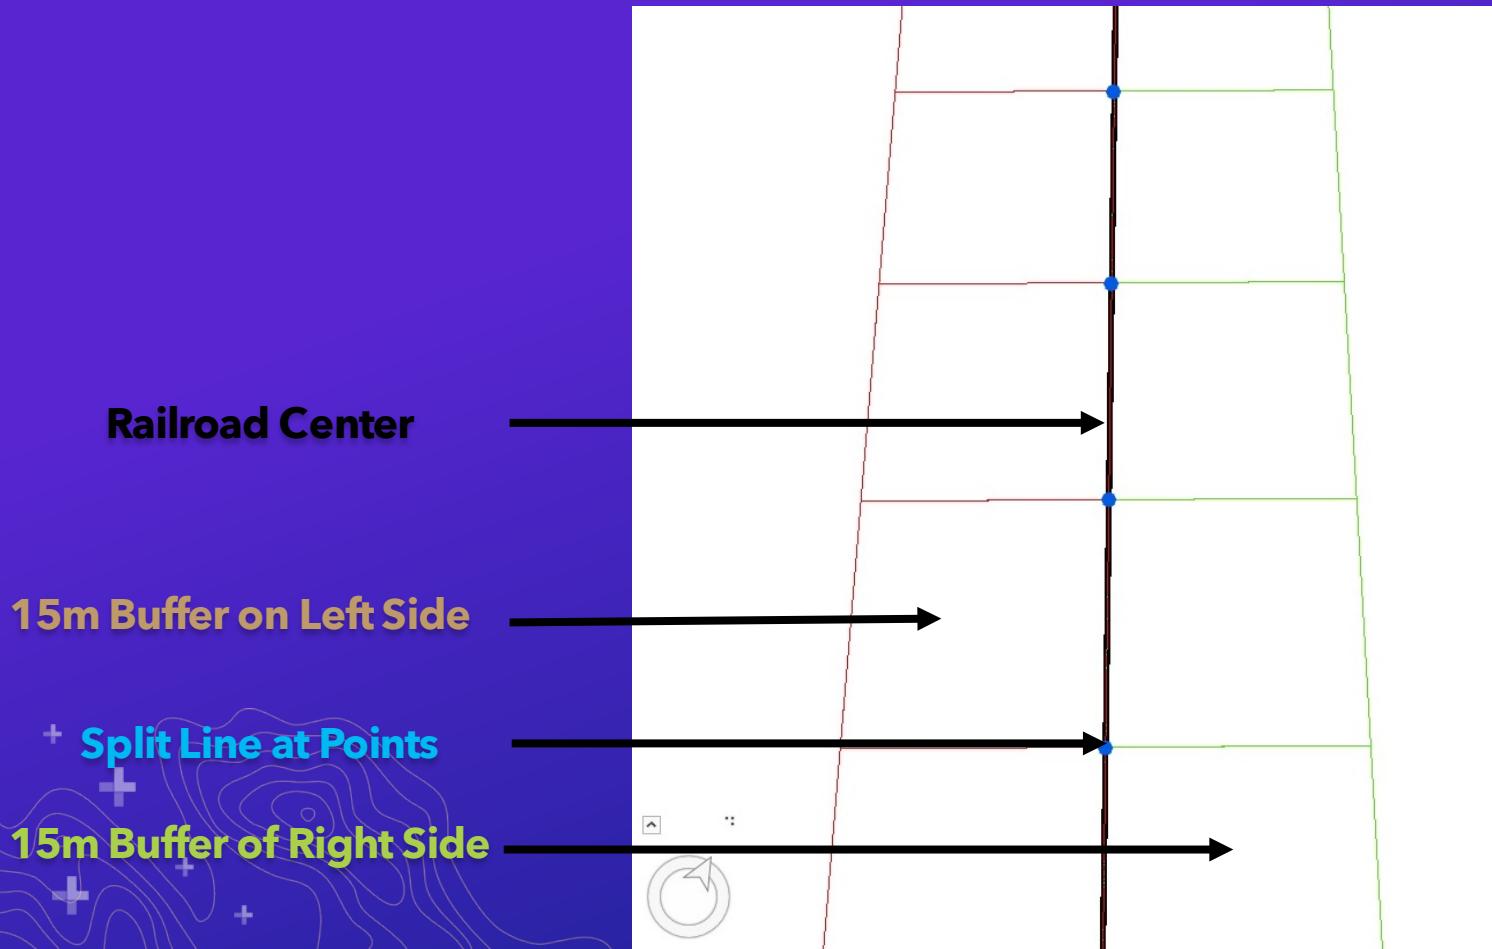
Data Processing

- Export Training Data for Deep Learning
- Create Unit of Analysis
- Identify Training Samples (Positive and Negative Samples)



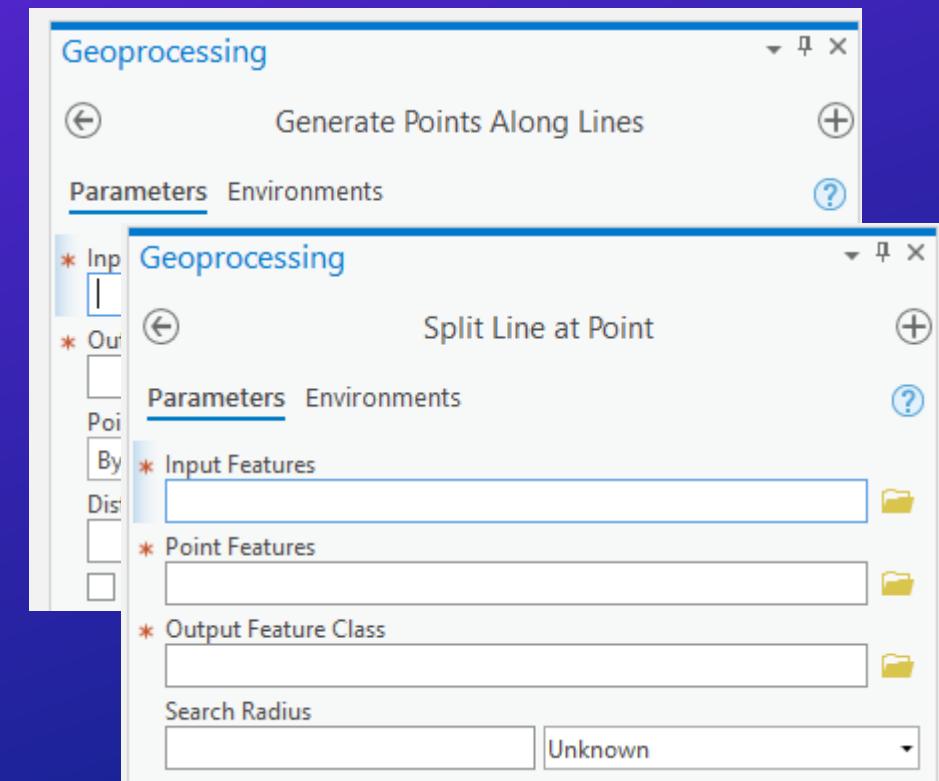
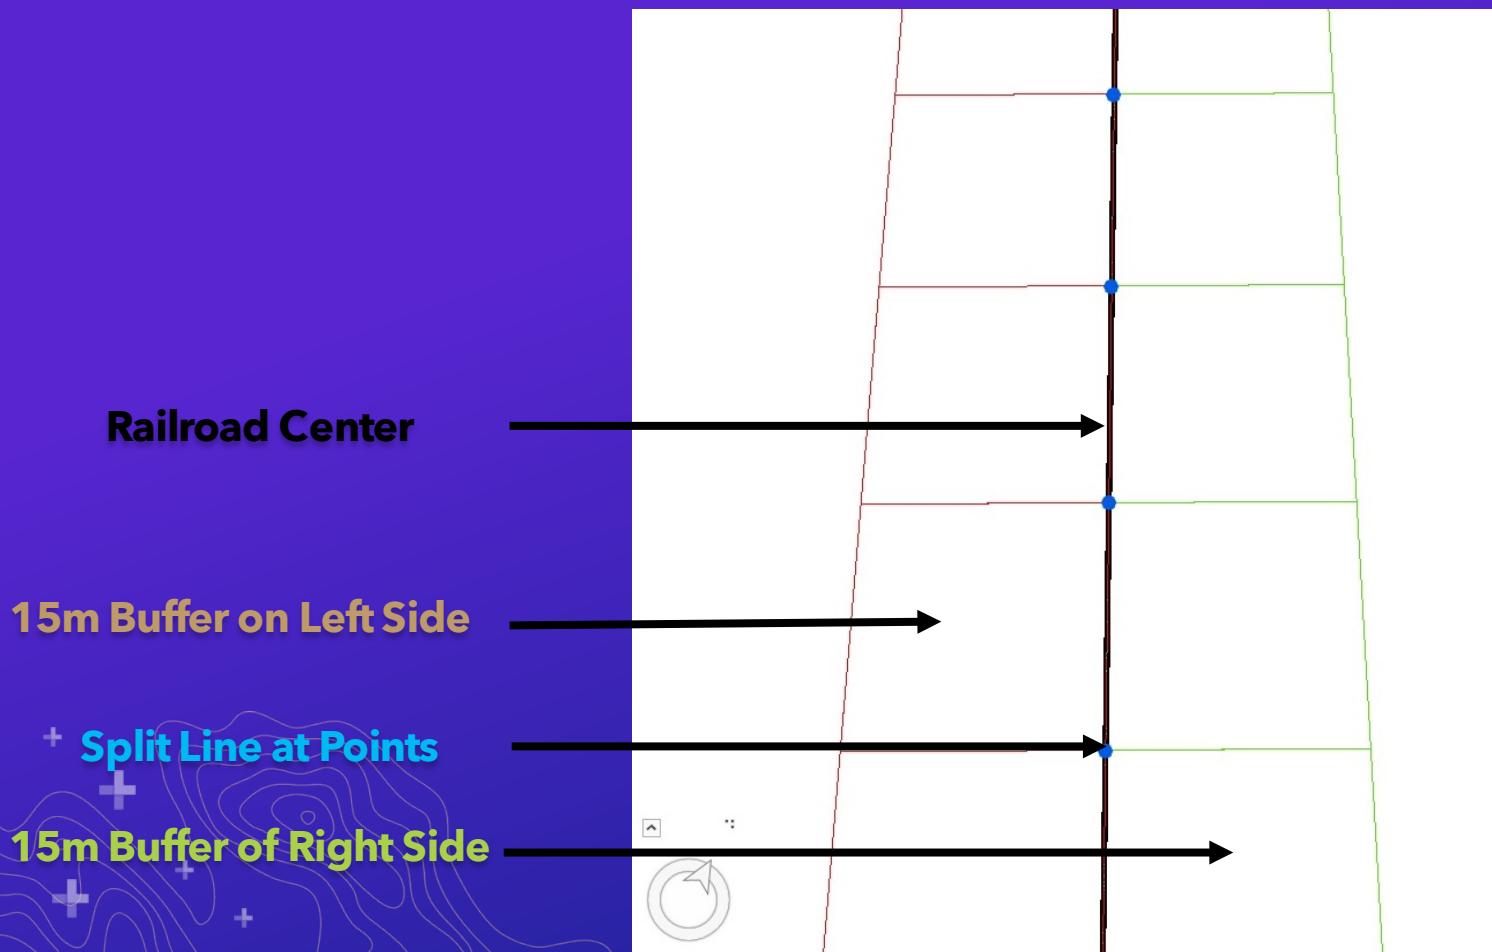
Data Processing

- Export Training Data for Deep Learning
- Create Unit of Analysis
- Identify Training Samples (Positive and Negative Samples)



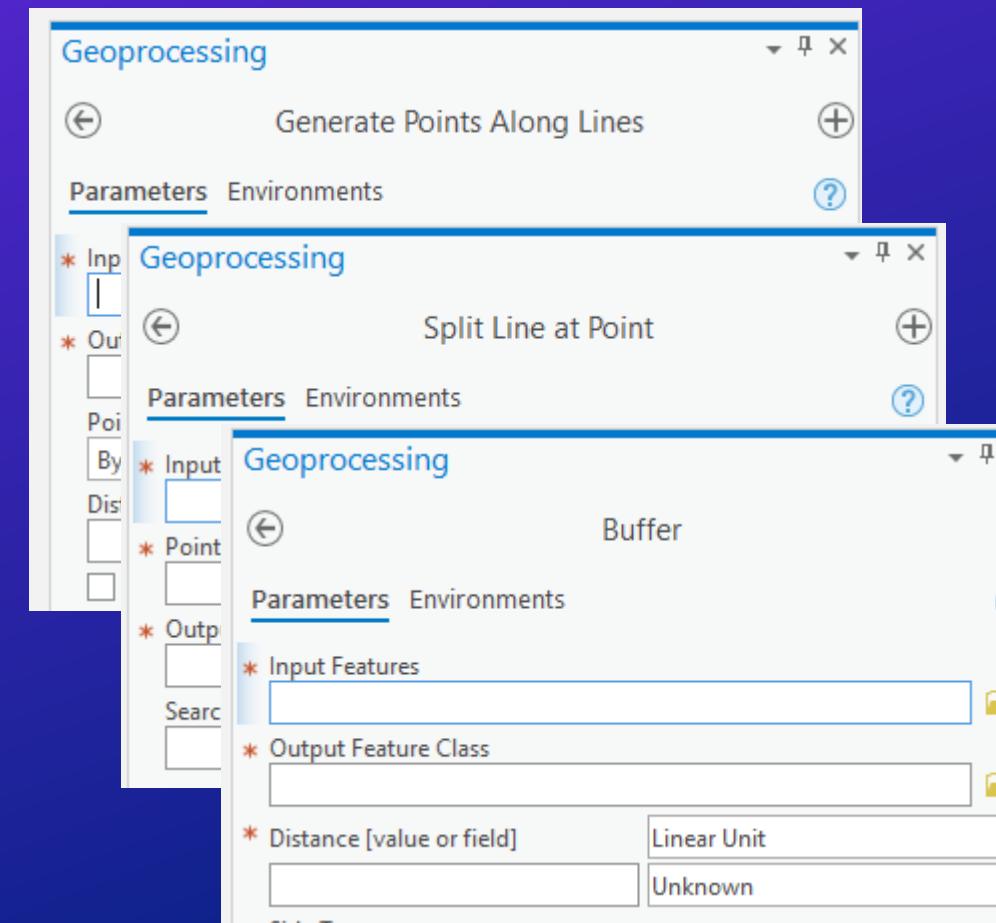
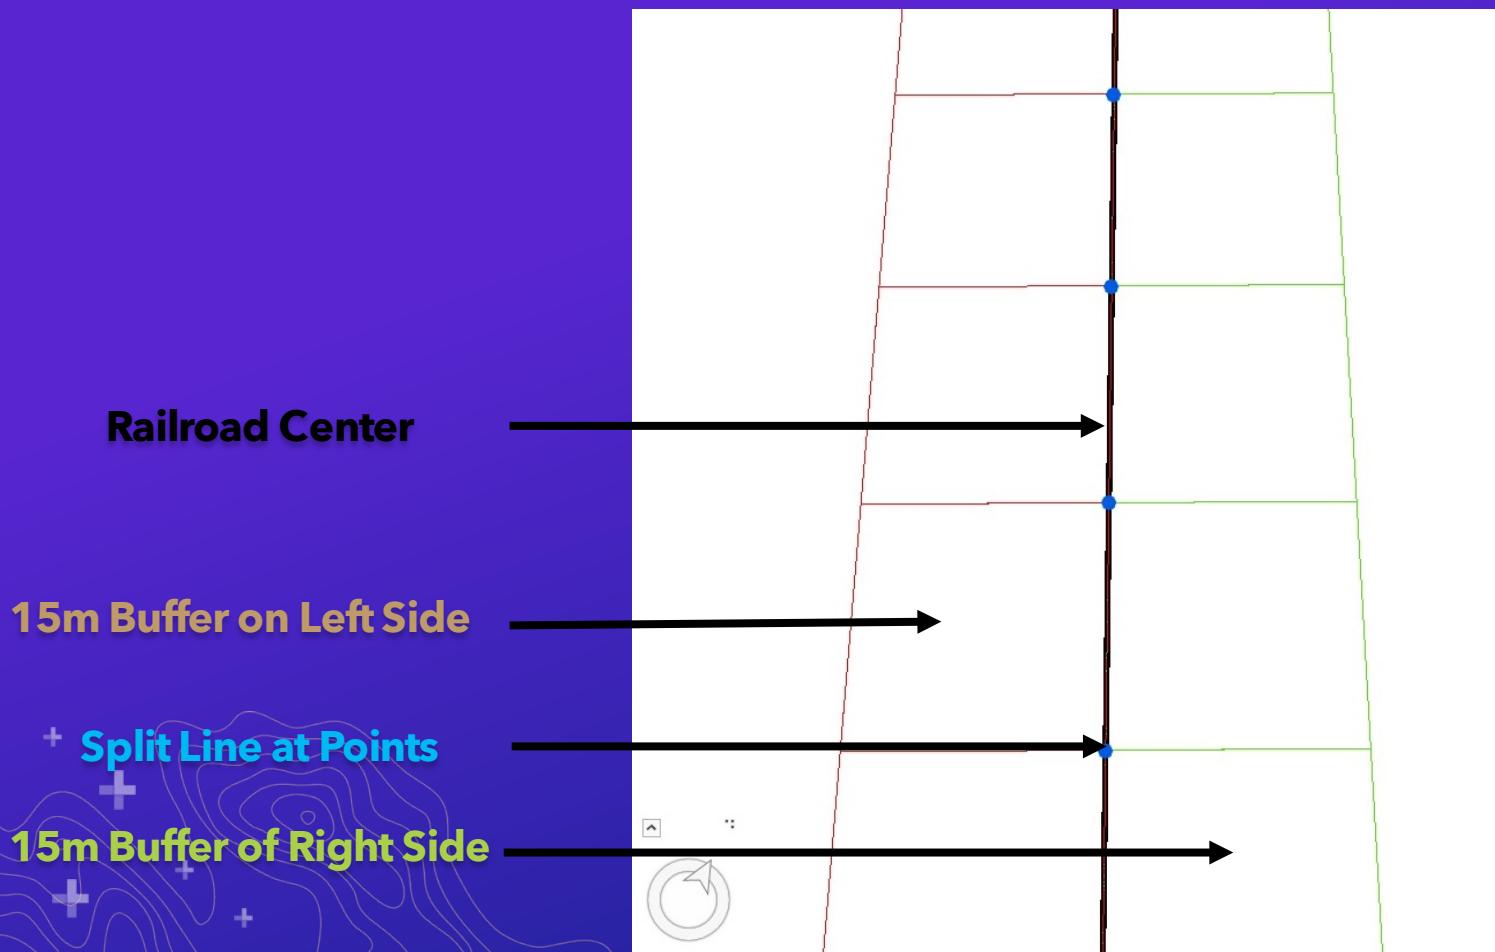
Data Processing

- Export Training Data for Deep Learning
- Create Unit of Analysis
- Identify Training Samples (Positive and Negative Samples)



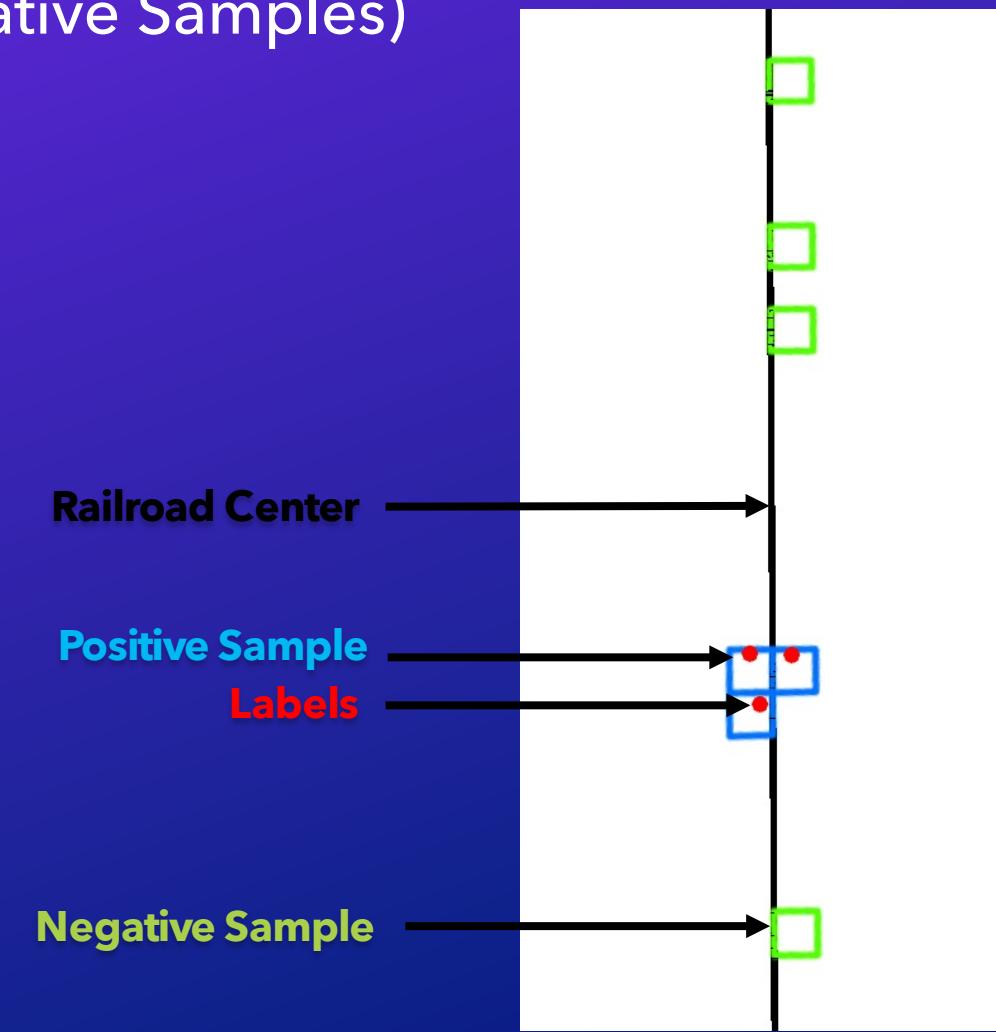
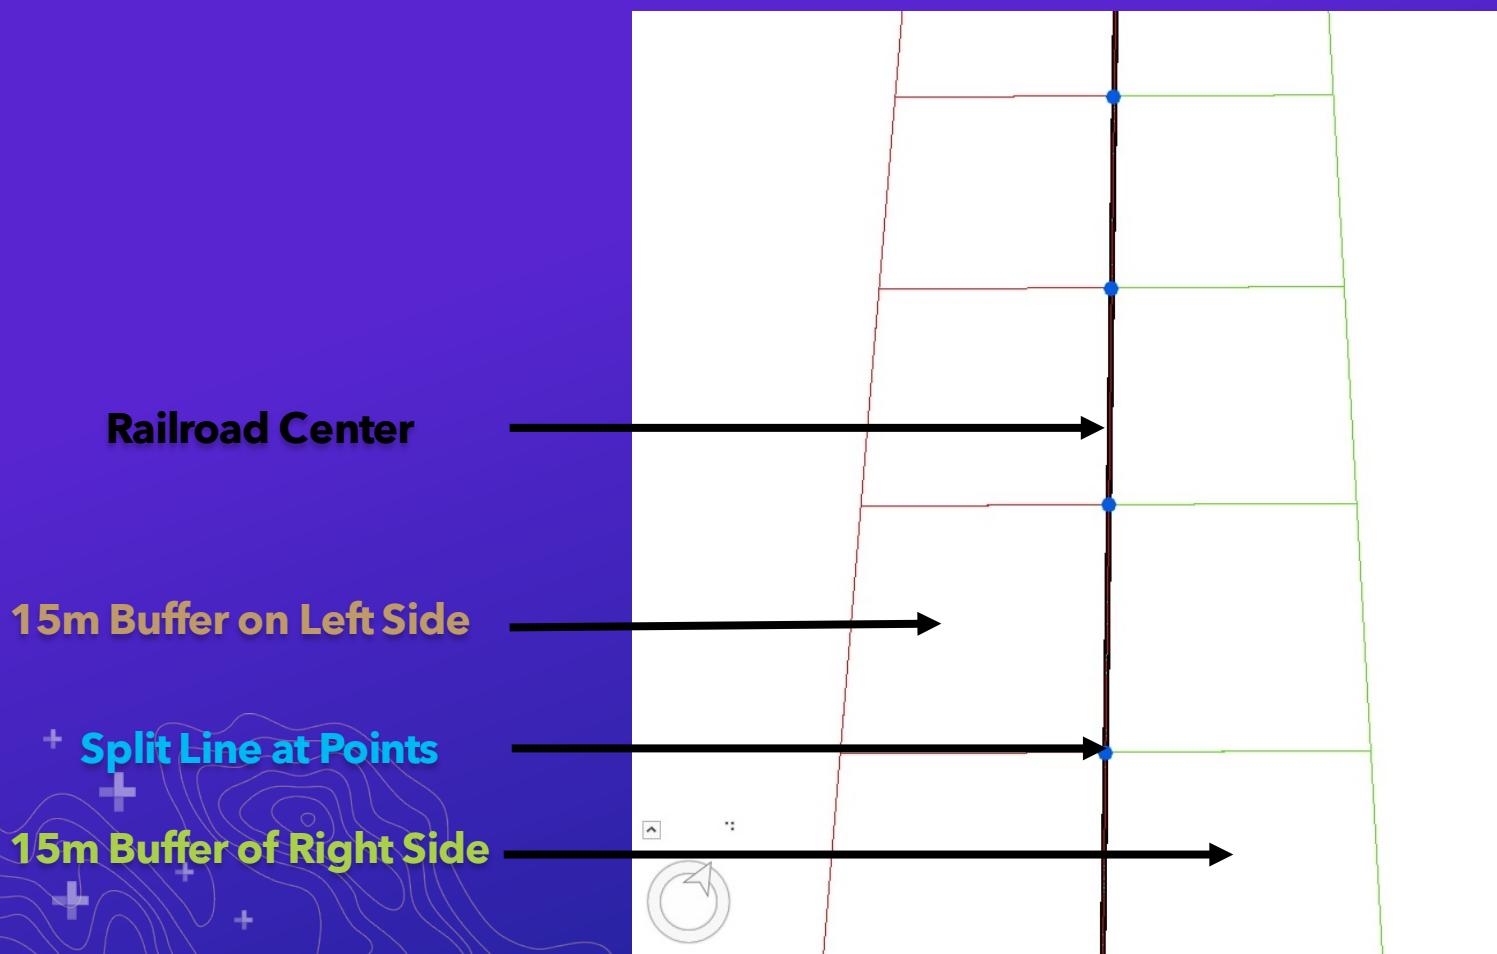
Data Processing

- Export Training Data for Deep Learning
- Create Unit of Analysis
- Identify Training Samples (Positive and Negative Samples)



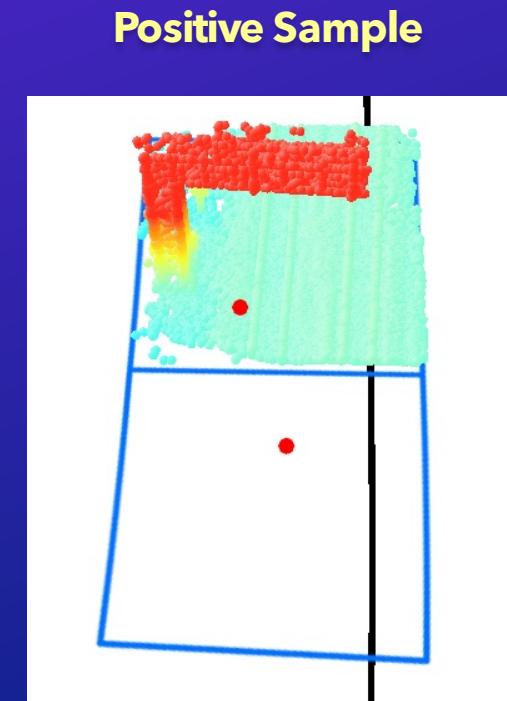
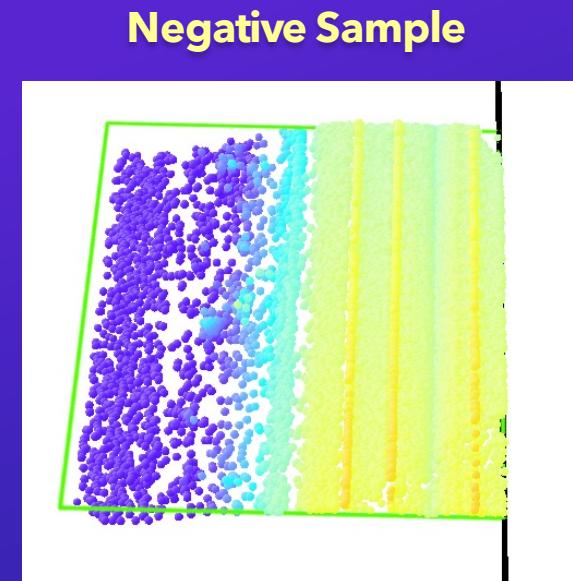
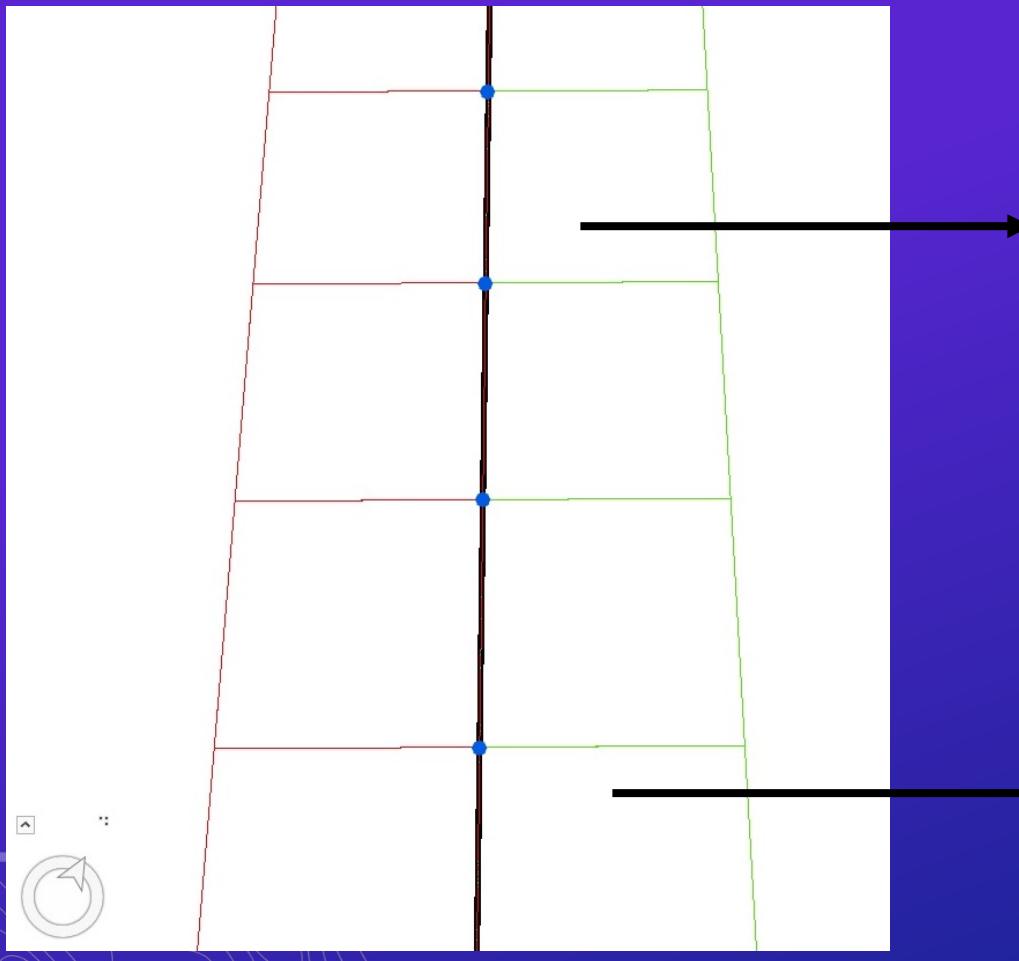
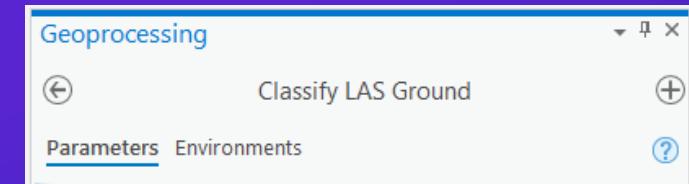
Data Processing

- Export Training Data for Deep Learning
- Create Unit of Analysis
- Identify Training Samples (Positive and Negative Samples)



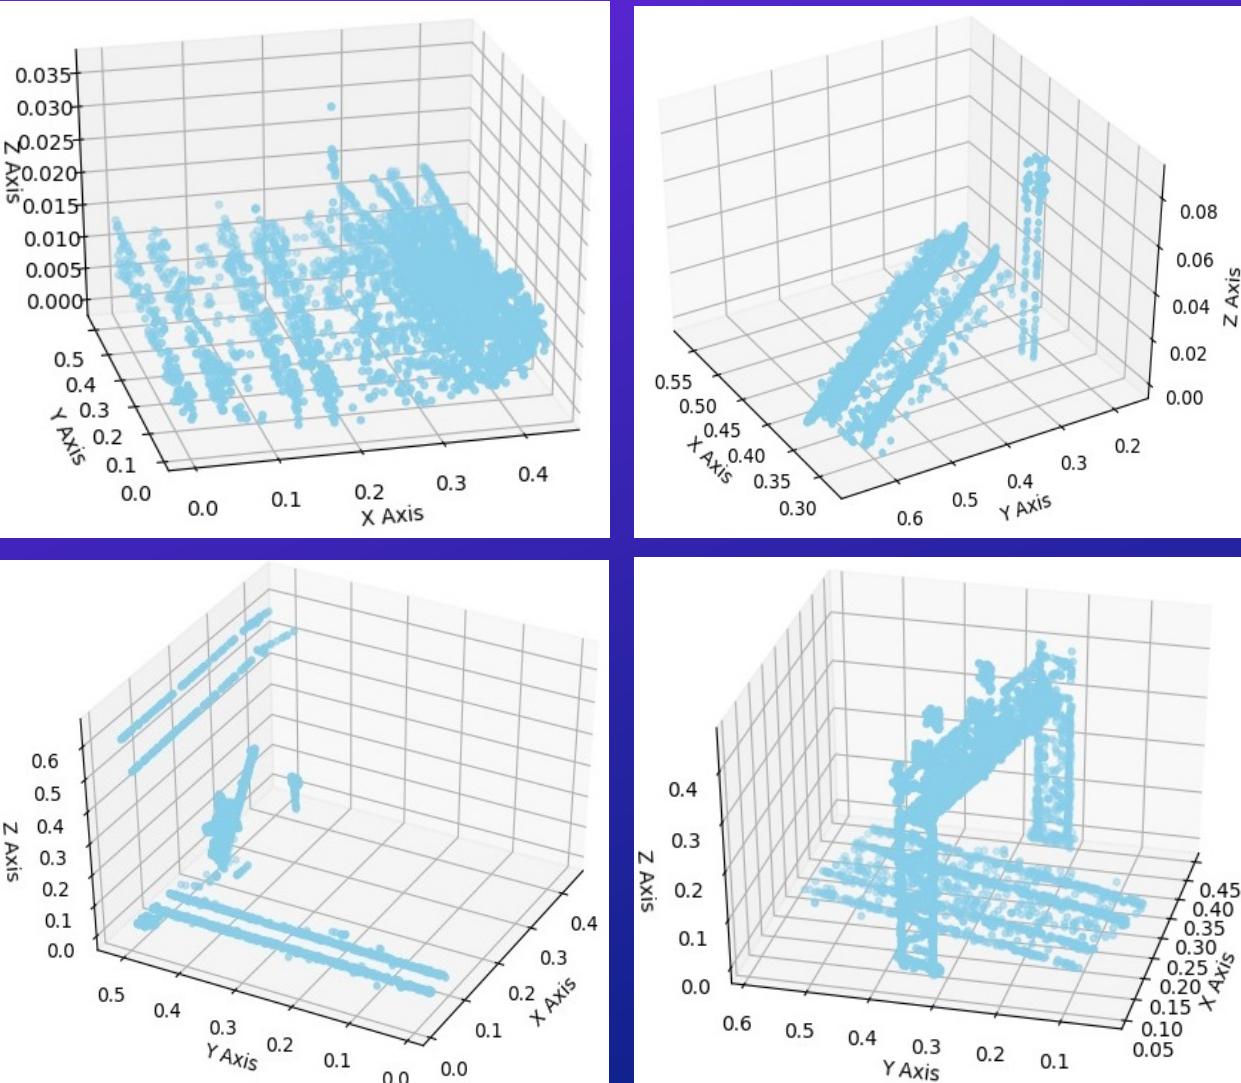
Data Processing

- Clip original LAS files to boxes



Data Processing

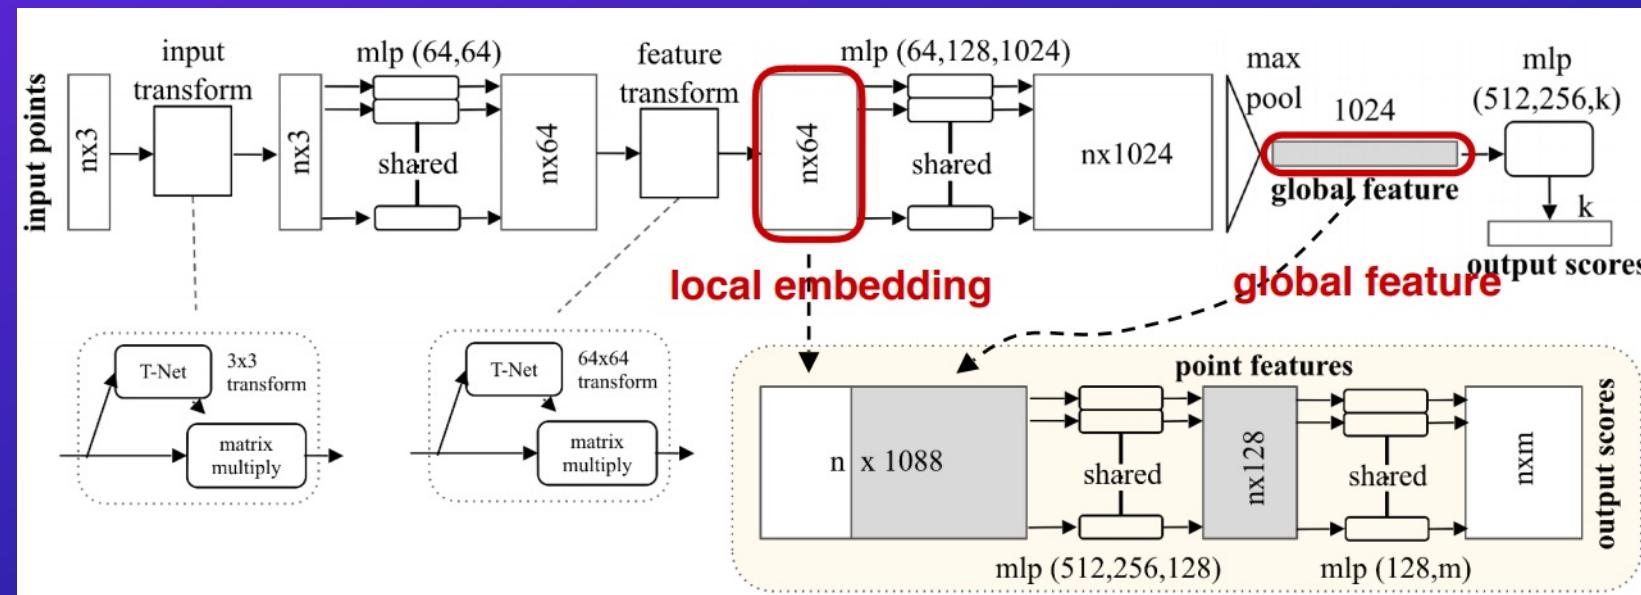
- Spatial Random Sampling
- Each 15m box had 50k - 200k points
- Optimal # of points for model ~4k points
- Train PointNet structure to use X, Y, relative Z, and Intensity



Model Training with GPU

- Training = Two Railways
- Testing = One Railway

Class	Training Run	Testing Run
Signal	242	115
Mile Post	302	103
Crossing	211	85
Switch	272	246
Negative	1667	654



Accuracy Assessment

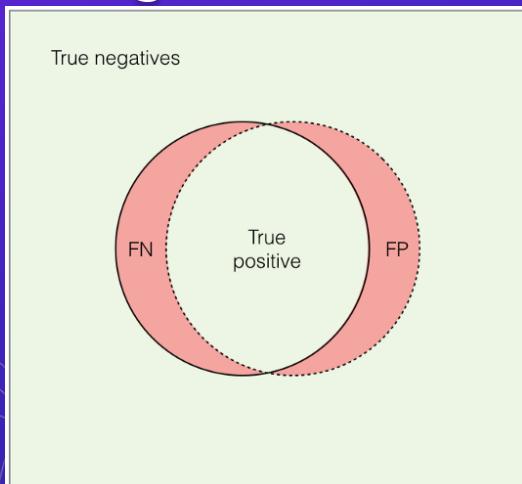
Training Data

Class	Precision	Recall
Negative	0.89	0.88
Post	0.72	0.76
Crossing	0.85	0.81
Signal	0.80	0.87
Switch	0.55	0.53

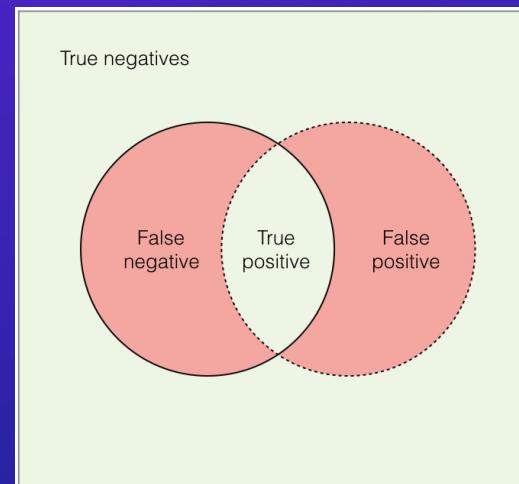
Testing Data

Class	Precision	Recall
Negative	0.86	0.85
Post	0.70	0.73
Crossing	0.74	0.72
Signal	0.76	0.75
Switch	0.48	0.51

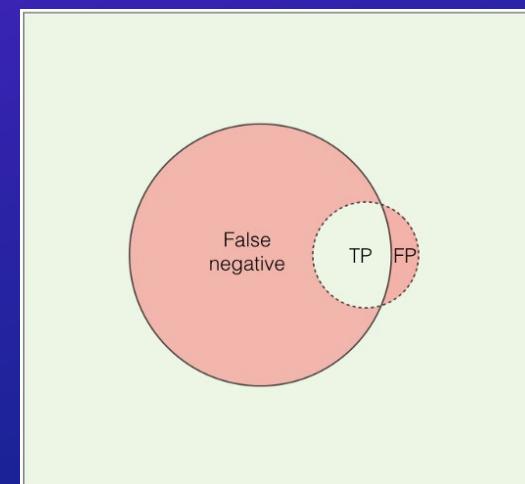
**High Precision
High Recall**



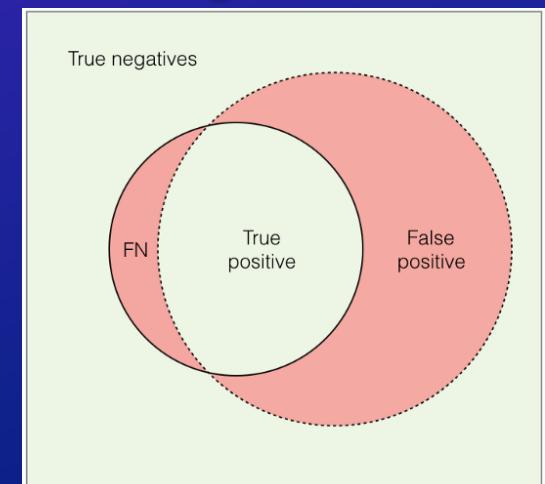
**Low Precision
Low Recall**



**High Precision
Low Recall**

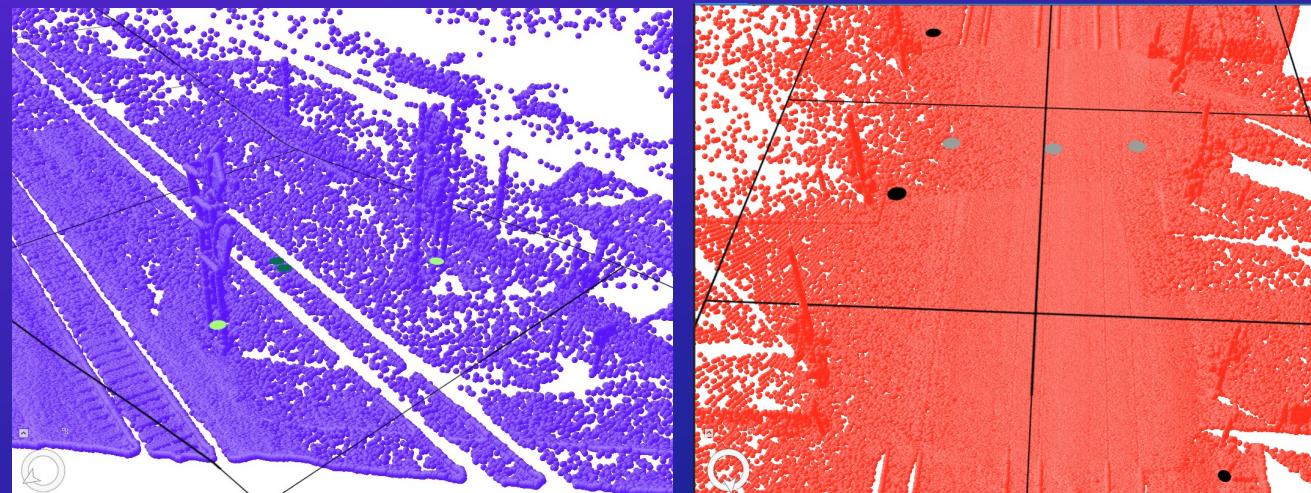
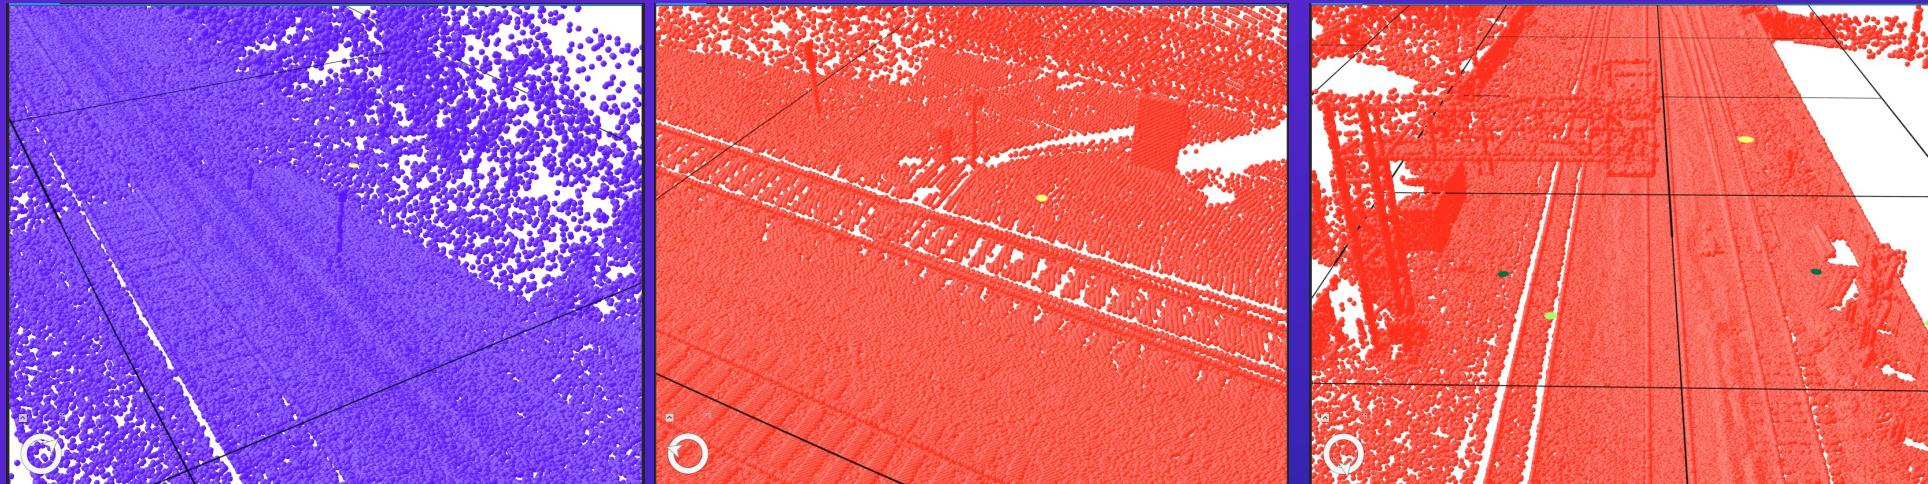


**Low Precision
High Recall**



Output Results

- Point feature class of identified asset locations
- Attribution classifying each point as specific asset type





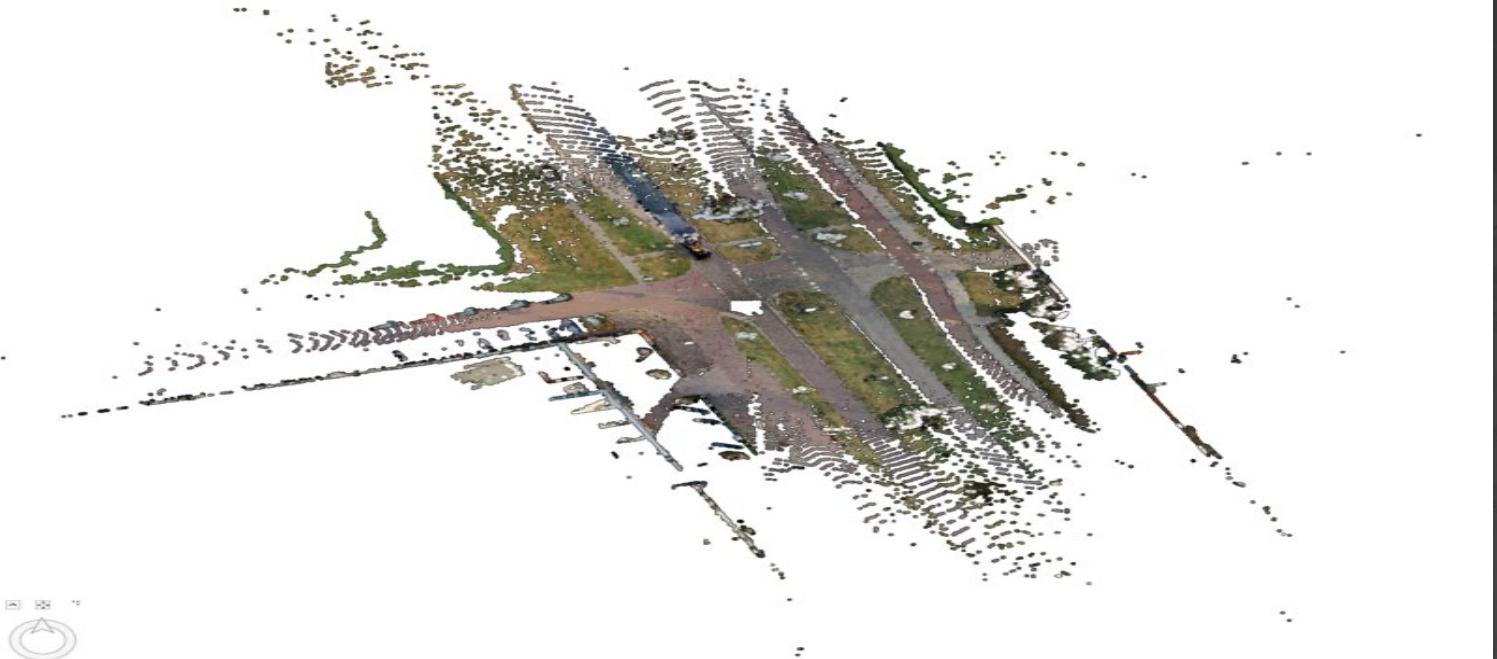
PointCNN & Feature Extraction: mobile point clouds



THE
SCIENCE
OF
WHERE

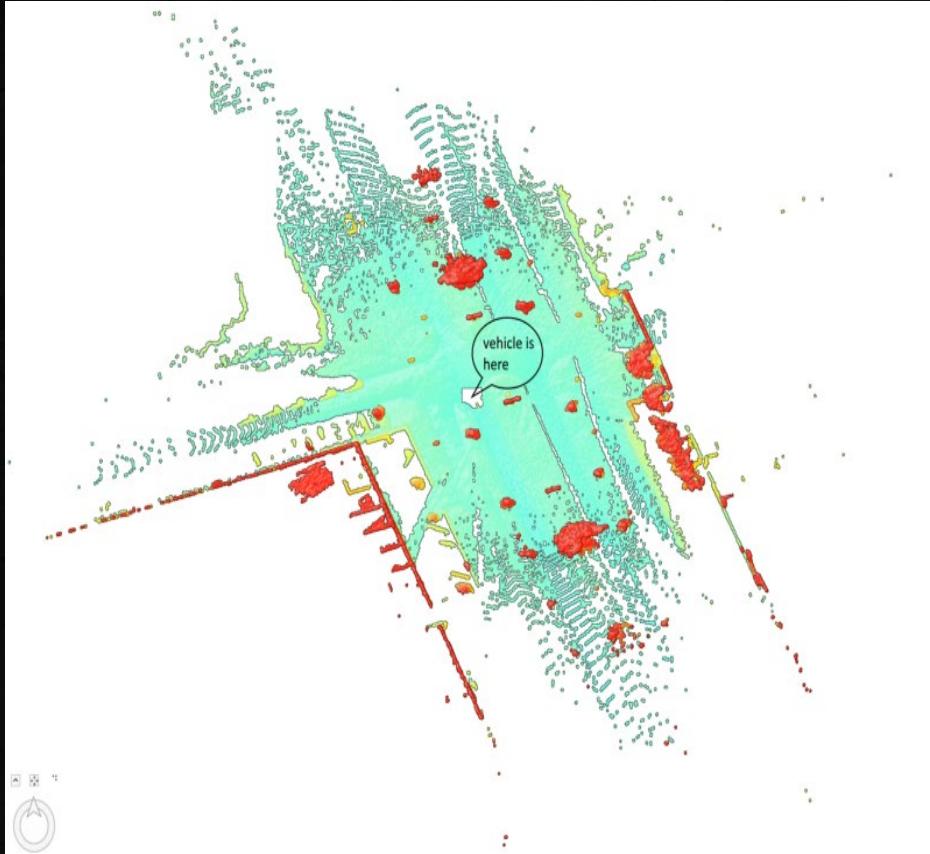
PoC: PointCNN in Mobile Point Clouds

CycloMedia's "CycloRama" sample point cloud with synced 360-imagery



PoC: PointCNN in Mobile Point Clouds

Mobile point clouds have two peculiarities making them harder to work with compared to airborne data:



1. Point Density

2. Gaps in Data

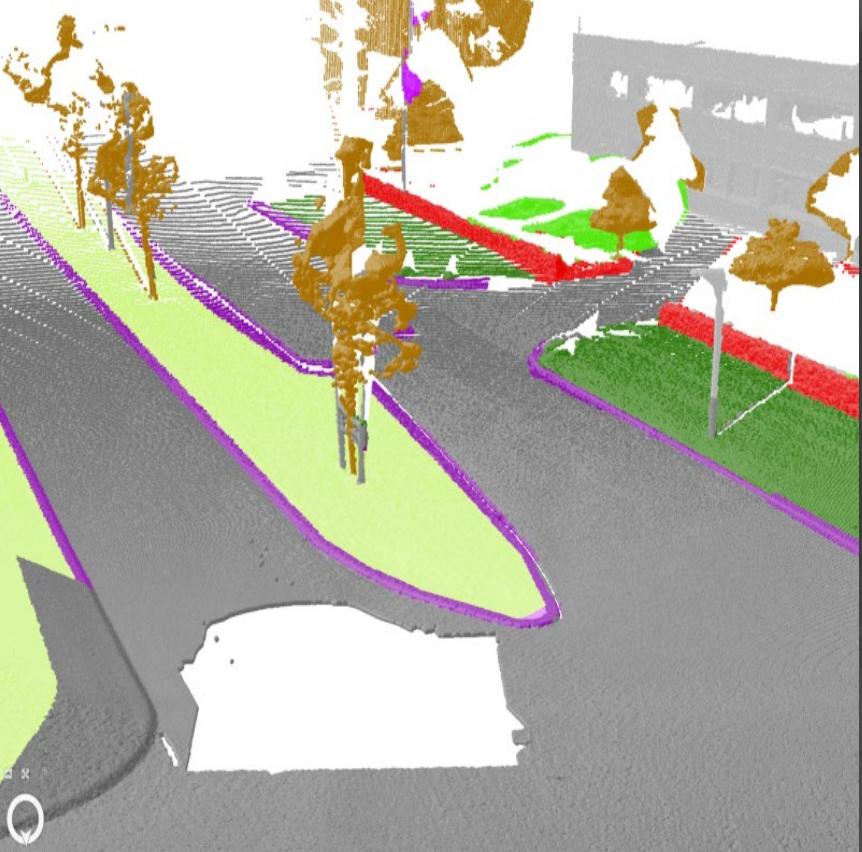


PoC: PointCNN in Mobile Point Clouds

An advantage: simpler to label and create training data

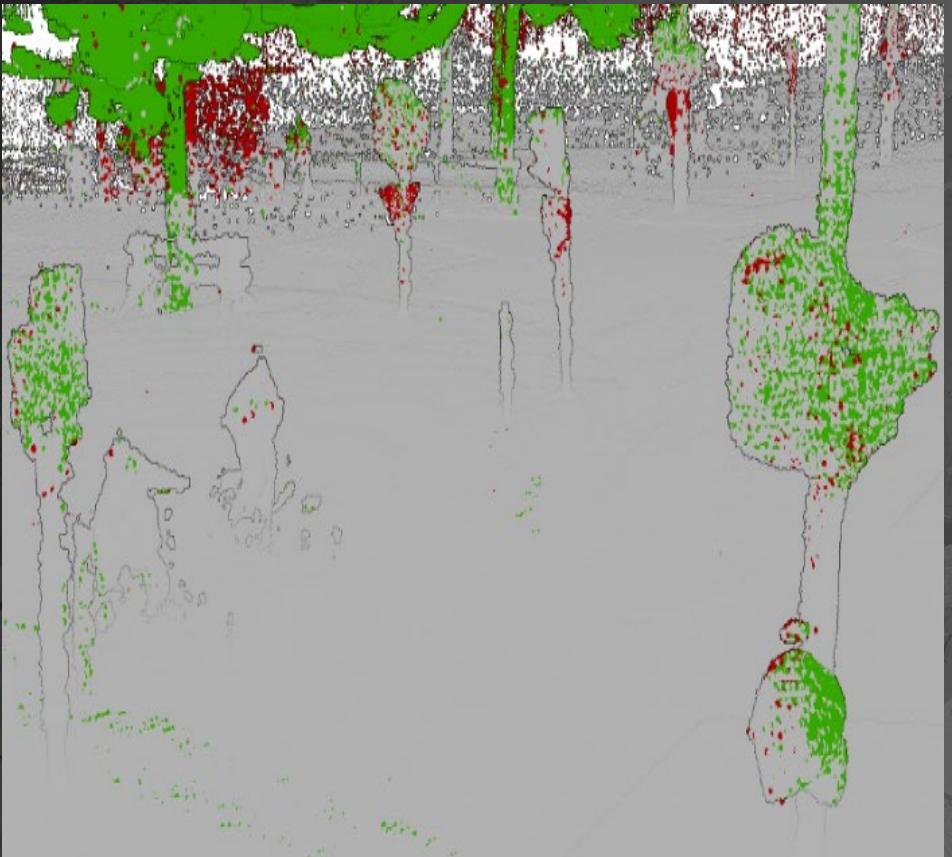
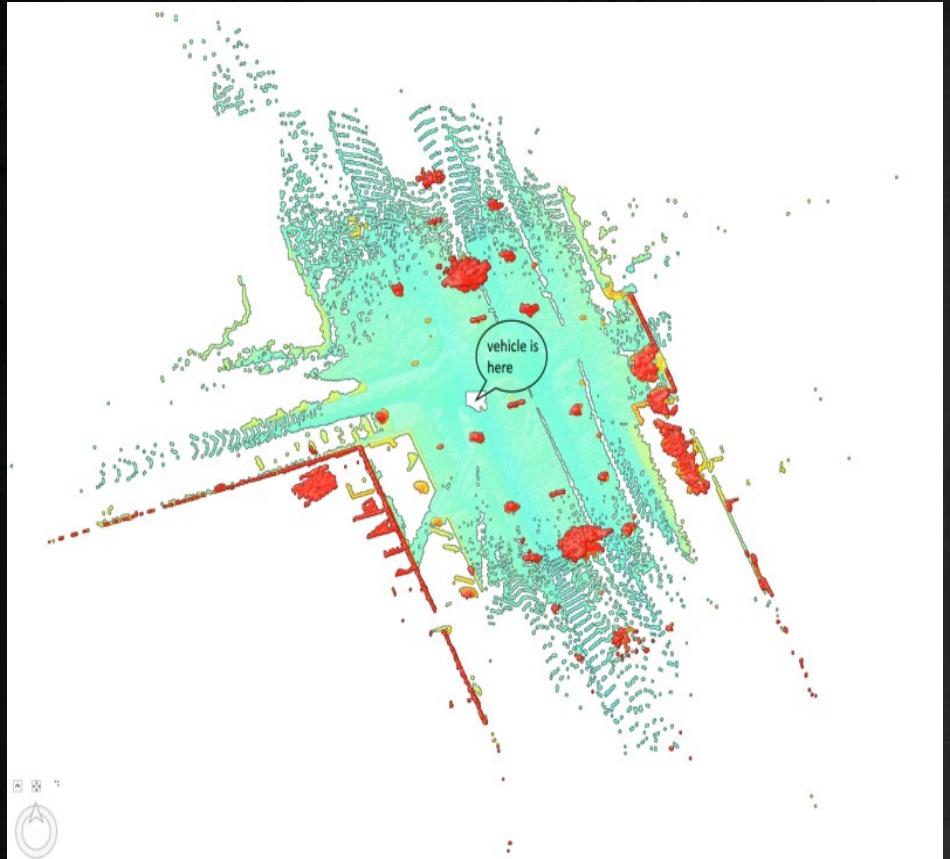


Classifying by
projecting 2D
semantic masks
onto the points



PoC: PointCNN in Mobile Point Clouds

Training on individual clouds or single consolidated cloud?



LAS Dataset Layer

CycloMedia-sample - Scene-RGB - ArcGIS Pro

Project Map Insert Analysis View Edit Imagery Share Appearance Data Classification

Dmitry (Конференция Esri в России и странах СНГ)

Scene-RGB

Scene-Class

41 ft | Selected Features: 0 | 41 ft | Selected Features: 0 | 41 ft | Selected Features: 0 |

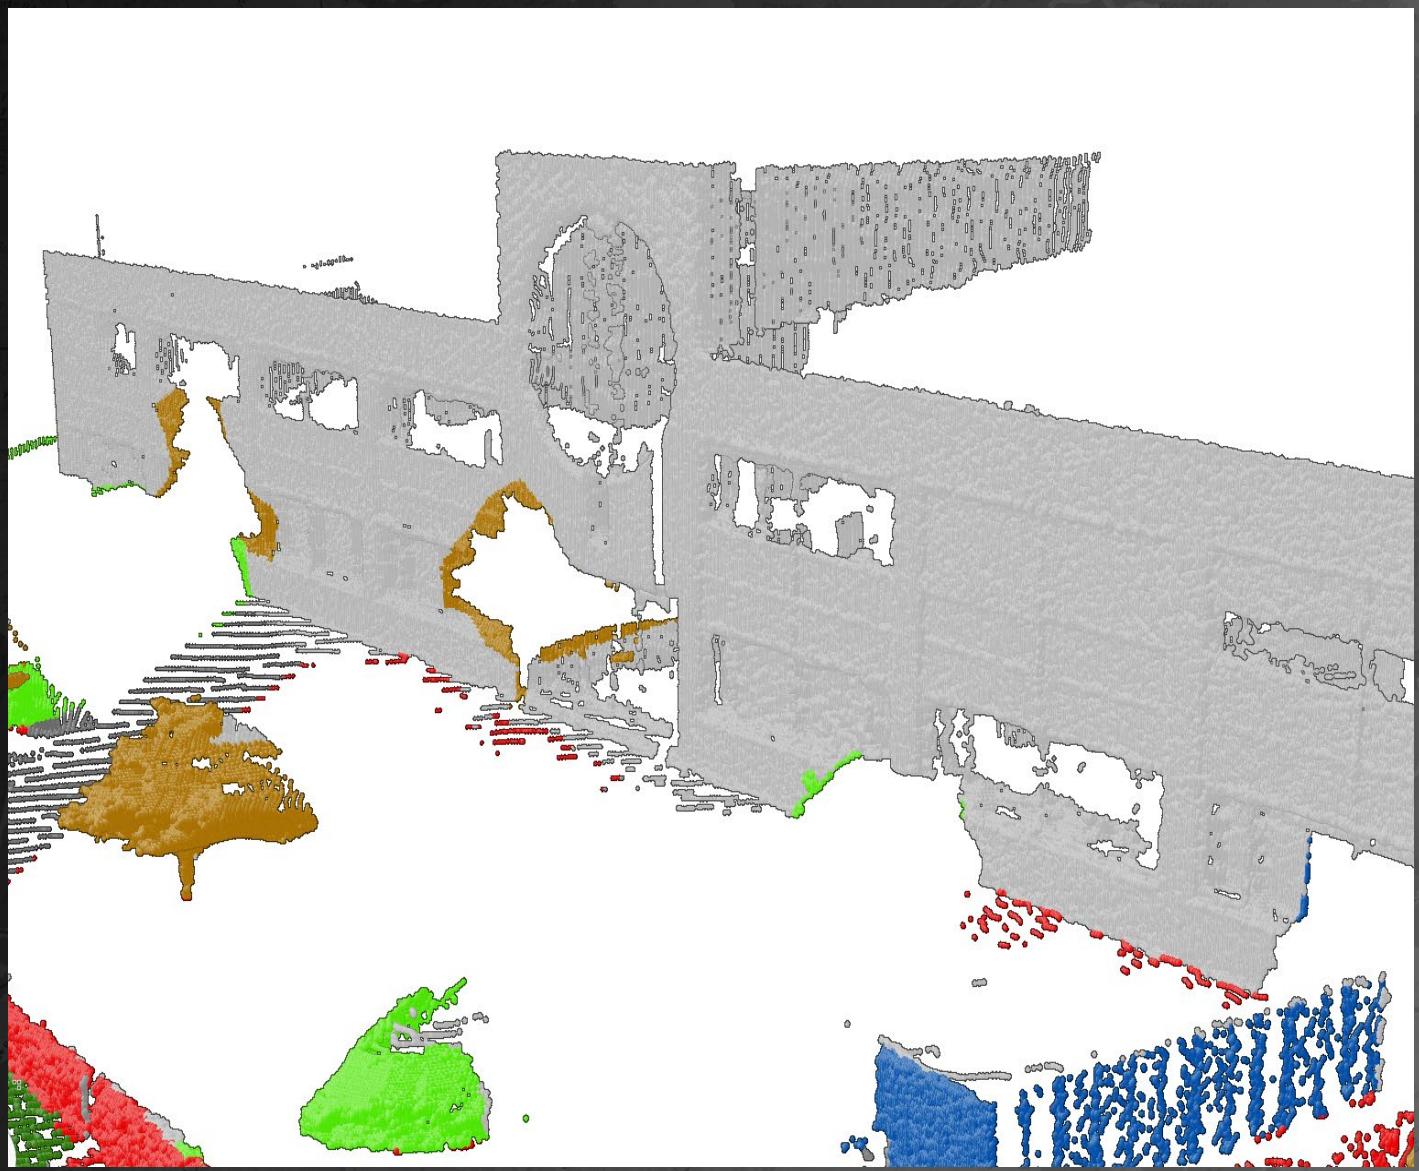
4.4013090°E 51.9145435°N 28.203 ft | 4.4013090°E 51.9145435°N 28.203 ft | 4.4013090°E 51.9145435°N 28.203 ft |

Contents Tasks Catalog Pop-up Symbology Geoprocessing Bookmarks

PoC: PointCNN in Mobile Point Clouds

Buildings are well represented
in the training set:

POINTS_COUNT	CLASS_NAME
539,065,936	road
361,835,379	building
284,781,152	vegetation_high
99,861,937	sidewalk
70,959,436	car
70,826,139	grass
55,858,858	parking
47,782,986	vegetation_low
43,682,759	barrier_wall
38,715,988	bike_lane
33,589,152	curb
32,879,547	barrier_separator
17,685,268	structure_bridge

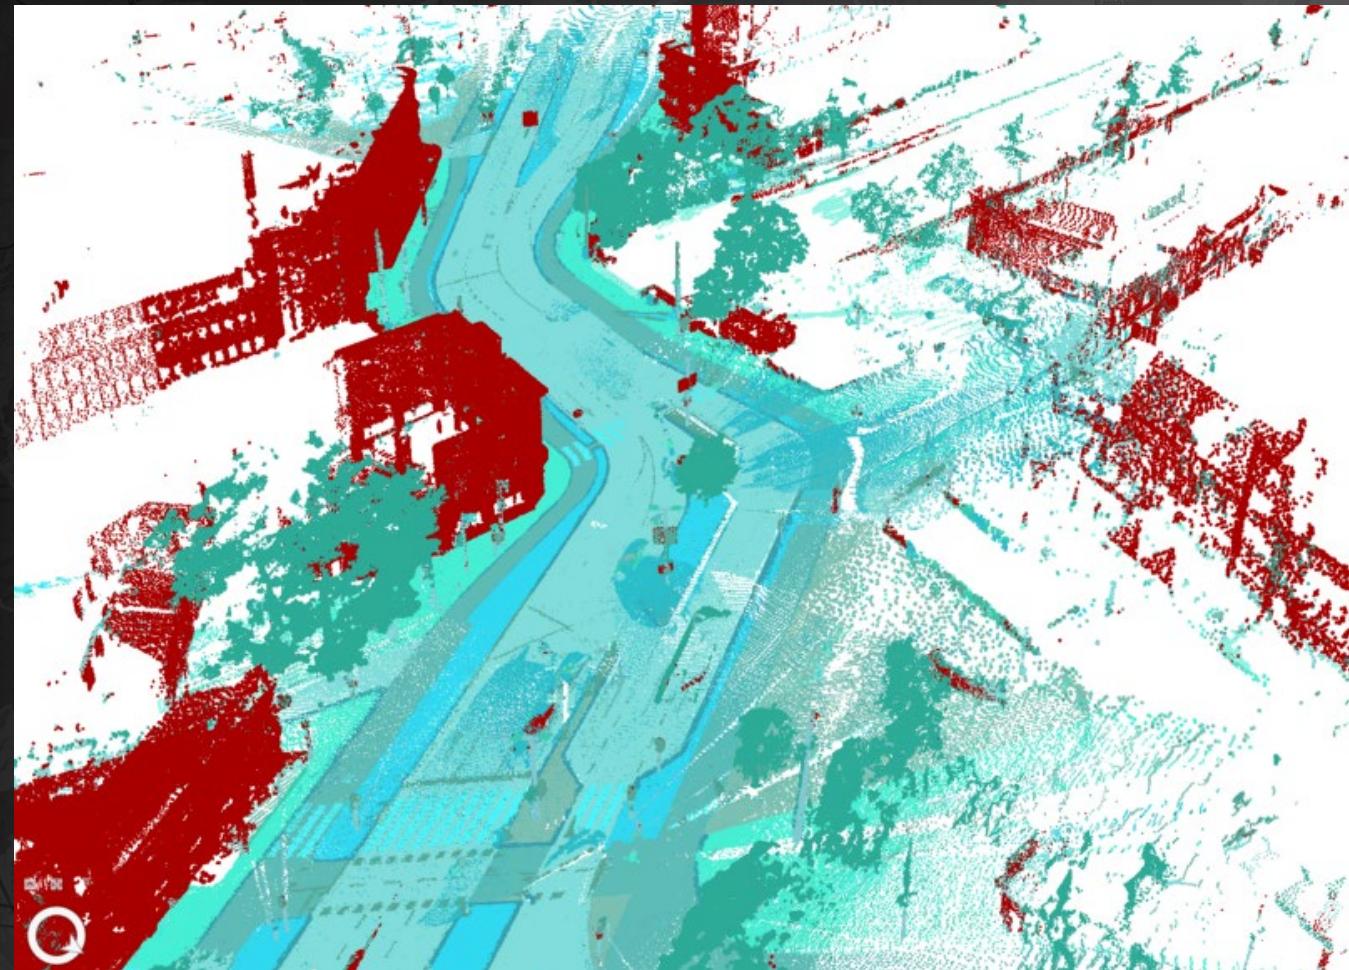


PoC: PointCNN in Mobile Point Clouds

After ~36 hours of training on
NVIDIA Quadro GV100:

@E150K

	OTHER	BUILDING
Precision:	[0.99126286	0.96047395]
Recall:	[0.9923109	0.95528204]
F1-score:	[0.9917866	0.95787096]



Project Map Insert Analysis View Edit Imagery Share Point Cloud Layer CycloMedia-sample - Scene-Class - ArcGIS Pro ? - × Dmitry (Конференция Esri в России и странах СНГ) Contents Tasks Catalog Pop-up Symbology Geoprocessing Bookmarks

Scene-RGB X

Scene-Class X

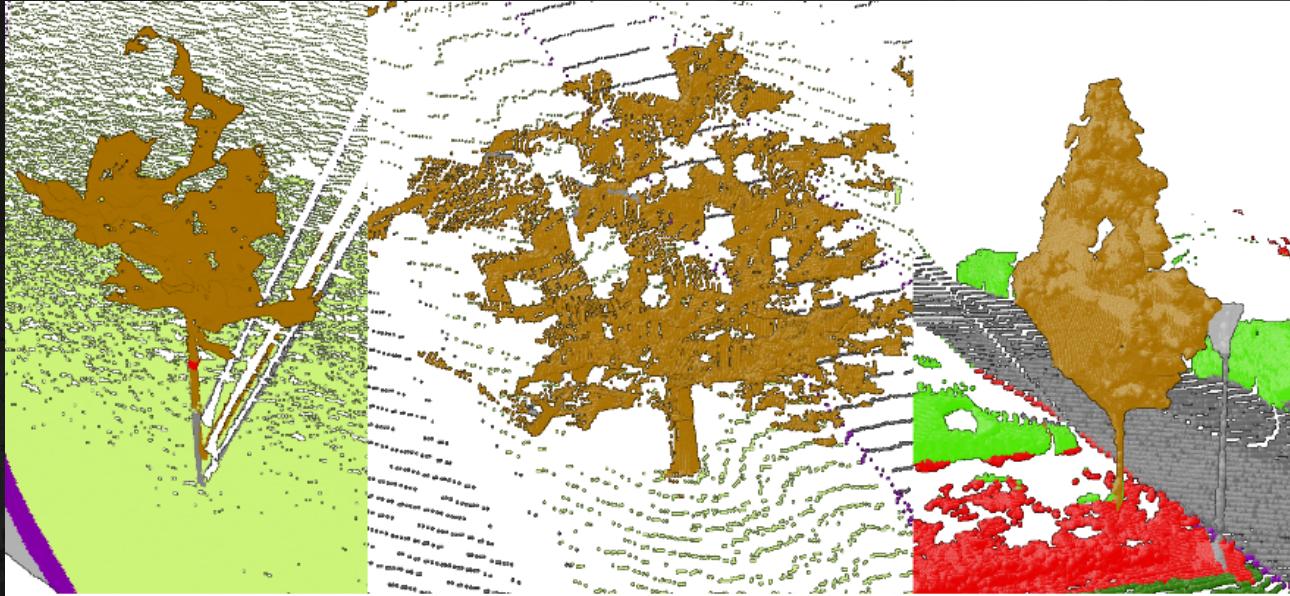
PointCNN:
Trees

74 ft | Selected Features: 0 | 4.3916330°E 51.9112411°N | 2.453 ft | Selected Features: 0 |

PoC: PointCNN in Mobile Point Clouds

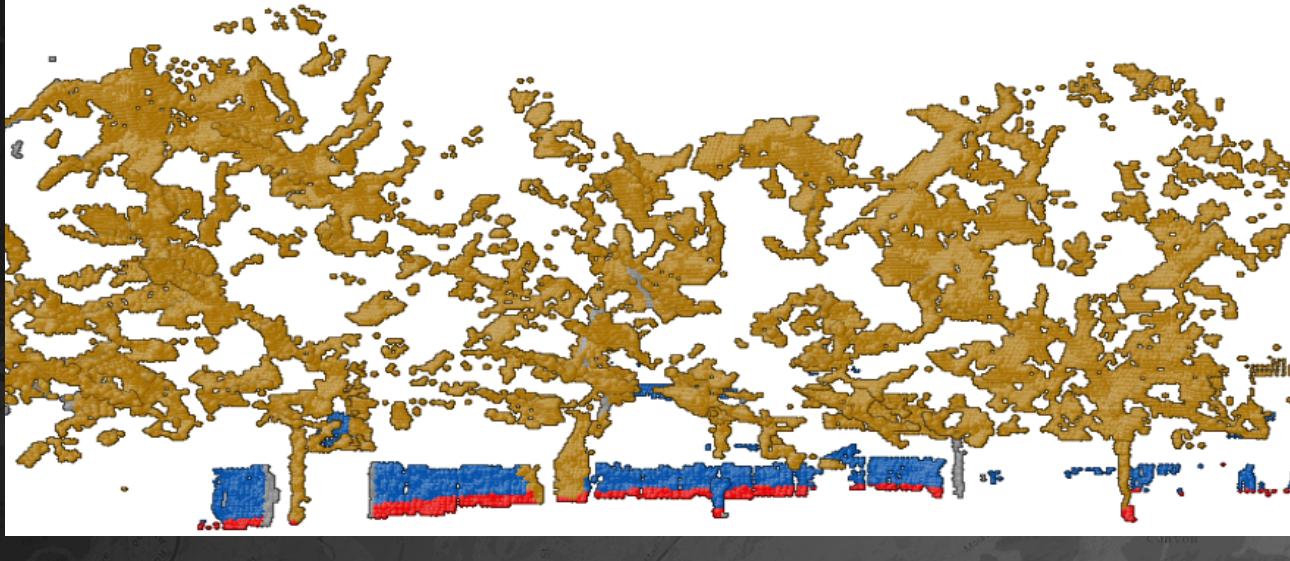
Trees are too well represented:

POINTS_COUNT	CLASS_NAME
539,065,936	road
361,835,379	building
284,781,152	vegetation_high
99,861,937	sidewalk
70,959,436	car
70,826,139	grass
55,858,858	parking
47,782,986	vegetation_low
43,682,759	barrier_wall
38,715,988	bike_lane
33,589,152	curb
32,879,547	barrier_separator
17,685,268	structure_bridge



But harder to learn than buildings:

- Species
- Age
- Density
- RGB ranges
- Canopy type

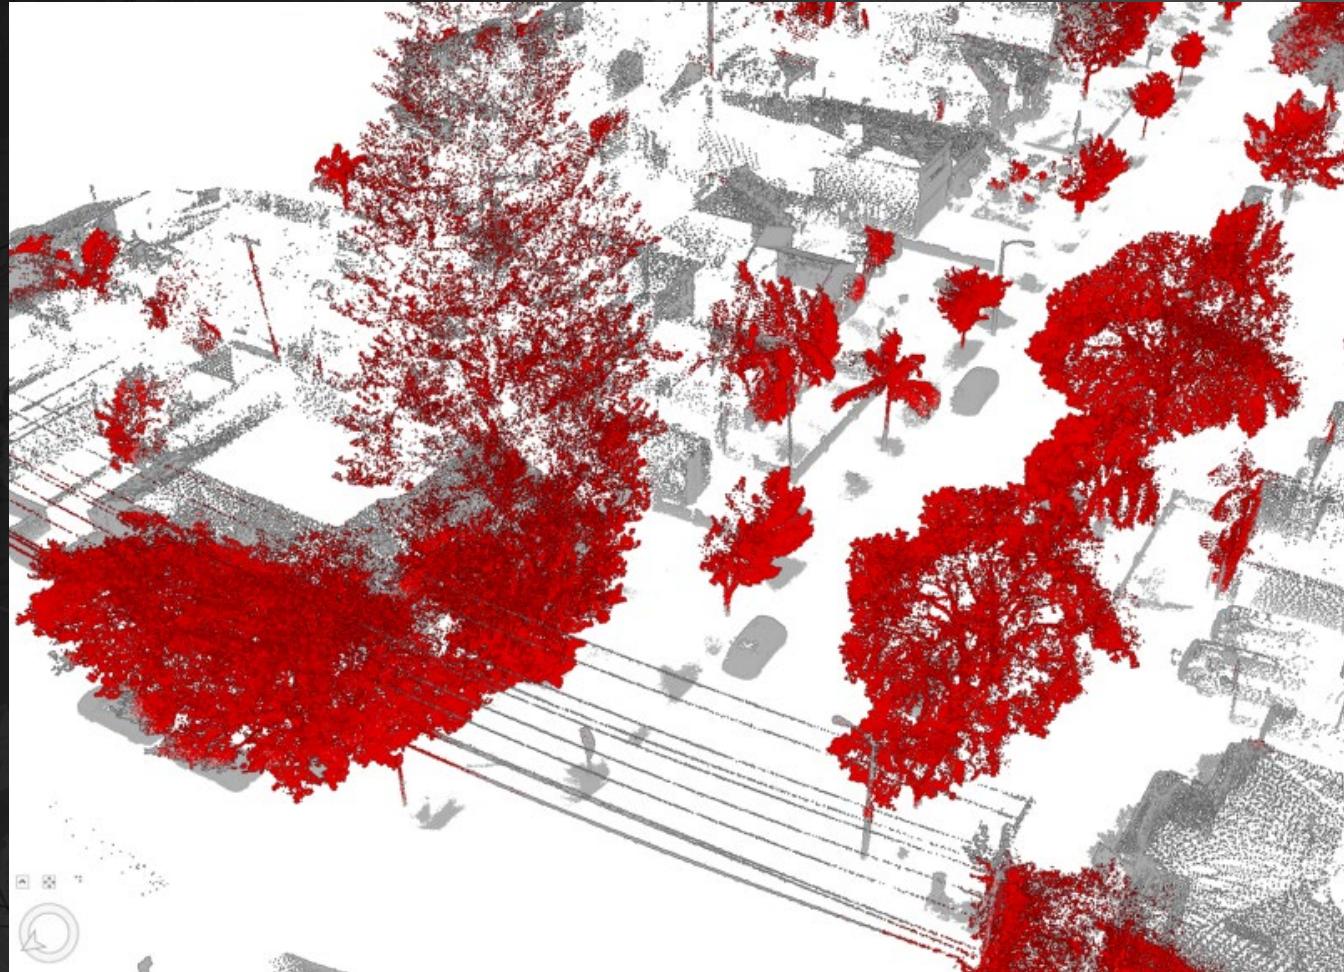


PoC: PointCNN in Mobile Point Clouds

After ~24 hours of training on GV100:

@E100K

	OTHER	TREE
Precision:	[0.98007521	0.94283636],
Recall:	[0.98718658	0.91327329],
F1-score:	[0.98361804	0.92781940]





PoC: PointCNN in Mobile Point Clouds

- Traffic Lights are harder:
 - Small number of samples
~580,000 points (0.032%)
 - Different types, attachment options
 - Located at most noisy intersections



PoC: PointCNN in Mobile Point Clouds

- Traffic Lights are harder:
 - Small number of samples
~580,000 points (0.032%)
 - Different types, attachment options
 - Located at most noisy intersections



PoC: PointCNN in Mobile Point Clouds

After ~72 hours of training on GV100:

@E300K

	OTHER	TRAFFIC LIGHT
Precision:	[0.99974482	0.51987179]
Recall:	[0.99939181	0.72079467]
F1-score:	[0.99956828	0.60406375]

Extracted Objects, after DBScan:

DBSCAN 30pts/0.5m

Precision = .655

Recall = .704

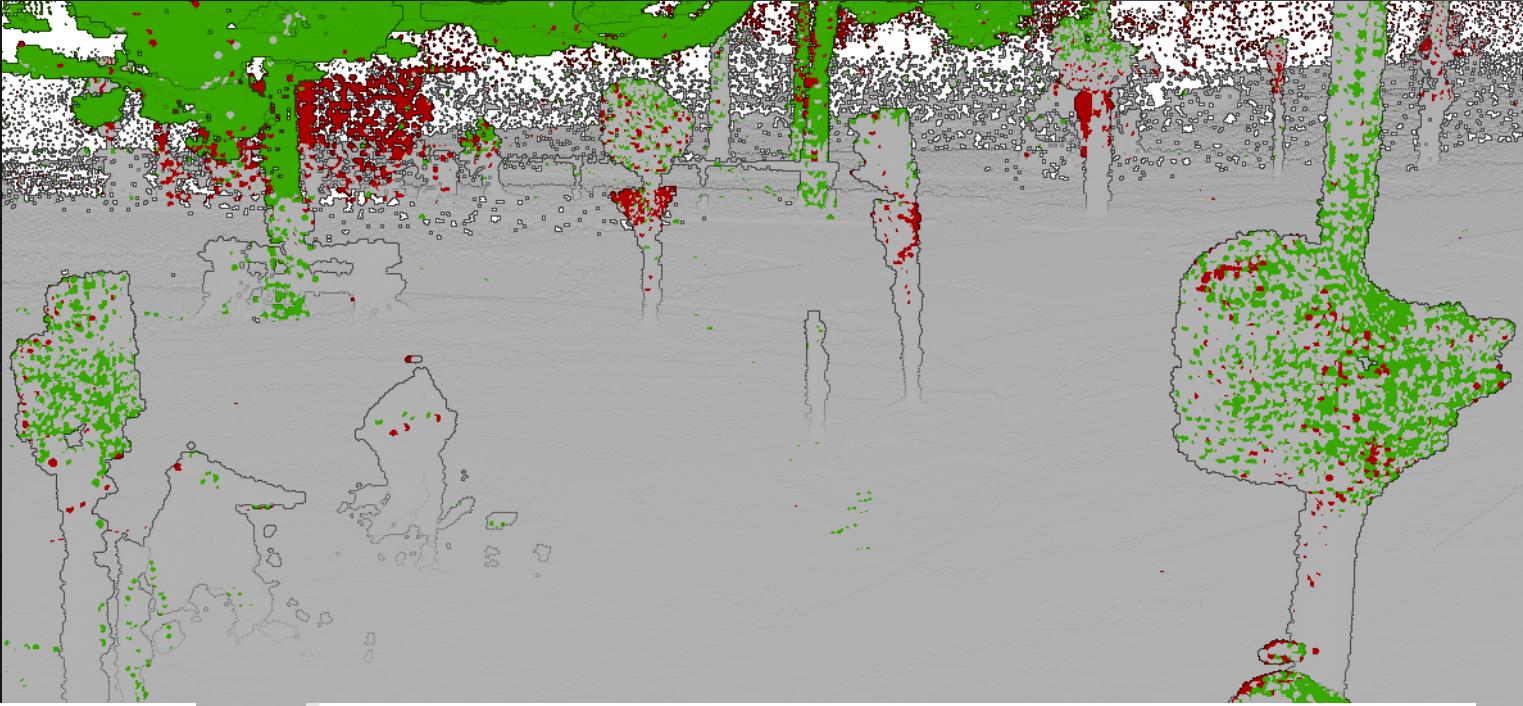
F1-Score = .679



PoC: PointCNN in Mobile Point Clouds

Denoising trick with PointCNN:

- If the **valuable semantics** is consistent, but **noise is random**, a trained model can be used as a "low-pass noise filter"



```
Rijnstraat
DBScan Settings: >=3 pts per cluster, .5m max between points
GroundTruth = 31 traffic lights
TruePositives = 21
FalsePositives = 31
FalseNegatives = 10
Precision = .404
Recall = .677
F1-Score = .506
```

Original Training set

```
Rijnstraat
DBSCAN Settings: 30pts per cluster / 0.5m max between points
TP = 19
FP = 10
FN = 12
Precision = .655
Recall = .704
F1-Score = .679
```

Training set with Buildings &
Trees relabeled by PointCNN
models

Project Map Insert Analysis View Edit Imagery Share Point Cloud Layer CycloMedia-sample - Scene-Class - ArcGIS Pro Dmitry (Конференция Esri в России и странах СНГ) ? Catalog Pop-up Symbology Geoprocessing Bookmarks

Contents Tasks

Scene-RGB X

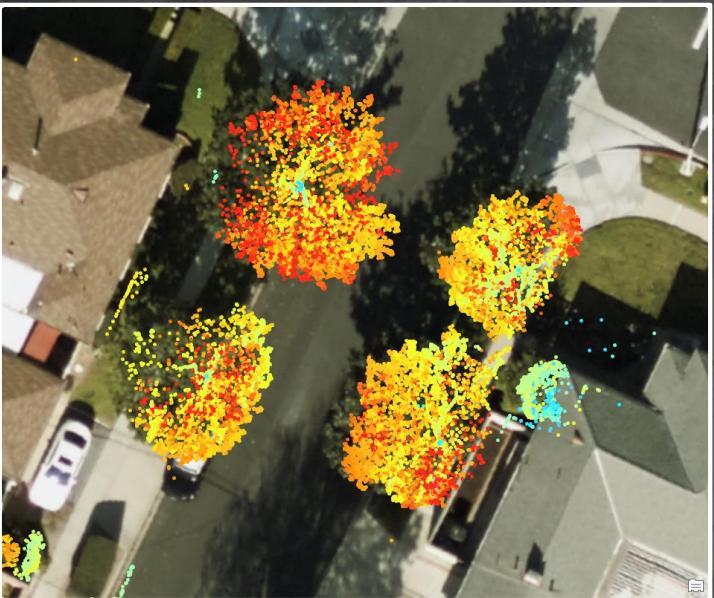
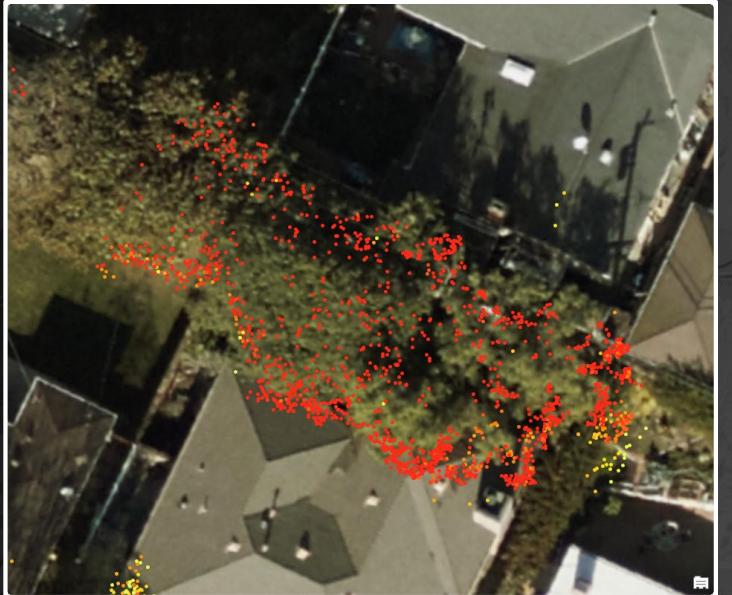
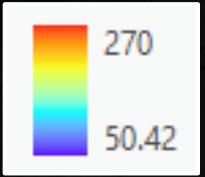
Scene-Class X

Feature Extraction: Trees

74 ft | Selected Features: 0 | 4.3916330°E 51.9112411°N | 2.453 ft | Selected Features: 0 |

Los Angeles data - Trees

- Number of LAS files: 24
- Number of LAS points: 33.5 million
- XYZ units: US foot
- Average point spacing: 0.11 feet
- Classes
 - 0 = Other
 - 1 = Trees



Feature extraction pipelines (#1 - Raster)



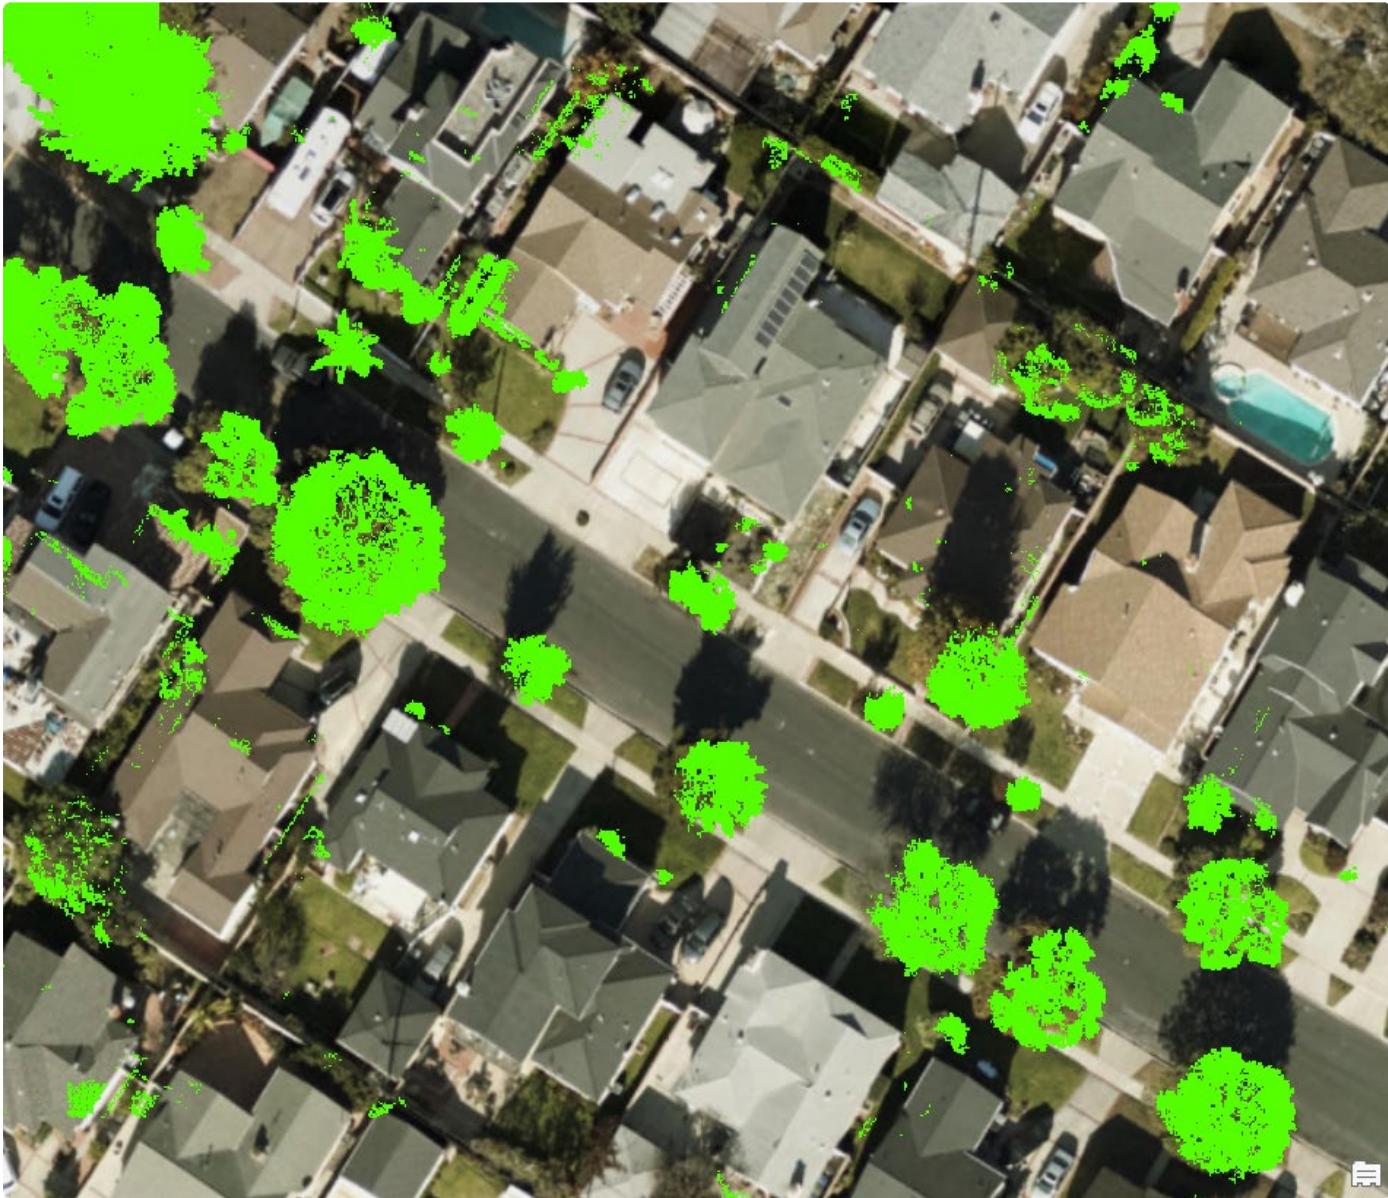
Feature extraction pipelines (#1 - Raster)

Suburban
Los Angeles



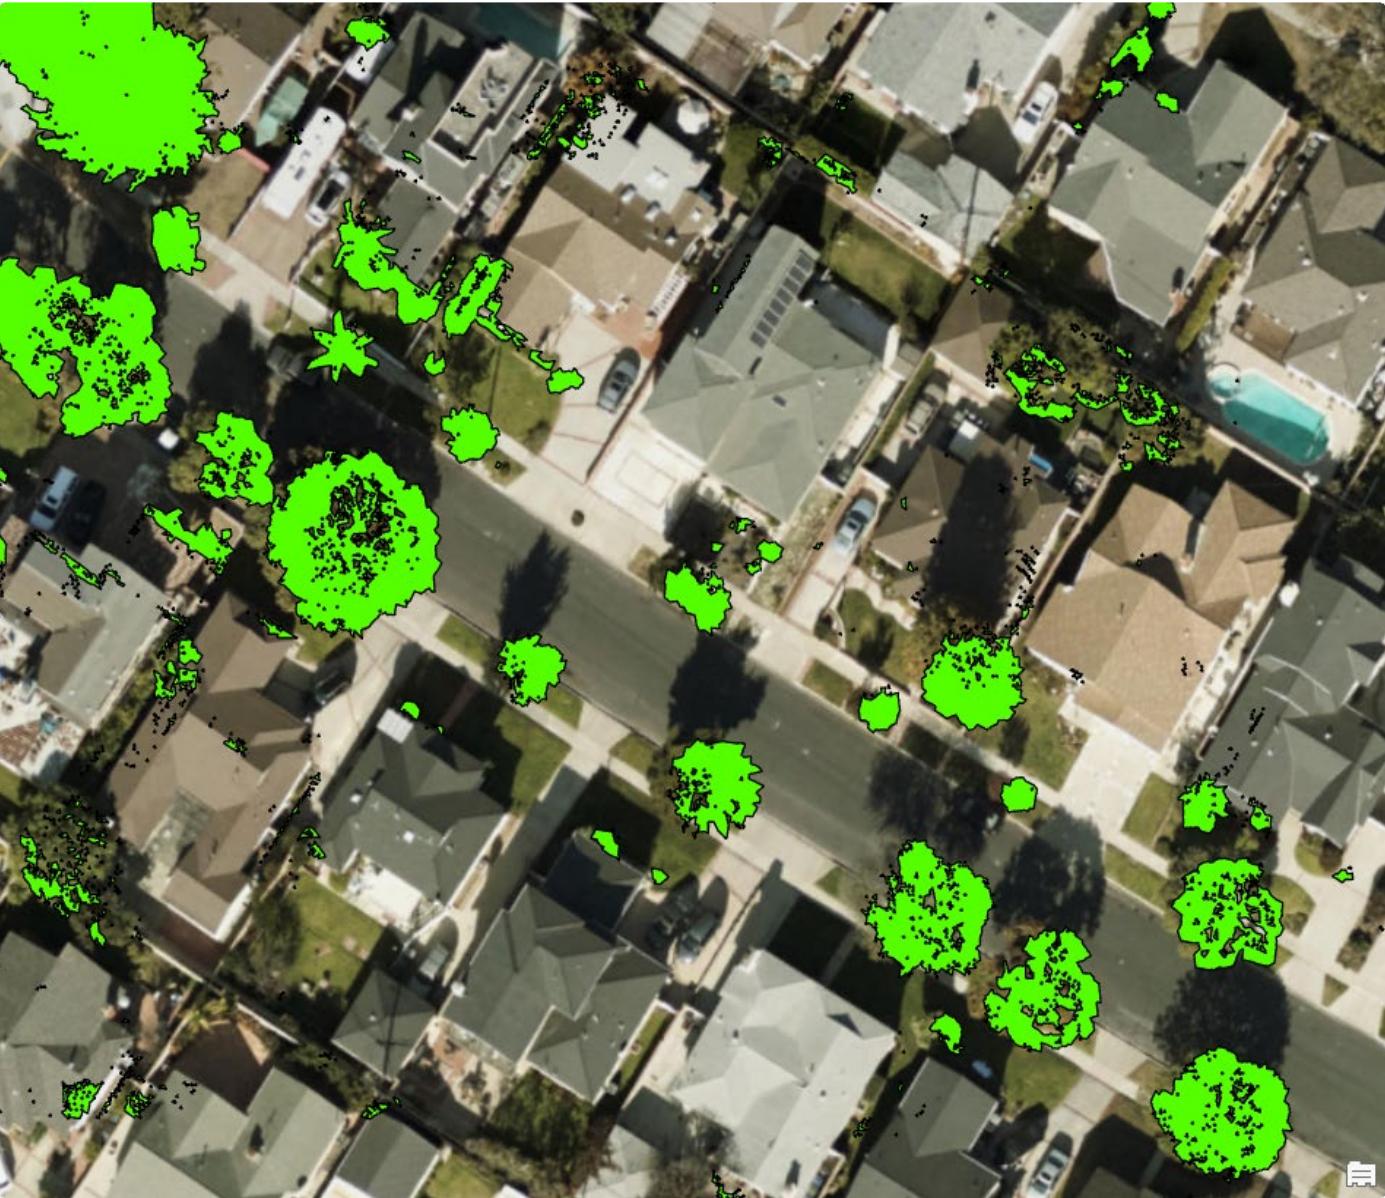
Feature extraction pipelines (#1 - Raster)

LiDAR to
raster



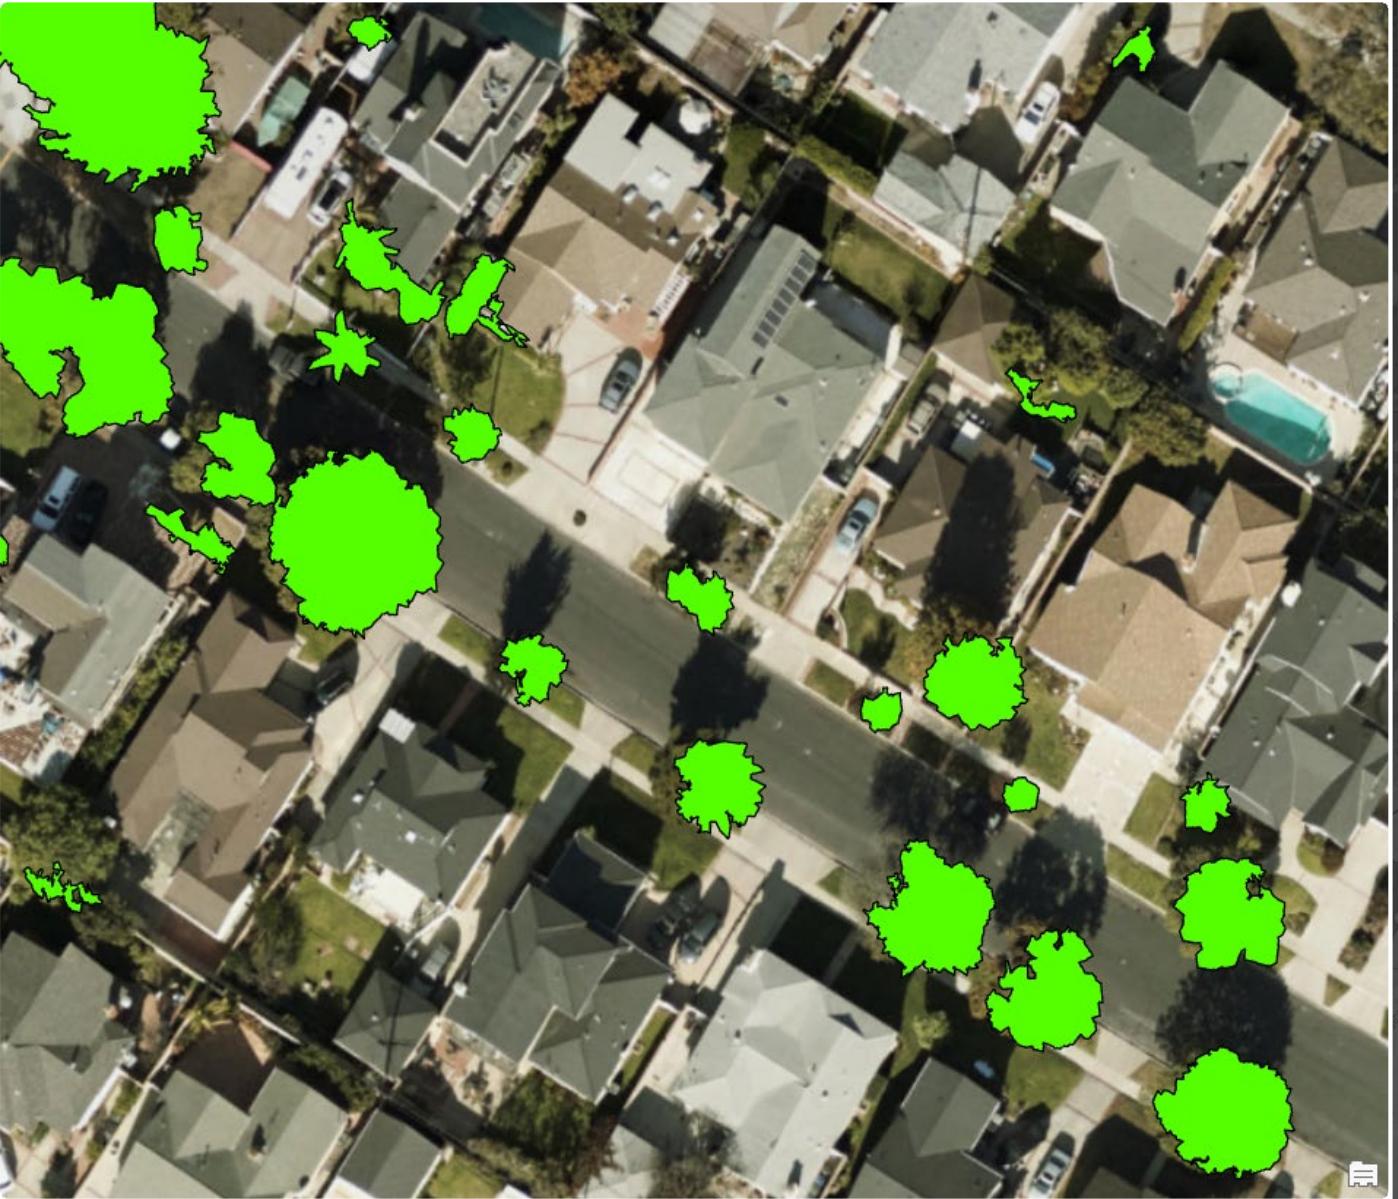
Feature extraction pipelines (#1 - Raster)

Raster to
polygons



Feature extraction pipelines (#1 - Raster)

Remove small polygons, fill holes



Feature extraction pipelines (#1 - Raster)

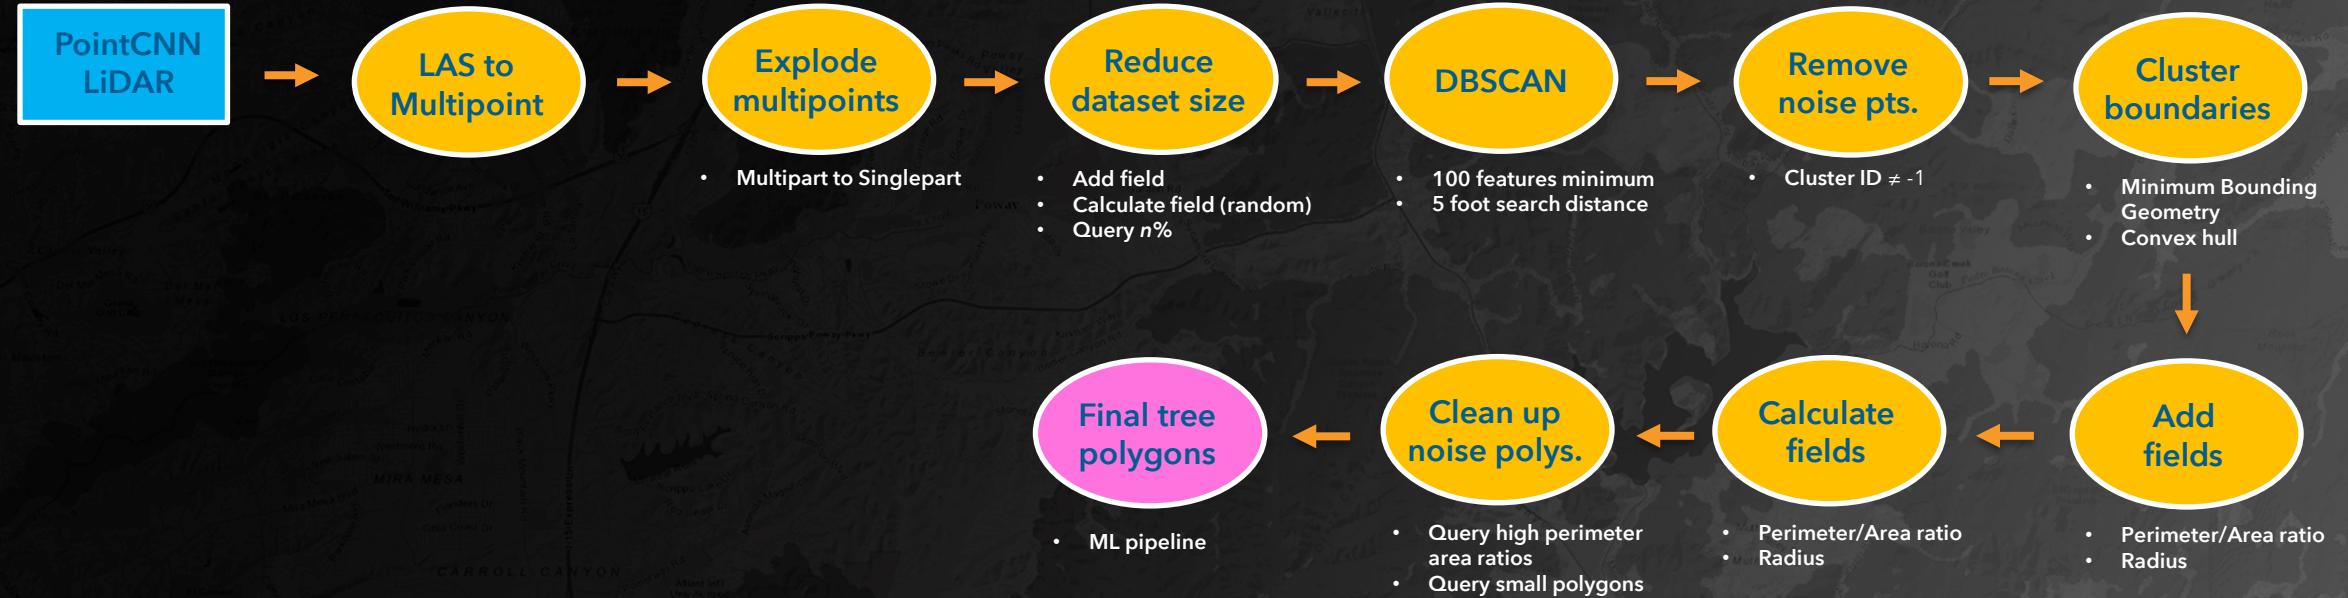


Clean up

- Remove high perimeter/area ratios

Feature to point

Feature extraction pipelines (#2 - ML)



Tasks

Pipeline #2 - Machine Learning

Tasks Messages

- Convert LAS files to vector points
- Reduce the size of the point feature class (Optional)
- Create point clusters
- Add and calculate fields
- Remove noise polygons
- Export final tree polygons to new feature class

Feature extraction pipelines (#2 - ML)

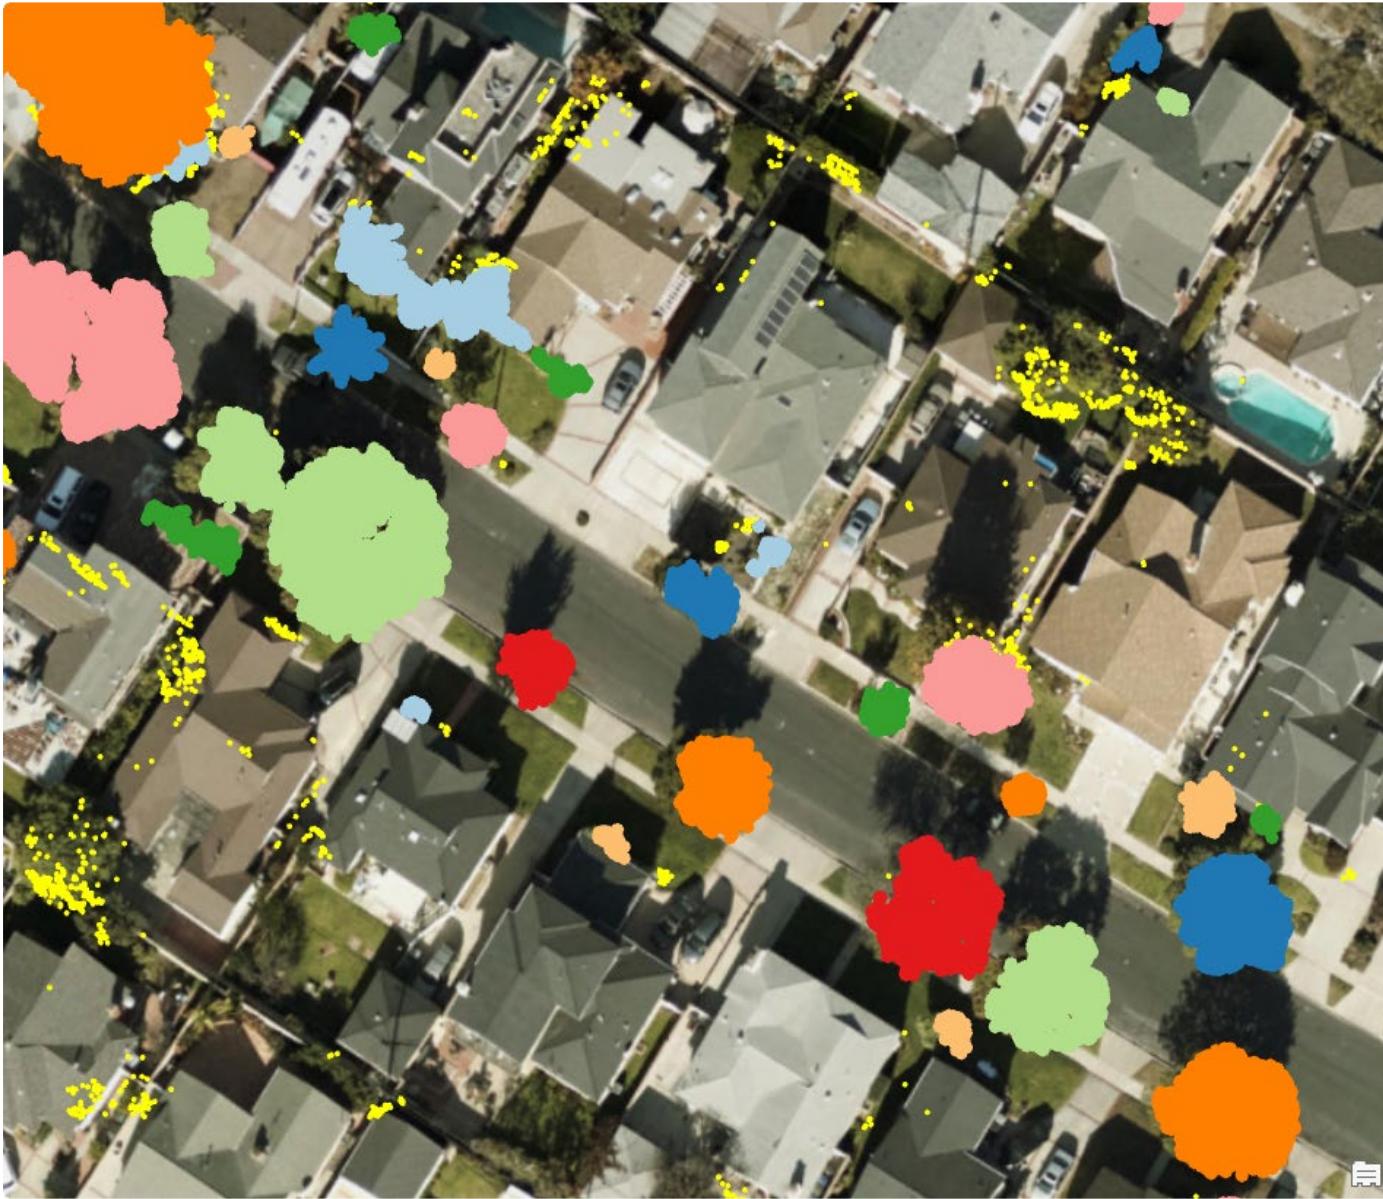
LiDAR
to Multipoints



Feature extraction pipelines (#2 - ML)

DBSCAN

- Min features: 100
- Search distance: 5 ft



Feature extraction pipelines (#2 - ML)

Minimum
bounding
geometry



Feature extraction pipelines (#2 - ML)

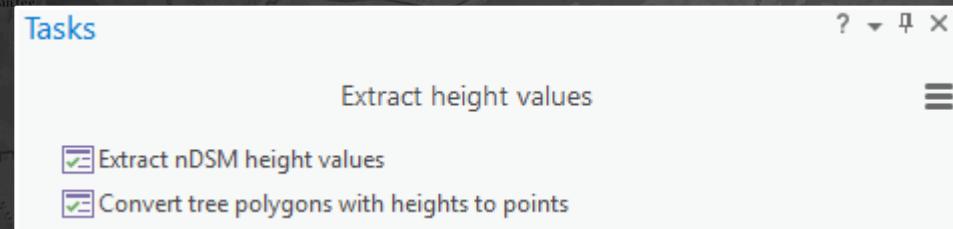
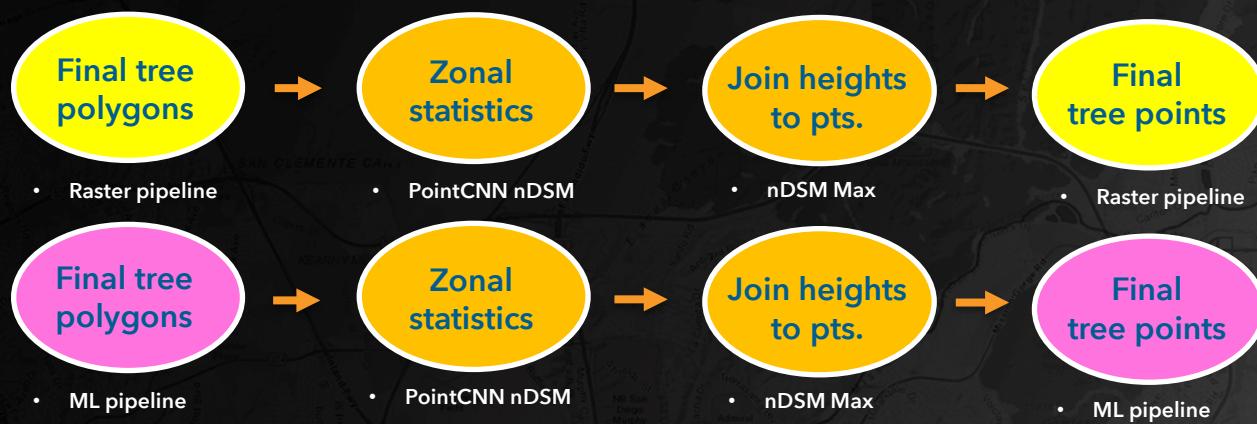
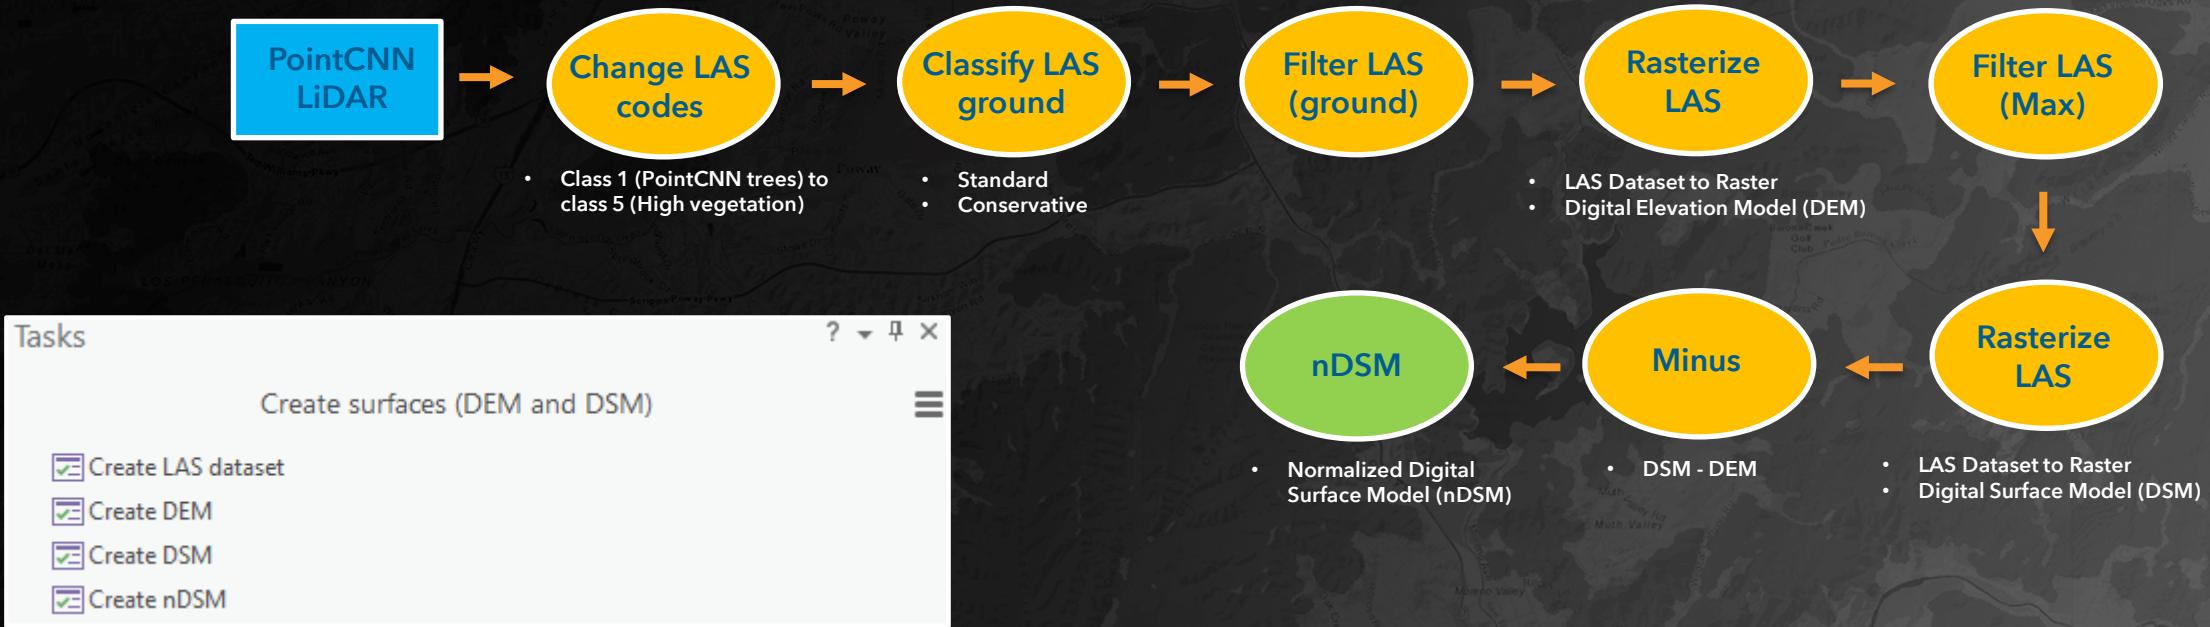
Clean up

- Remove high perimeter/area ratios
- Remove small polygons

Feature to point



Feature extraction pipelines (Creating heights)

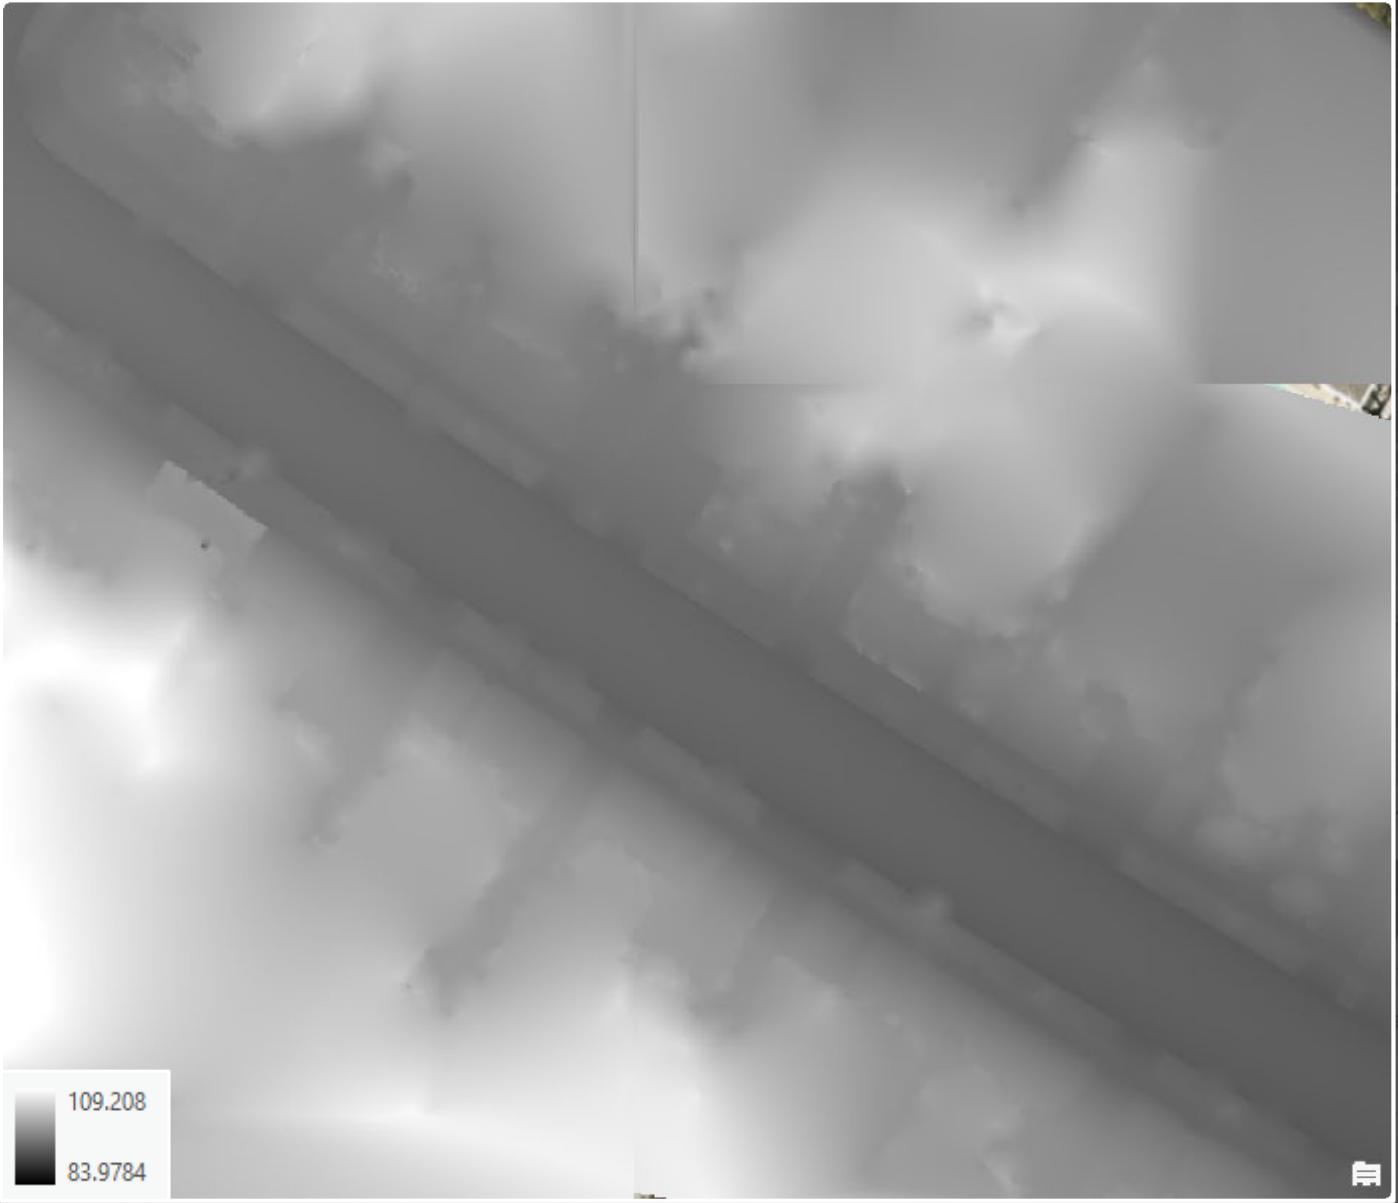


Feature extraction pipelines (Surfaces)



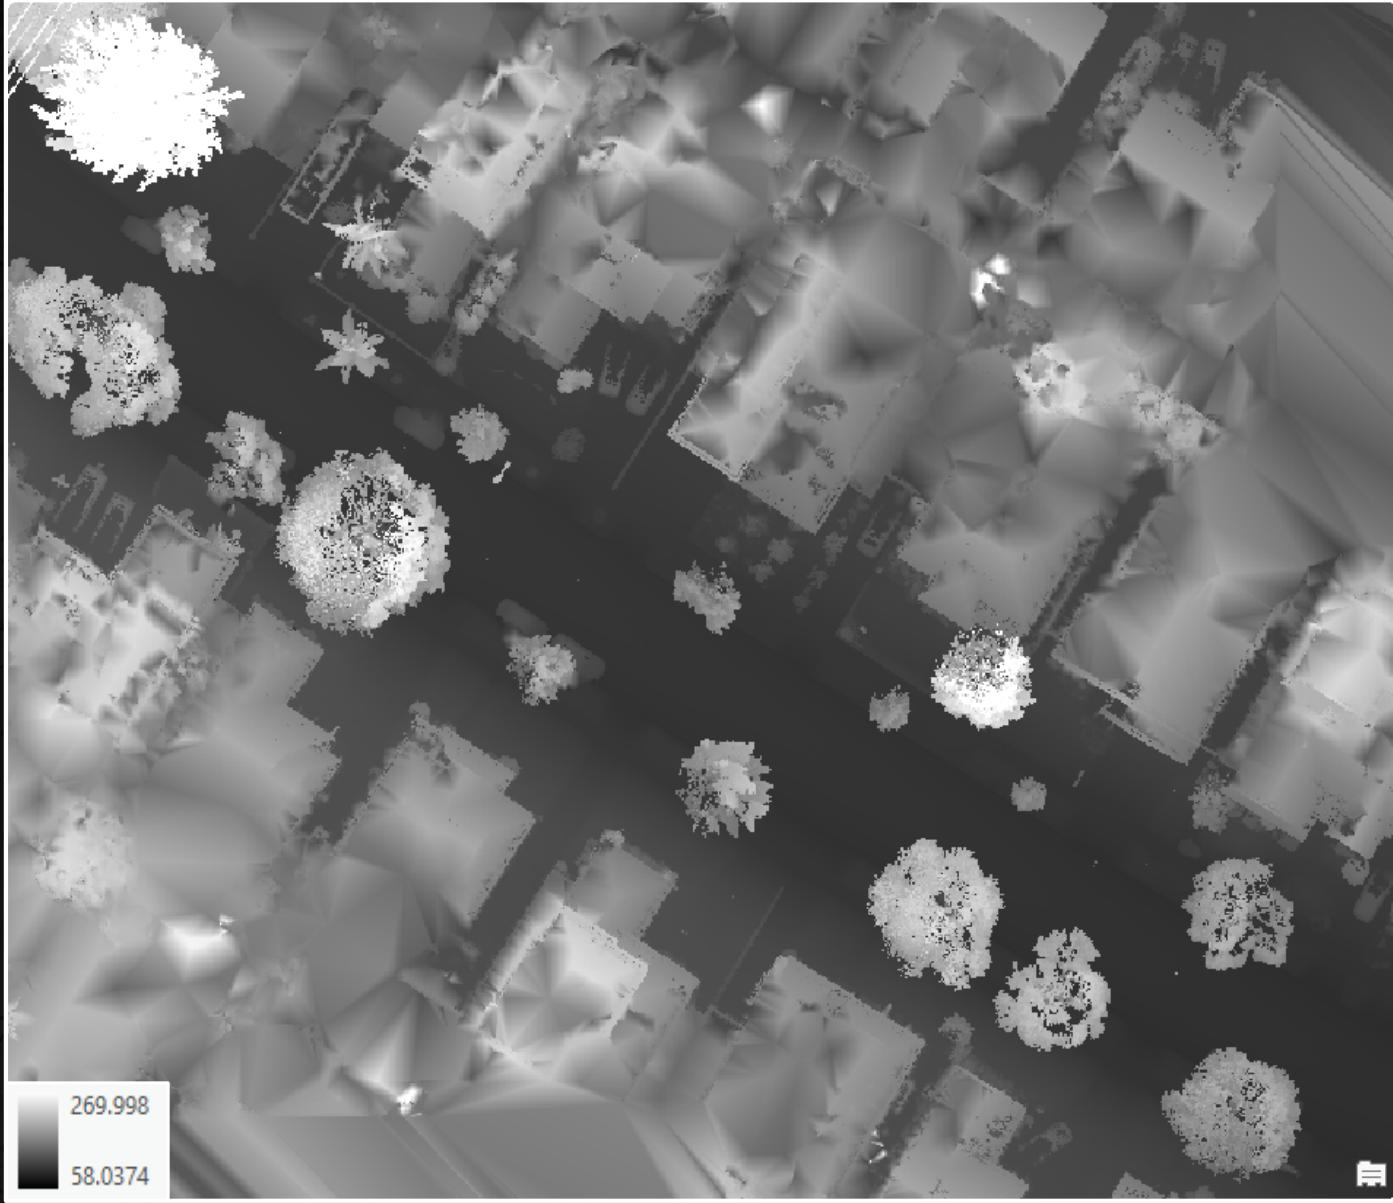
Feature extraction pipelines (Surfaces)

Digital Elevation
Model (DEM)



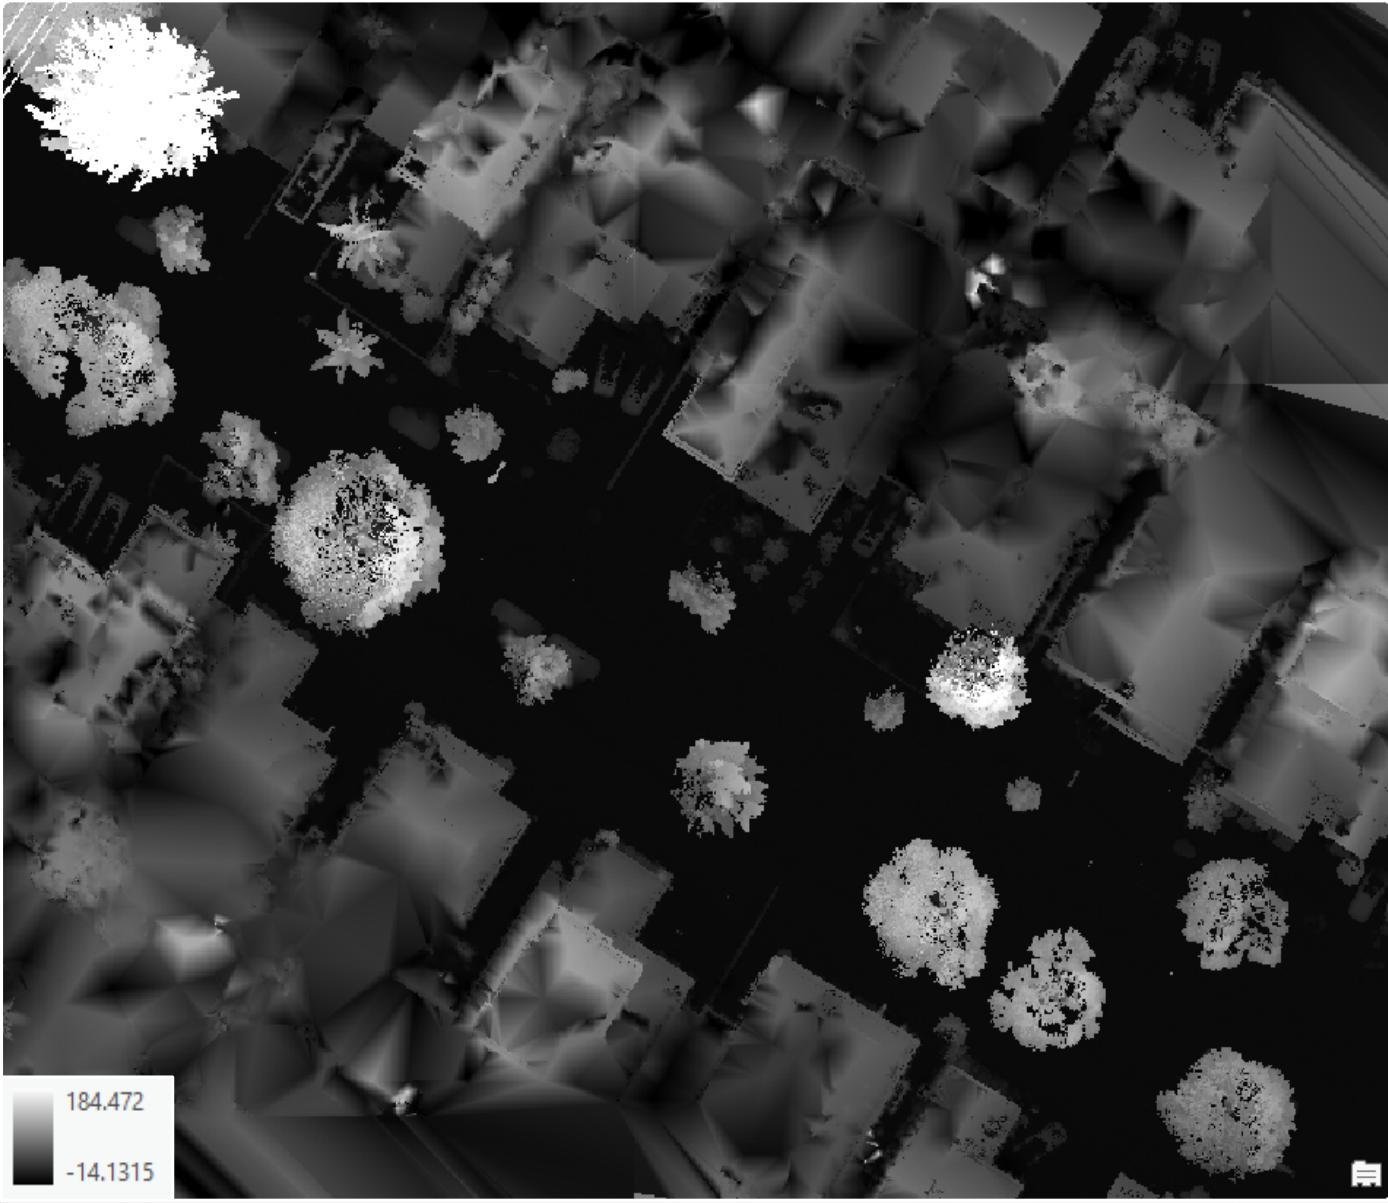
Feature extraction pipelines (Surfaces)

Digital Surface
Model (DSM)

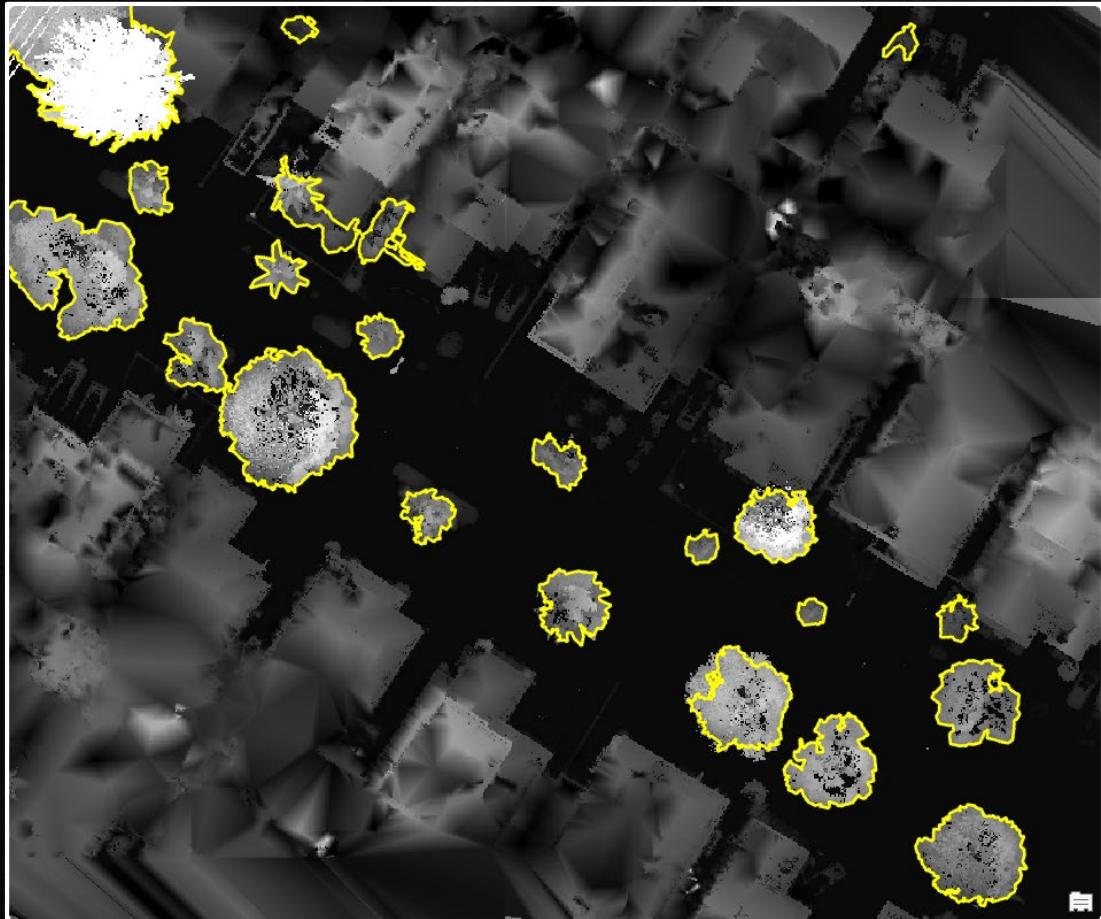


Feature extraction pipelines (Surfaces)

Normalized
Digital Surface
Model (nDSM)

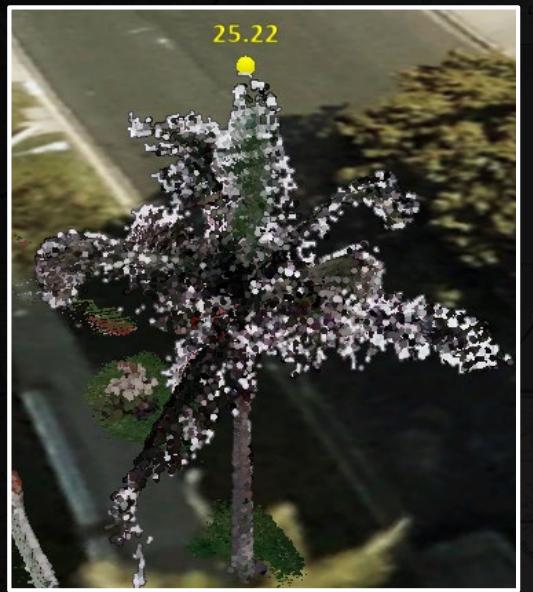


Feature extraction pipelines (Extracting heights)



Field:	Add	Calculate	Selection:	Zoom To	Switch	Clear	Delete	Copy	MIN	MAX	MEAN
OBJECTID	Shape	Id	gridcode	ORIG_FID	PerimArea	Radius	Shape_Length	Shape_Area	MIN	MAX	MEAN
37	Polygon	2661	5	2661	0.410733	03.515994	650.41017	1583.536807	4.462326	154.2267	27.415965
9	Polygon	556	5	556	0.138295	86.932646	546.213924	3949.637901	0.013176	109.498703	39.667253
19	Polygon	1082	5	1082	0.693453	12.51784	78.65191	113.420674	0.003136	76.517021	66.666903
73	Polygon	5680	5	5680	0.747342	13.868811	87.140311	116.600382	-0.350594	47.808449	30.809626
61	Polygon	4796	5	4796	0.619071	63.564123	399.385163	645.13616	-0.002556	46.717812	30.278498
48	Polygon	4111	5	4111	0.288888	27.710872	802.431075	2777.651806	-0.396576	43.958366	16.804566
26	Polygon	1681	5	1681	0.251077	07.481045	675.323325	2689.701185	-0.33046	43.733139	21.36026
2	Polygon	115	5	115	0.229437	34.146386	214.548068	935.108598	-0.010002	43.346733	20.196509
69	Polygon	5540	5	5540	0.377664	49.883333	313.426225	829.906863	-0.047104	42.973679	26.045426
27	Polygon	1685	5	1685	0.277736	21.643755	135.991725	489.6445	-0.196877	42.370506	25.789876
23	Polygon	1505	5	1505	0.385087	45.064626	283.149397	735.287419	0.002182	40.509964	27.601931
68	Polygon	5159	5	5159	0.402617	73.061744	459.060475	1140.190739	-0.042221	39.826958	21.757074
43	Polygon	3526	5	3526	0.40172	64.971799	408.22985	1016.205348	-0.011345	38.565872	17.673984
4	Polygon	209	5	209	0.60037	48.334549	303.694931	505.846151	-1.052994	37.414551	17.958374

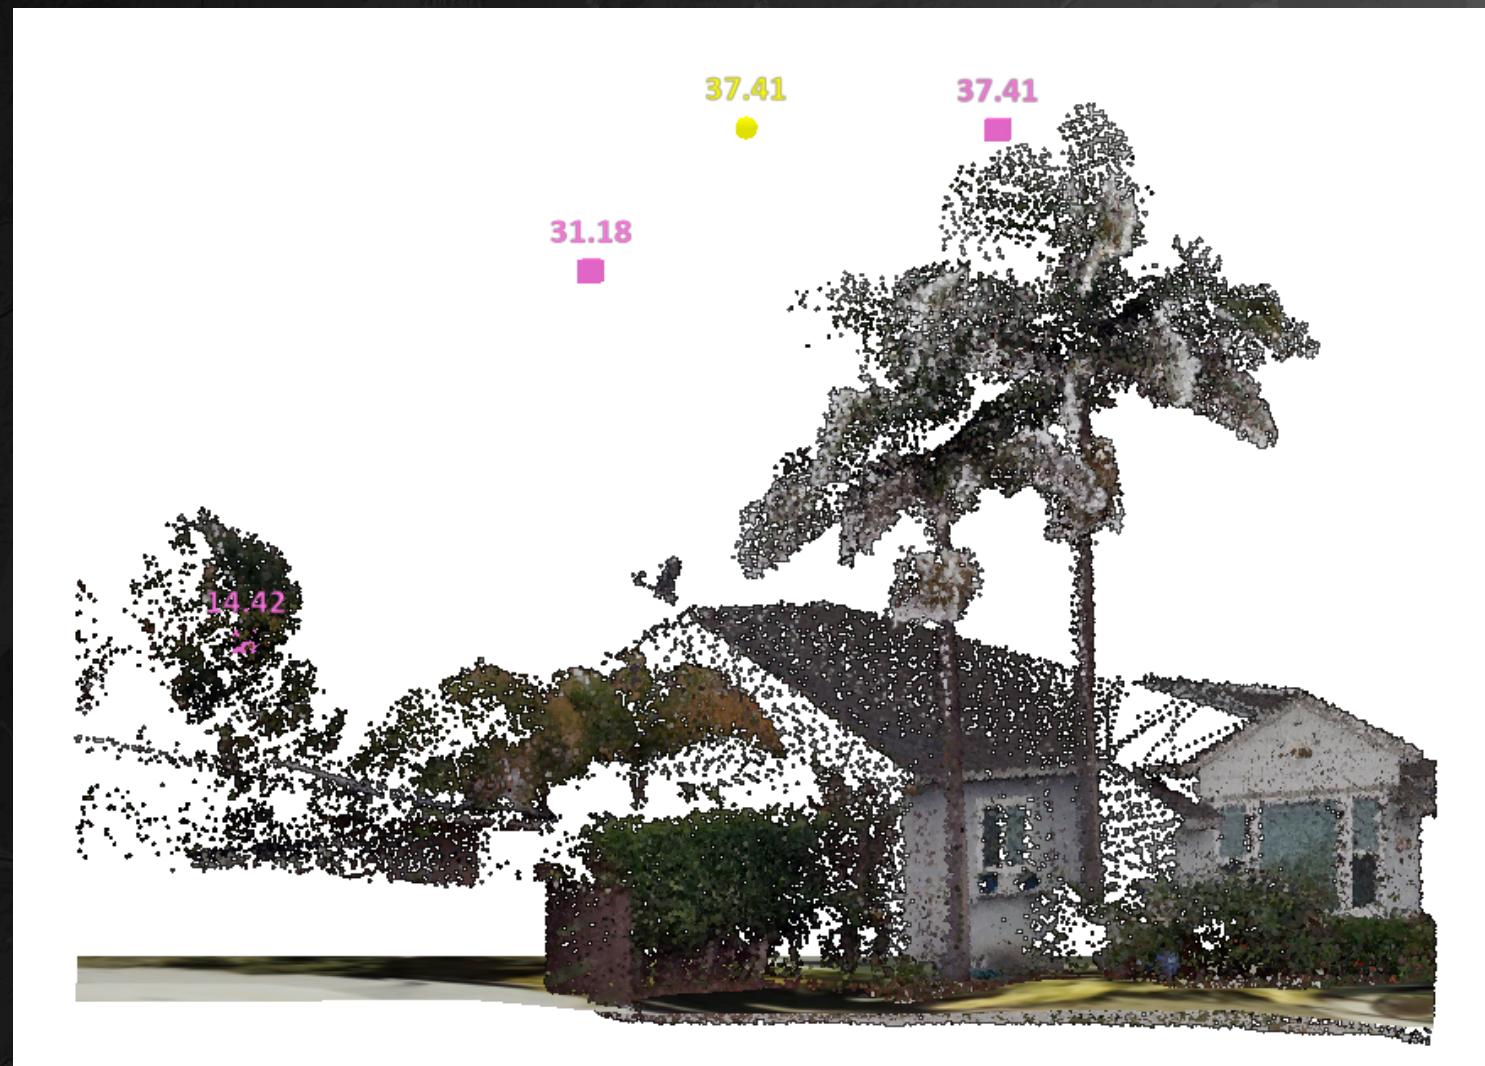
Results (2D and 3D)



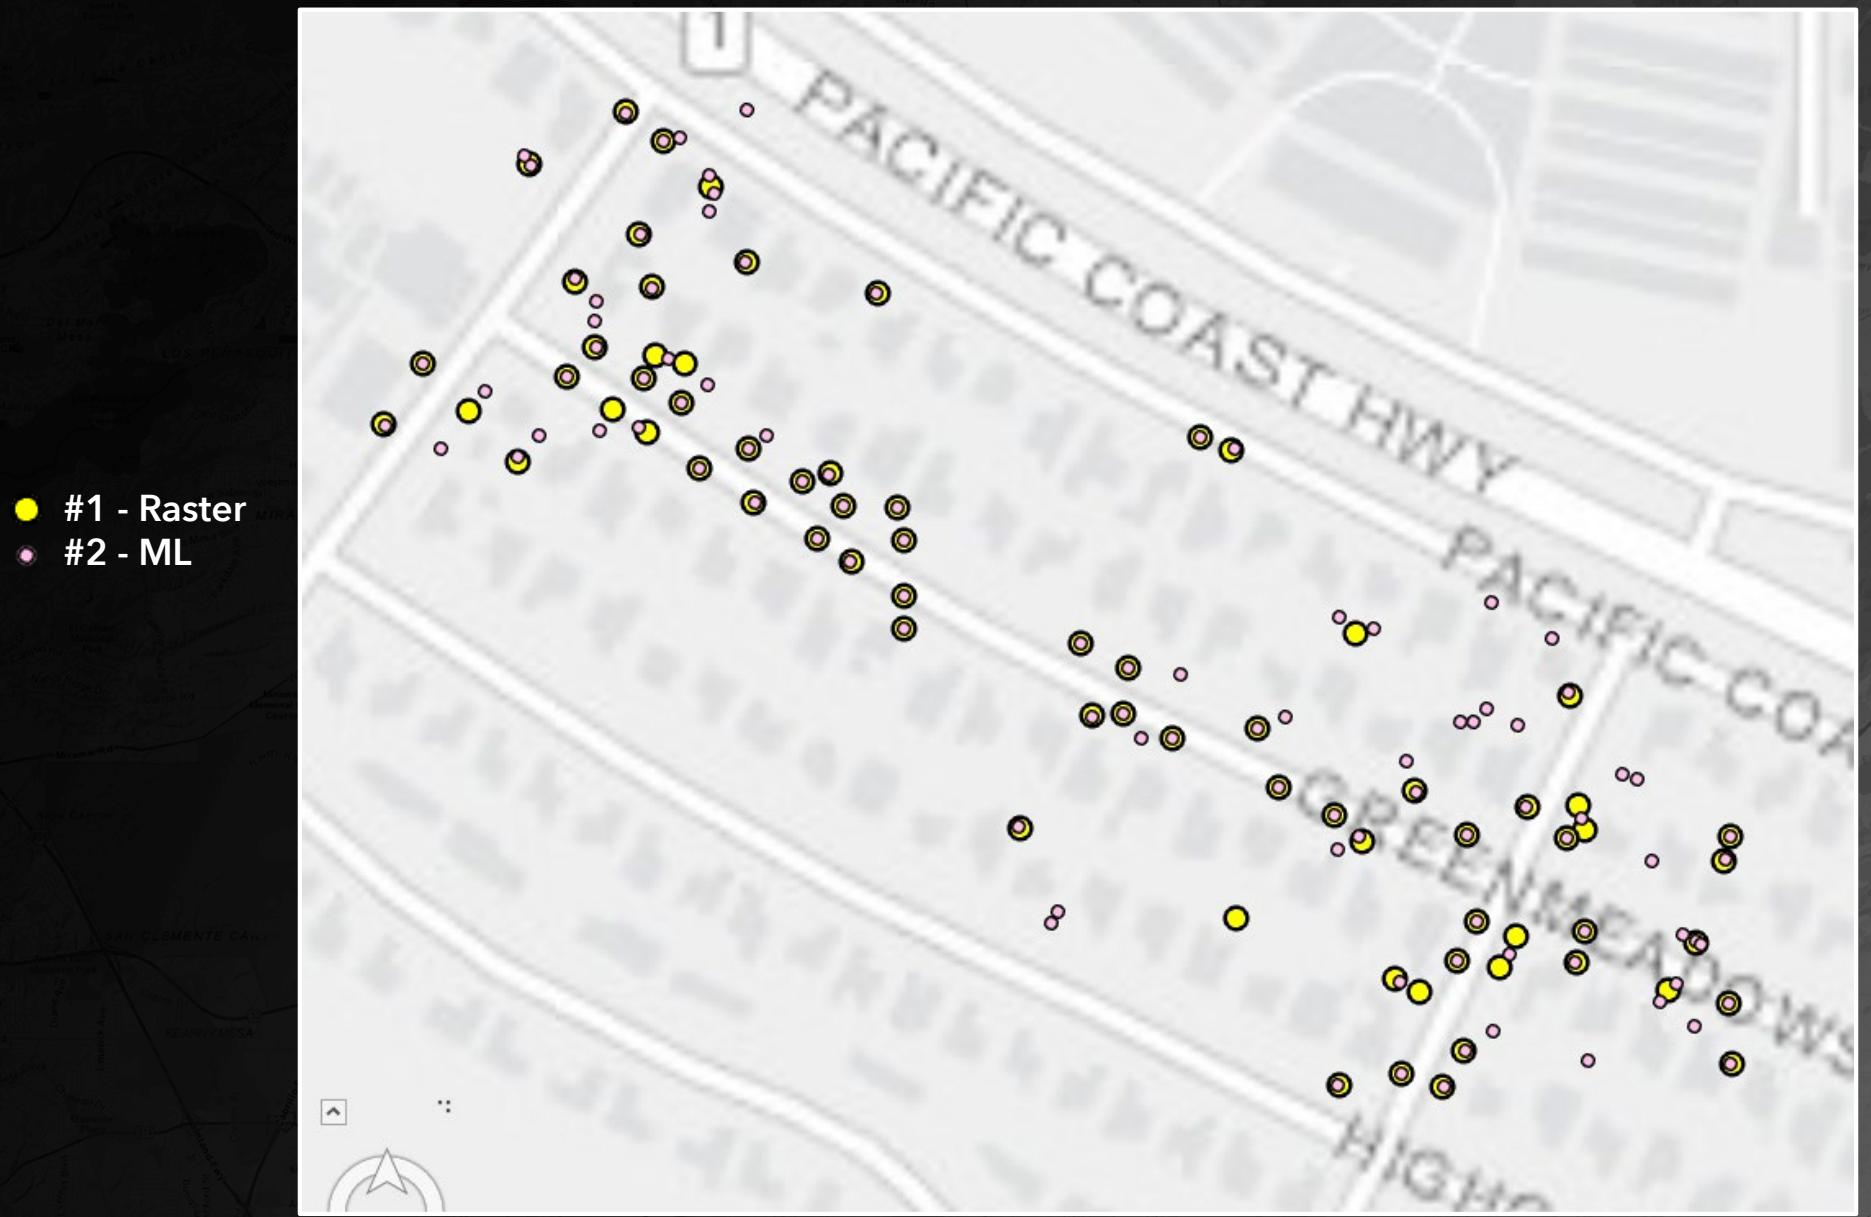
Results (2D and 3D)



Results (2D and 3D)

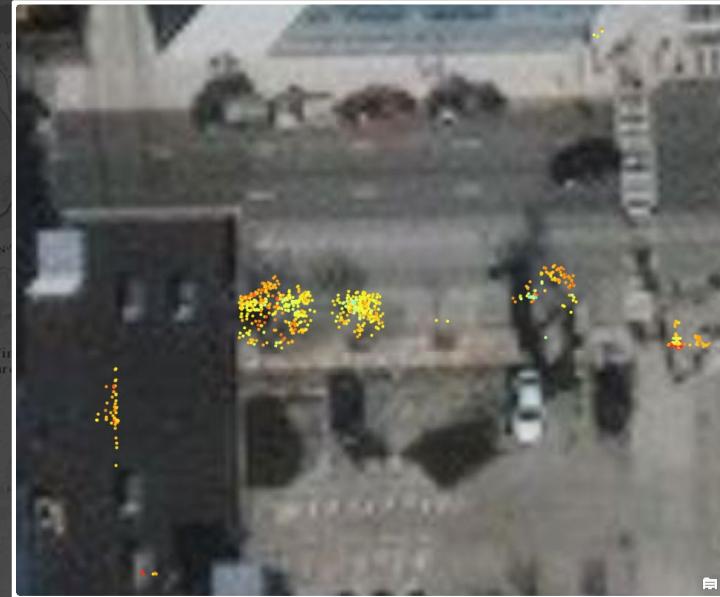
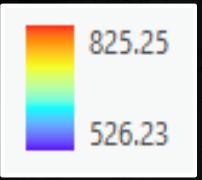


Results



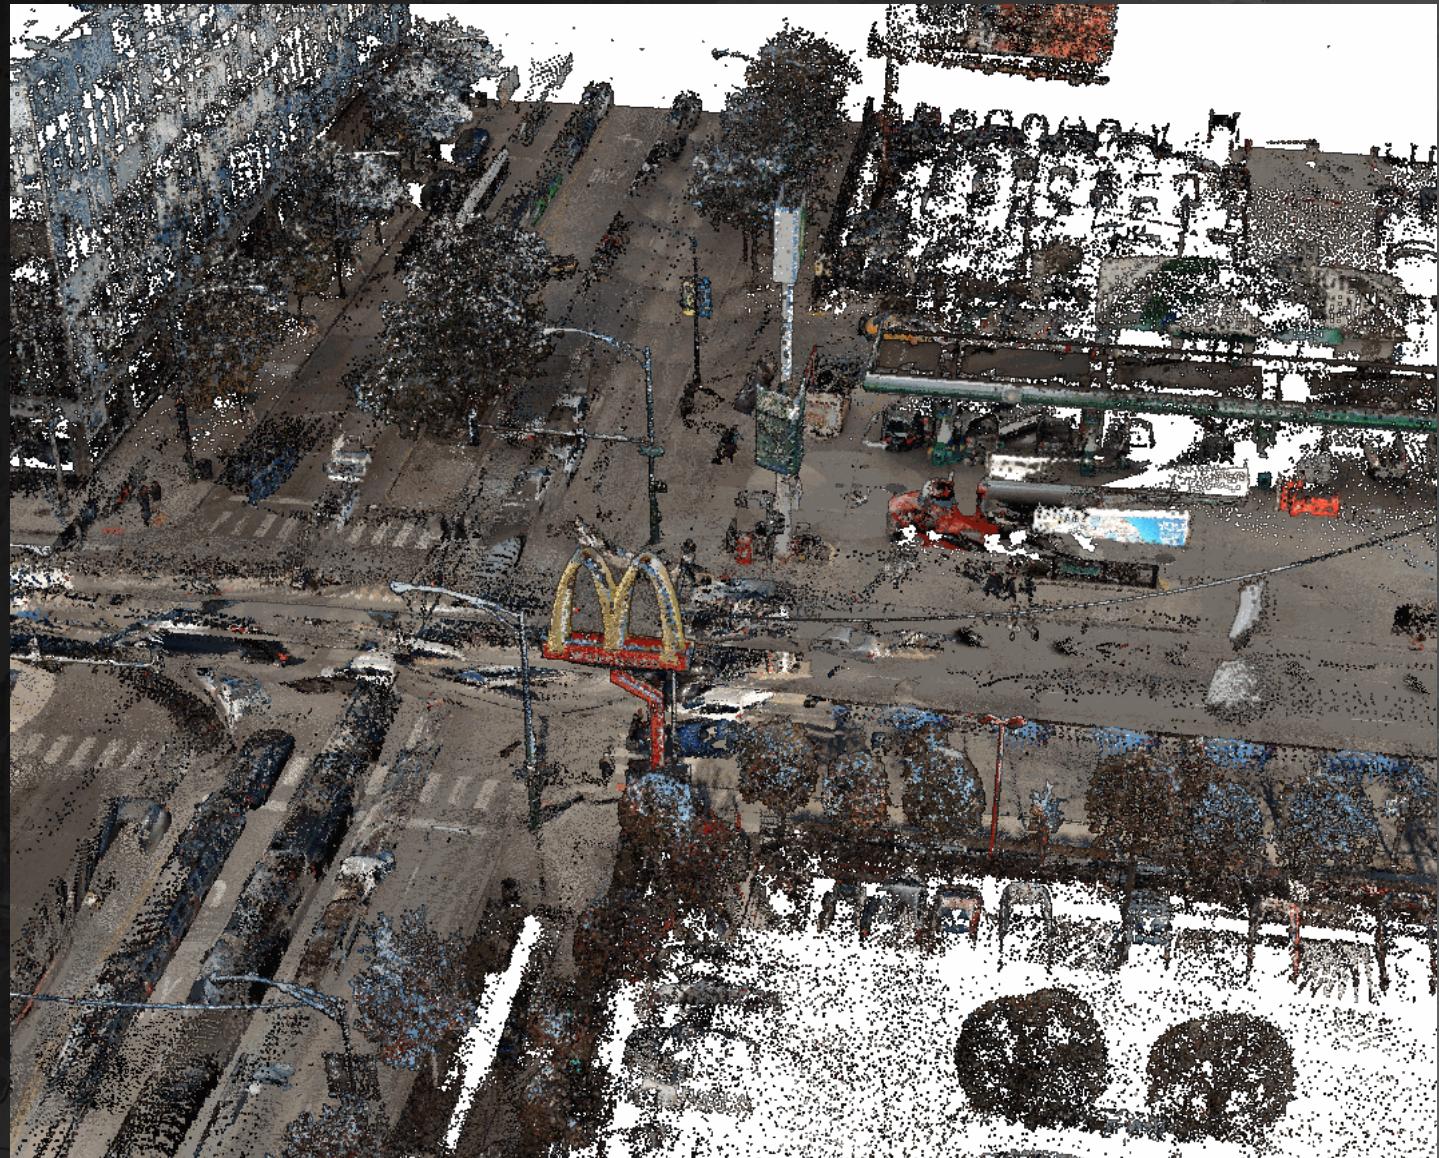
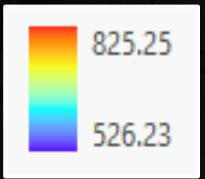
Chicago data - Trees - overfitting for season

- Number of LAS files: 24
- Number of LAS points: 60 million
- XYZ units: US foot
- Average point spacing: 0.26 feet
- Classes
 - 0 = Other
 - 1 = Trees



Chicago data - Trees - overfitting for season

- Number of LAS files: 24
- Number of LAS points: 60 million
- XYZ units: US foot
- Average point spacing: 0.26 feet
- Classes
 - 0 = Other
 - 1 = Trees





Netherlands data - Traffic lights

Schiedam

- Number of LAS files: 12
- Number of LAS points: 42.3 million
- XYZ units: Meters
- Average point spacing: 0.03 meters
- Classes
 - 0 = Other
 - 1 = Traffic lights



Amsterdam

- Number of LAS files: 35
- Number of LAS points: 136.6 million
- XYZ units: Meters
- Average point spacing: 0.04 meters
- Classes
 - 0 = Other
 - 1 = Traffic lights



Rijnstraat, Amsterdam

- Number of LAS files: 45
- Number of LAS points: 188.6 million
- XYZ units: Meters
- Average point spacing: 0.03 meters
- Classes
 - 0 = Other
 - 1 = Traffic lights



Feature extraction pipelines (#2 - ML)

Schiedam

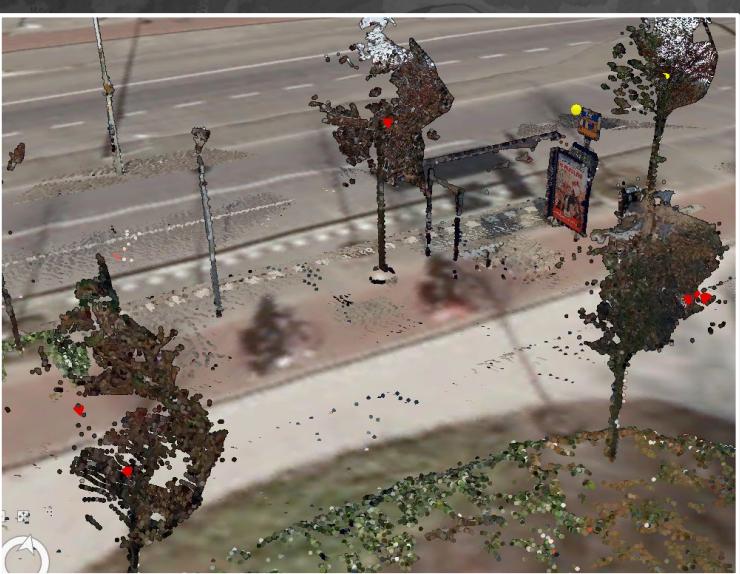
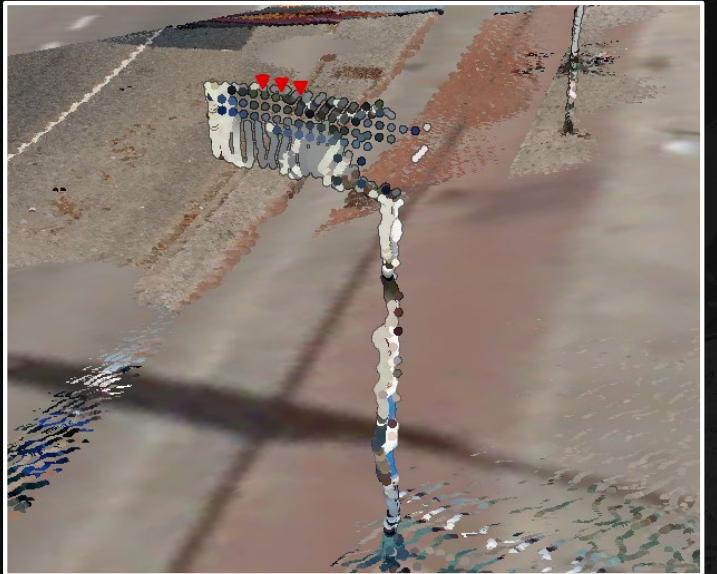


Amsterdam



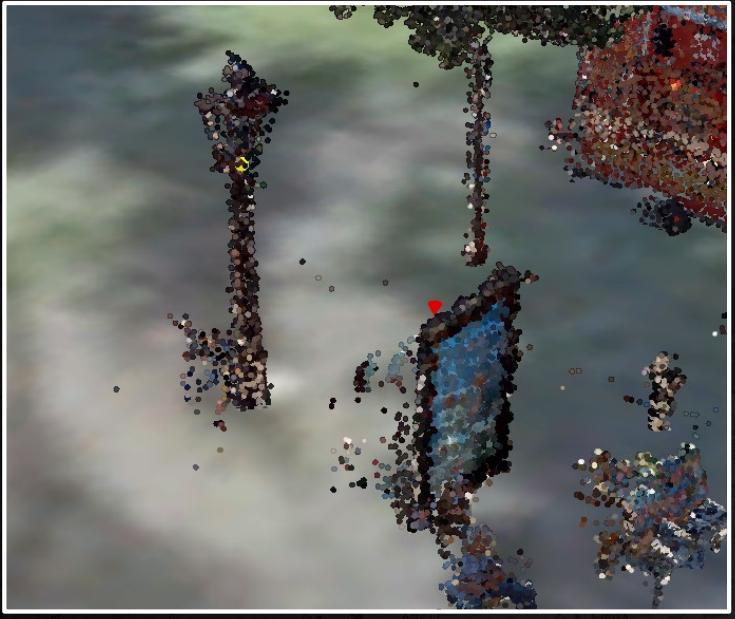
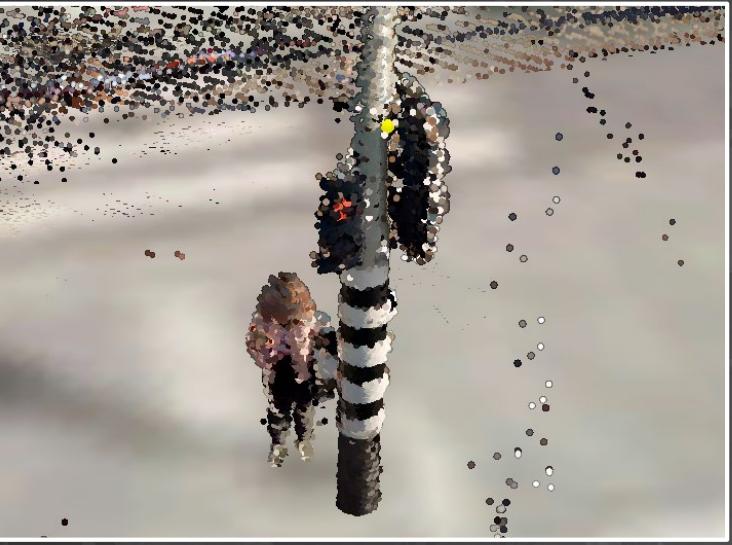
Schiedam - Traffic lights (3D)

- #1 - Traffic light
- #2 - Noise



Amsterdam - Traffic lights (3D)

- #1 - Traffic light
- #2 - Noise



Amsterdam - Traffic lights (3D)

- #1 - Traffic light
- #2 - Noise



Extracted Objects, after DBScan:

DBSCAN 30pts/0.5m

Precision = .655

Recall = .704

F1-Score = .679



Want to know more?



medium.com/geoaicloud

- Reconstructing 3D buildings from aerial LiDAR with AI
- 3D cities: Deep Learning in three-dimensional space
- PointCNN: replacing 50,000 man hours with AI
- Object extraction from Mobile LiDAR point clouds with Machine Learning
- ...and much more